



CLOUDTRUCKS

Data Scientist — Logistics Intelligence & Machine Learning

ENGAGEMENT 18 Months	PROJECTED START Oct 26, 2026	PROJECTED END 2028
TARGET RATE \$55-\$68/hr W2	WORK MODE Hybrid	EXTENSION No
LOCATION San Francisco, CA	CLIENT REF. ID CLT-2026-IPR-517.	WORK AUTHORIZATION Valid US EAD

01 / CONTEXT

The assignment

CloudTrucks operates a technology platform for owner-operators and small trucking carriers, spanning load booking, dispatch, payments, compliance, and business management.

The Data Science function supports operational and product decisions where transportation data, driver behavior, load economics, and marketplace signals need to be converted into usable predictions and recommendations.

PRODUCTION-FACING FOCUS Load recommendations • schedule optimization • operational forecasting • product-behavior measurement

02 / OWNERSHIP

What the consultant will own

MODEL & DATA • Build and evaluate predictive models using transportation, load, driver, customer, and operational data. • Develop features representing historical load behavior, pricing, route characteristics, driver preferences, utilization, and operational outcomes. • Support ranking/recommendation use cases for load selection and related driver-facing product decisions. • Analyze model outputs against business KPIs such as revenue opportunity, utilization, booking behavior, and operational efficiency. • Establish leakage-safe training, validation, and holdout strategies for production-oriented modeling.	PRODUCTION & IMPACT • Compare baseline statistical/ML approaches with tree-based models and document model trade-offs. • Use Python and SQL for exploratory analysis, feature engineering, model development, and validation. • Apply model explainability techniques where stakeholders need to understand drivers behind predictions or recommendations. • Work with engineers to translate validated models into production-ready scoring or data workflows. • Monitor model performance post-deployment and investigate data drift, degradation, or unexpected business behavior. • Design experiments and measurement frameworks to determine whether model-driven product changes improve business outcomes.
--	---

03 / HOW THE WORK FITS TOGETHER

Role architecture



ROLE PROFILE

The consultant will work with Product, Operations, and Engineering stakeholders to turn operational questions into measurable models and analyses — not operate as a research-only data science resource.

01 OPERATIONAL QUESTION	02 DATA + FEATURES	03 MODEL + VALIDATION	04 PRODUCT / OPERATIONS	05 MEASURE + MONITOR
-------------------------------	--------------------------	-----------------------------	-------------------------------	----------------------------

04 / NON-NEGOTIABLES

Required qualifications

CORE QUALIFICATIONS	EXECUTION & COLLABORATION
• Bachelor's or Master's degree in Data Science, Computer Science, Statistics, Business Analytics, Mathematics, or a related quantitative discipline. • 1-3 years of hands-on experience with Python-based data science or equivalent graduate-level project experience. • Strong Python skills with Pandas, NumPy, and Scikit-learn. • Strong SQL capability for joining, aggregating, and analyzing production-style datasets. • Demonstrated experience building and validating supervised machine-learning models. • Practical understanding of feature engineering, model evaluation, cross-validation, and holdout testing.	• Experience with regression, classification, or gradient-boosting methods. • Ability to explain model results to non-technical Product or Operations stakeholders. • Git/GitHub experience and comfort working in a collaborative engineering environment.

05 / DIFFERENTIATORS

Preferred signals

◆ XGBoost or another gradient-boosting framework.	◆ SHAP or another model-explainability method.	◆ Experience with recommendation, ranking, optimization, marketplace, or logistics problems.
◆ Exposure to experimentation/A-B testing and KPI measurement.	◆ Experience with dbt, Airflow, or cloud data platforms.	◆ Transportation, freight, marketplace, or supply-chain analytics exposure.

06 / EVALUATION PATH

Interview / submission flow

1	Client Data Science Interview Technical discussion covering feature engineering, model validation, leakage, evaluation metrics, and business interpretation.
2	Applied Case / Working Session Candidate receives a logistics-oriented problem and walks through data preparation, modeling approach, validation strategy, and expected business impact.
3	Cross-Functional Interview Discussion with Product/Operations and Engineering stakeholders focused on communication, experimentation, and production collaboration.

ROLE PROFILE

4

Client Manager

Final fit, availability, engagement expectations, and project alignment.