

SOWMYA KONDETI
DATA ENGINEER

Texas | (469) 920-0698 | kondetisowmya24@gmail.com

PROFESSIONAL SUMMARY

Data Engineer with 4+ years of experience designing and maintaining scalable ETL/ELT pipelines, data lakes, and cloud data warehouses across AWS and Azure. Experienced with Python, SQL, PySpark, AWS Glue, S3, EMR, Redshift, Lambda, Athena, IAM, CloudWatch, Airflow, Databricks, and Snowflake. Strong background in data quality, governance, security, performance optimization, CI/CD, and production troubleshooting.

TECHNICAL SKILLS

- **Programming:** Python, SQL, T-SQL, PySpark, Bash
- **Data Engineering:** ETL/ELT, Data Pipelines, Data Ingestion, Data Transformation, Data Validation, Data Quality, Batch Processing, Event-Driven Processing
- **Big Data & Orchestration:** Apache Spark, PySpark, Azure Databricks, Apache Airflow, Delta Lake, Kafka, RabbitMQ, AutoSys
- **Cloud:** Microsoft Azure (Azure SQL, Storage Accounts, Azure Functions, App Services, Service Bus, Application Insights, Azure Monitor), AWS (S3, RDS, EC2, Glue, VPC, CloudFormation)
- **Databases & Warehousing:** SQL Server, Azure SQL, PostgreSQL, MySQL, Snowflake
- **DevOps:** Azure DevOps, Git, CI/CD, Docker, Jenkins, Datadog, Application Insights
- **BI, Development & Testing:** Power BI, REST APIs, FastAPI, PyTest, Unit Testing, Integration Testing, Debugging, SDLC

PROFESSIONAL EXPERIENCE

Data Engineer | Reidy Inc

Oct 2024 - Present

- Built scalable Python, SQL, and Apache Airflow ETL/ELT pipelines to ingest multi-GB MLS and CoreLogic datasets from SFTP into Hive-partitioned Azure Data Lake Storage Gen2, incorporating file sensors, automated retries, and idempotent processing.
- Developed PySpark and Apache Spark transformations on Azure Databricks to standardize schemas, deduplicate records, and integrate property listings, sales, and comparable datasets.
- Implemented a Databricks Delta Lake medallion architecture with bronze, silver, and gold layers to provide analytics-ready data for appraisal, smart-search, and reporting services.
- Orchestrated SFTP ingestion, Spark processing, and Snowflake loading with Azure Data Factory, supporting historical backfills, incremental loads, monitoring, and failure alerts.
- Created Snowflake fact and dimension tables using SQL to support property valuation, comparable-sales analytics, and investment reporting.
- Implemented automated Python, PySpark, and SQL data-quality checks for schema validation, null values, duplicate records, and source-to-target reconciliation.
- Optimized Spark and Databricks workloads using partition tuning, broadcast joins, caching, and cluster autoscaling.
- Managed code and automated deployments using Git, Azure DevOps, CI/CD, and Docker across development, UAT, and production environments.

Data Engineer | Apple

Feb 2024 - Oct 2024

- Built scalable AWS ETL/ELT pipelines using Python, PySpark, AWS Glue, and Amazon S3 to process high-volume structured, semi-structured, and unstructured datasets.
- Used Amazon EMR to run large-scale Spark transformations and historical backfills, applying partition tuning, broadcast joins, and caching to improve performance.
- Developed event-driven AWS Lambda functions to validate incoming files, capture metadata, and trigger downstream Glue processing workflows.
- Managed dataset metadata through the AWS Glue Data Catalog and used Amazon Athena to validate and analyze curated data stored in S3.
- Loaded curated datasets from S3 into Amazon Redshift and developed optimized SQL models for reporting and downstream analytics.
- Applied least-privilege IAM roles, encryption, and access controls to protect sensitive data and support governance requirements.
- Monitored Glue, EMR, Lambda, and pipeline workloads through Amazon CloudWatch and Datadog, troubleshooting failures, latency, and delayed data loads.

- Collaborated with analysts, data architects, platform teams, and security teams to translate business requirements into AWS data solutions and participated in Git code reviews.

Data Engineer | Accenture (Bristol Myers Squibb)

Oct 2021 - Jul 2022

- Developed and optimized 150+ SQL and PL/SQL scripts with a six-engineer team, improving data-processing performance by 20%.
- Automated data-ingestion and batch-scheduling workflows using AutoSys, reducing manual scheduling effort by 20% and improving job reliability.
- Implemented Slowly Changing Dimension Type 2 in Snowflake to preserve historical data changes and support audit traceability for pharmaceutical compliance.
- Consolidated four product trackers and more than 10 KPI reports into Power BI dashboards, improving reporting consistency and stakeholder visibility.
- Built SQL and PL/SQL reconciliation checks to compare source and target record counts, identify load exceptions, and improve data accuracy across scheduled batch pipelines.
- Monitored and troubleshooted AutoSys job dependencies and Snowflake data loads, resolving failures to meet scheduled reporting timelines.

Data Engineer | Adobe

May 2020 - Oct 2021

- Built batch and incremental ETL pipelines using Python, SQL, and Apache Airflow to ingest application and product-usage data into Amazon S3 and PostgreSQL.
- Developed PySpark transformations for structured and semi-structured datasets, applying schema standardization, cleansing, enrichment, and partitioning for downstream analytics.
- Designed dimensional data models and curated fact and dimension tables to support scalable reporting, reusable analytics, and efficient incremental processing.
- Developed Kafka and RabbitMQ consumers for event-driven ingestion, implementing schema validation, deduplication, checkpointing, and failure handling.
- Optimized complex PostgreSQL queries, indexes, views, and stored procedures to improve data-access performance and reduce reporting latency.
- Implemented source-to-target reconciliation, automated data-quality checks, pipeline monitoring, and CI/CD controls to support reliable production releases.

EDUCATION

Master of Science in Advanced Data Analytics | University of North Texas, Denton, TX

Jul 2022 - Dec 2023