



MACHINE LEARNING ENGINEER

Healthcare AI & Data Engineering



CLIENT	EMIDS	LOCATION (CLIENT)	Nashville, TN
WORK ARRANGEMENT	Remote (U.S.)	EMPLOYMENT TYPE	W2 — Staffing Partner
ASSIGNMENT TERM	24 Months	START	Immediate / Upon Selection
COMPENSATION	\$48–\$58 / hr W2 (DOE)	EXTENSION	Possible — Subject to Funding
REQ. ID	EM-2026-00UT62-ML	REQ. WORK LOCATION	Nashville, TN

CLIENT / PROJECT CONTEXT

The project supports development of machine-learning capabilities against healthcare datasets and business workflows where data quality, model reliability, explainability, and controlled deployment are critical. The engineering team works with structured and unstructured healthcare information, develops analytical and predictive models, evaluates model behavior, and translates research-grade approaches into repeatable engineering workflows.

The role is intended for an engineer who can work directly in Python and SQL, develop and evaluate machine-learning models, perform feature engineering and data preparation, and contribute to the integration of analytical models into production-oriented healthcare technology environments.

WHAT YOU WILL OWN

- Develop, train, validate, and evaluate supervised and unsupervised machine-learning models using Python and established ML frameworks.
- Prepare healthcare and operational datasets for modeling, including data cleaning, transformation, validation, feature engineering, and exploratory analysis.
- Build reproducible analytical workflows using Python, Pandas, NumPy, scikit-learn and related libraries.
- Apply appropriate statistical and machine-learning evaluation methods for classification, prediction, model comparison, and error analysis.
- Work with SQL-based data sources and support extraction, transformation, aggregation, and validation of model-ready datasets.
- Investigate model performance issues including overfitting, class imbalance, data leakage, feature quality, and generalization.
- Contribute to deep-learning workloads using PyTorch or TensorFlow where applicable to the project.
- Support development of NLP/LLM-based capabilities involving healthcare text and other unstructured information.
- Contribute to retrieval-augmented generation workflows, including document preparation, retrieval, evidence grounding, and evaluation where applicable.
- Develop model explanations and analytical outputs that can be reviewed by technical and healthcare-domain stakeholders.
- Support experimentation, model comparison, hyperparameter tuning, cross-validation, and documentation of results.
- Work with engineering teams to package analytical components for downstream applications and services.



- Maintain source-controlled analytical code and technical documentation using Git-based development practices.
- Participate in code reviews, technical discussions, defect resolution, and sprint-level delivery activities.
- Communicate model assumptions, analytical findings, limitations, and performance results to engineers and non-technical stakeholders.

REQUIRED QUALIFICATIONS

Education

- Bachelor's degree in Computer Science, Data Science, Statistics, Mathematics, Engineering, or a related quantitative discipline.
- Master's degree in Data Science, Computer Science, Machine Learning, or a related field is preferred.

Technical

- Strong hands-on Python programming experience.
- Working knowledge of SQL and relational data.
- Hands-on experience with Pandas, NumPy and scikit-learn.
- Practical understanding of supervised machine learning, model evaluation, feature engineering, and statistical analysis.
- Experience building and evaluating classification or predictive models.
- Experience with at least one deep-learning framework such as PyTorch or TensorFlow.
- Experience using Git in a collaborative development environment.
- Ability to work with large, imperfect, or heterogeneous datasets.

Healthcare / Data

- Experience working with healthcare, clinical, claims, research, insurance, patient, provider, or other health-related datasets is strongly preferred.
- Understanding of healthcare data quality and the need for traceability, privacy, and controlled use of sensitive data is preferred.

PREFERRED QUALIFICATIONS

- Experience with NLP, transformer models, LLMs, RAG, or other generative-AI technologies.
- Experience with explainable AI, model interpretability, uncertainty estimation, or counterfactual analysis.
- Experience with PySpark or distributed data processing.
- Experience with Power BI or Tableau for analytical communication.
- Exposure to Microsoft Azure or another major cloud platform.
- Experience with Docker or containerized ML environments.
- Exposure to ML pipelines, experiment tracking, model monitoring, or MLOps practices.
- Experience converting research or prototype ML work into reusable engineering components.
- Healthcare interoperability exposure, including familiarity with healthcare data standards, is a plus.

CANDIDATE EVALUATION

The client will evaluate candidates primarily on:

- Python and machine-learning fundamentals.
- Ability to work with and prepare real-world datasets.



- Model development, validation, and performance analysis.
- Understanding of statistical reasoning and model limitations.
- Healthcare/research data experience.
- Ability to explain technical decisions clearly.
- Practical exposure to deep learning, NLP, LLM or RAG systems.
- Ability to work within a software-engineering delivery process rather than only an academic/research environment.

CLIENT INTERVIEW PROCESS

ROUND 1 — Emids Technical Interview

ROUND 2 — ML/AI Engineering Interview

ROUND 3 — Client Manager / Delivery Discussion

FINAL STEP — Staffing Partner Compliance / Onboarding

Employment verification, employment eligibility/work authorization documentation, background screening as required by the client, W2 onboarding, and project start coordination.

WORK AUTHORIZATION

Candidates must be legally authorized to work in the United States for the duration of the assignment. Potentially workable categories may include:

- U.S. Citizen
- Permanent Resident / Green Card
- F-1 OPT
- F-1 STEM OPT Extension
- H-4 EAD
- Other valid employment authorization accepted by the staffing employer and client
- H-1B candidates subject to an employer-sponsored transfer and client/staffing-partner approval

International students on F-1 status must have valid employment authorization appropriate to the assignment. CPT eligibility must be confirmed by the candidate's DSO and the staffing employer before placement.

Employment authorization is subject to verification and client/vendor compliance requirements. No placement is contingent solely on future immigration approval.

ASSIGNMENT CONDITIONS

- Remote position; candidate must work from a U.S. location approved for the assignment.
- Standard W2 employment through the staffing partner.
- Initial assignment is 24 months.
- Extension may be considered based on client funding, project continuation, performance, and business requirements.
- Start is targeted as soon as the selected candidate completes the client and staffing-partner onboarding process.
- Compensation is dependent on demonstrated technical experience and final approved bill/rate structure.