

Ajay Gundla

[LinkedIn](#) | ajaygundla.896@gmail.com | +15182506128

PROFESSIONAL SUMMARY:

- Experience in data analytics and engineering, specializing in SQL, DBT, and Control M orchestration, transforming raw datasets into actionable insights to drive strategic business decisions across multiple industries.
- Built complex ETL pipelines integrating cloud platforms, ensuring seamless ingestion, transformation, and loading of structured and semi-structured data with high accuracy and scalability for enterprise reporting and analytics.
- Created modular DBT models to standardize data transformations, enabling automated testing and documentation while improving data consistency, reducing downstream data errors, and accelerating analytics delivery timelines.
- Developed and optimized intricate SQL queries over large datasets to extract, transform, and aggregate business-critical metrics, improving query performance and reducing runtime for critical operational reports.
- Executed batch job scheduling and orchestration in Control M, ensuring timely processing, alert management, and automated error handling for mission-critical data pipelines supporting multiple departments.
- Formulated data quality frameworks to continuously validate, reconcile, and monitor data integrity across multiple source systems, proactively identifying discrepancies and reducing manual reconciliation efforts.
- Shaped data pipeline architecture and implementation strategies across AWS, Azure, and GCP, enhancing performance, reliability, and maintainability while aligning with cloud-native best practices.
- Constructed scalable, reusable data models that supported analytics for finance, healthcare, and insurance industries, enabling faster insights and supporting complex decision-making processes with high data fidelity.
- Devised end-to-end data solutions including requirement gathering, source system analysis, pipeline design, and deployment, providing measurable improvements in reporting accuracy and operational efficiency.
- Rolled out automated CI/CD pipelines for DBT transformations, integrating Git version control, streamlining deployment, and reducing manual intervention across data engineering teams.
- Partnered with business stakeholders to translate functional requirements into actionable technical solutions, enabling analytics teams to generate high-value insights for strategy and operational decision-making.
- Mentored junior engineers and analysts in SQL optimization, DBT modeling, and Control M orchestration, fostering a culture of knowledge sharing and technical excellence across cross-functional teams.
- Enhanced performance of large-scale SQL transformations through indexing strategies, query refactoring, and execution plan optimization, reducing compute costs while improving operational throughput.

- Applied Agile principles to manage sprint planning, backlog grooming, and iterative delivery of data engineering initiatives, ensuring consistent alignment with business objectives and timely project completion.
- Improved observability of data pipelines by implementing logging, monitoring, and alerting solutions that tracked job execution, data anomalies, and performance bottlenecks in real-time.
- Crafted documentation for data lineage, transformation logic, and job orchestration workflows, ensuring comprehensive knowledge transfer and simplifying maintenance for technical and non-technical stakeholders.
- Structured robust error-handling mechanisms for Control M jobs and ETL processes, reducing manual intervention and enabling automated recovery and retry for critical workflows.
- Developed dashboards and metrics reports that empowered stakeholders to track pipeline health, business KPIs, and transformation performance across multiple data sources and platforms.
- Executed migration of on-premise ETL workflows to cloud-native architectures in AWS, Azure, and GCP, enhancing scalability, reliability, and cost efficiency across enterprise data systems.
- Initiated proactive performance tuning and query optimization initiatives that reduced runtime, minimized storage costs, and improved responsiveness for end-users accessing analytics applications.
- Deployed advanced analytics frameworks to integrate structured and semi-structured data sources, enabling more comprehensive insights for product, marketing, and operational teams.
- Guided cross-functional teams in defining and implementing data governance standards, ensuring compliance, security, and regulatory alignment in multi-industry deployments.
- Conceptualized and implemented data model versioning and CI/CD processes in Git, providing traceability, rollback capabilities, and consistency for enterprise-scale transformations.
- Systematized testing procedures for DBT models, introducing unit tests, schema checks, and automated validation to maintain high data quality standards and reduce regression issues.
- Facilitated collaboration between engineering, analytics, and business teams to prioritize pipeline development, optimize workflows, and ensure alignment with strategic initiatives.
- Laid the foundation for predictive and prescriptive analytics by integrating clean, structured data pipelines with advanced modeling frameworks, enabling forward-looking business insights and operational planning.

TECHNICAL SKILLS:

- **Big Data Ecosystem:** Spark, Hive, Kafka, HDFS, Airflow, Spark Streaming, NiFi, Sqoop, Flume, MapReduce, Yarn, Zookeeper
- **Programming languages:** Java, Python, Scala, Pyspark, Spring boot, SQL, Unix, T-SQL
- **ETL:** Semarchy xDM, Informatica Power Center 10.x/9.6, Talend 5.6, AWS Glue, SSIS
- **Cloud Environment:** Azure (Databricks, Data lake, Data factory, Blob storage, SQL, HDInsight, Fabric, CosmosDB), GCP (Bigquery, Cloud storage, DataProc, Functions, GKE, Spanner), AWS (EMR, Redshift, S3, Glue, Lambda, Snowflake, Athena, RDS, EC2, Cloud Watch, VPC, SQS, Jenkins)
- **Reporting Tools:** PowerBI, EHR, Tableau, Semarchy
- **Containerization and Orchestration Tools:** Docker, Kubernetes, Docker Hub, Jenkins
- **Databases & Tools:** MySQL, MS SQL Server 2014/2012, Oracle 11g/10g, Teradata 15/14, Erwin 8/9
- **Python Libraries / Machine Learning Algorithms:** Pandas, Scikit-Learn, NumPy, Matplotlib, Decision Tree, Logistic Regression, Random Forest, Gradient Boosting
- **Data Warehousing & BI:** Cloudera, SSIS, Star and Snowflake schema, Splunk, Facts and Dimensions tables
- **Software Life Cycle/ Methodologies:** Agile models, SDLC, Waterfall

- **Certifications:** Azure AZ-900, ServiceNow CAD and SaaS Analyst, Tableau certified, Azure DP-203

PROFESSIONAL EXPERIENCE:

Northern Trust: July 2023 – Current (Location: Chicago, IL)

Role: Sr. Data Analyst

Responsibilities:

- Built scalable AWS-based ETL pipelines to extract and transform transactional and reference data, ensuring high integrity and accuracy for financial reporting and regulatory compliance across global operations.
- Developed DBT modular transformations for complex banking datasets, standardizing metrics across departments and enabling automated testing and validation of all data models.
- Executed Control M batch jobs for nightly ETL workflows, including scheduling, error handling, and alerting to ensure uninterrupted data availability for downstream reporting and analytics teams.
- Applied SQL optimization techniques to reduce query runtime for multi-terabyte datasets, enabling faster operational reporting and analytics across risk, finance, and treasury functions.
- Conceptualized and implemented data pipeline architecture using AWS Redshift, S3, Glue, and Lambda, achieving automated ingestion, transformation, and analytics-ready dataset creation for multiple business units.
- Guided cross-functional teams to align data engineering solutions with business requirements, improving collaboration, reducing development bottlenecks, and ensuring timely delivery of critical financial reports.
- Constructed comprehensive logging, monitoring, and alerting mechanisms in Control M and CloudWatch, improving observability and reducing manual intervention for batch ETL failures.
- Rolled out automated CI/CD pipelines for DBT models using Git, enabling version control, rollback capability, and repeatable deployments across development, testing, and production environments.
- Improved data quality frameworks with automated validation checks, reconciliation scripts, and anomaly detection, reducing manual review time while ensuring compliance with internal and regulatory standards.
- Partnered with data analysts to translate functional requirements into optimized SQL transformations and DBT models, supporting complex financial dashboards and operational decision-making.
- Laid the foundation for advanced analytics by integrating structured and unstructured banking data from multiple sources, enabling richer insights for risk and treasury management.
- Mentored junior engineers on SQL performance tuning, DBT development, and Control M orchestration best practices, accelerating onboarding and improving overall team productivity.
- Introduced schema evolution handling and version control in Redshift tables to manage changing business requirements without impacting downstream analytics.
- Streamlined job dependency management in Control M for ETL workflows spanning multiple AWS accounts, reducing delays and increasing reliability of nightly processing.
- Enhanced reporting efficiency by creating reusable DBT macros, templates, and standardized transformations, improving maintainability and reducing redundant development effort.
- Executed data reconciliation between legacy and new ETL pipelines, ensuring consistency, accuracy, and completeness of financial and operational data across systems.

- Directed the implementation of encryption, access controls, and compliance standards in AWS data pipelines to secure sensitive banking data across environments.
- Coordinated with business users to refine KPIs and metrics definitions, translating operational needs into actionable SQL queries and pipeline transformations.
- Improved cost efficiency by optimizing S3 storage, partitioning strategies, and Redshift table designs, balancing performance and operational costs effectively.
- Championed automation of Control M alerts, retries, and failure notifications, reducing operational overhead while increasing the reliability of batch processes.
- Built end-to-end observability dashboards integrating Redshift, CloudWatch, and Control M metrics, providing real-time insights into pipeline health and performance.
- Produced data lineage documentation, mapping sources to transformed datasets to enhance transparency, traceability, and auditability of financial data.
- Deployed serverless AWS Lambda functions to handle real-time transformations and error handling, complementing batch ETL workflows and enhancing system responsiveness.
- Fine-tuned Redshift query execution plans and storage configurations to support large-scale aggregation, filtering, and joins for reporting and analytical applications.
- Systematized testing of DBT transformations with unit tests and snapshot tests, ensuring high-quality and accurate datasets for downstream financial analyses.
- Enabled cross-team knowledge sharing through documentation, workshops, and collaborative design sessions, fostering alignment between analytics, engineering, and business teams.

Environment: AWS, Redshift, S3, Glue, Lambda, DBT, Control M, SQL, Git, Python, CloudWatch, Airflow, Snowflake, Spark, Jenkins, Tableau, Docker, PySpark, Linux.

Johnson & Johnson Health: September 2020 – June 2023 (Location: New Brunswick, NJ)

Role: Sr Data Analyst

Responsibilities:

- Developed Azure-based data pipelines to aggregate patient, clinical, and operational datasets, ensuring high-quality data availability for analytics, reporting, and strategic decision-making across multiple business units.
- Applied advanced DBT modeling techniques to structure healthcare datasets, enabling modular, testable, and version-controlled transformations with improved maintainability.
- Rolled out Control M scheduling for batch ingestion workflows, orchestrating cross-system dependencies and reducing processing time for critical healthcare reporting.
- Executed complex SQL queries across Azure SQL Data Warehouse to extract clinical metrics, optimize reporting, and generate actionable insights for patient care and operational efficiency.
- Devised data architecture solutions incorporating Azure Data Lake, Synapse Analytics, and Databricks to enable scalable storage, transformation, and analytics for large-scale health datasets.
- Enhanced ETL reliability by implementing error handling, logging, and automated retries in Control M, reducing downtime and ensuring consistent data delivery for analytics teams.
- Mentored junior engineers on cloud-based pipeline design, DBT modeling, and Control M orchestration, improving team proficiency and accelerating project delivery.
- Conceptualized and deployed data validation frameworks to reconcile multiple sources, ensure integrity, and reduce manual verification in sensitive health data environments.

- Built dashboards integrating Azure Synapse, Power BI, and Python, enabling clinical and operational stakeholders to monitor KPIs and make data-driven decisions efficiently.
- Introduced CI/CD practices for DBT deployments using Git and Azure DevOps, standardizing pipeline releases and reducing errors during production updates.
- Partnered with healthcare analysts to translate business requirements into scalable SQL and DBT solutions, providing timely and accurate reporting on clinical outcomes and operational metrics.
- Streamlined data ingestion from electronic health records (EHR) and other clinical systems, implementing robust pipelines that supported analytical reporting and compliance audits.
- Applied query optimization and partitioning strategies for large healthcare datasets in Azure SQL DW, reducing processing times and improving end-user report responsiveness.
- Developed reusable DBT macros and standardized transformations to simplify healthcare data modeling and ensure consistent calculation of critical patient metrics.
- Enhanced observability of pipeline performance through automated logging, monitoring, and alerting across Azure Data Factory and Control M workflows.
- Executed migration of legacy on-premise ETL workflows to Azure cloud platforms, ensuring secure, scalable, and cost-effective data management for healthcare analytics.
- Championed data governance initiatives, implementing access controls, lineage tracking, and compliance measures for sensitive health information and regulatory adherence.
- Coordinated with cross-functional teams to optimize reporting pipelines and reduce latency in providing actionable insights to business and clinical stakeholders.
- Applied advanced Python scripting for data transformation, anomaly detection, and pipeline automation to reduce manual intervention and increase operational efficiency.
- Refined batch and real-time processing approaches to balance performance, cost, and reliability across Azure-based healthcare data pipelines.
- Facilitated alignment between analytics, engineering, and business teams, ensuring timely delivery of data solutions that directly supported clinical operations and research projects.
- Laid the groundwork for predictive analytics initiatives by providing clean, structured datasets suitable for machine learning and advanced modeling in healthcare contexts.

Environments: Azure Data Lake, Azure Synapse, Databricks, DBT, Control M, SQL, Git, Python, Power BI, Azure DevOps, Airflow, Pandas, Spark, Cosmos DB, Kubernetes, Linux, Docker, PySpark, Jupyter.

Equifax: November 2017 – August 2020 (Location: Atlanta, GA)

Role: Data Analyst

Responsibilities:

- Constructed data pipelines in GCP to ingest, cleanse, and transform insurance claims, policy, and customer datasets, ensuring high-quality analytics-ready data for reporting and operational insights.
- Developed DBT transformation models for standardizing and enriching insurance data, enabling consistent calculation of metrics across claims, underwriting, and risk analytics.
- Applied SQL optimization techniques across BigQuery datasets, improving performance of reporting queries while supporting cross-departmental analytics and operational decision-making.
- Executed Control M batch job orchestration to manage ingestion and transformation workflows, ensuring timely and reliable delivery of analytics datasets.
- Mentored junior analysts and engineers on GCP architecture, DBT modeling, and ETL best practices, building team capability and fostering knowledge transfer.

- Designed data quality checks, reconciliation scripts, and automated alerts to monitor data integrity, reducing errors and improving confidence in insurance reporting metrics.
- Partnered with business teams to translate functional requirements into technical solutions, ensuring accurate reporting of claims, policy renewals, and customer activity.
- Rolled out reusable DBT macros and modular transformations to streamline development and reduce redundant coding across insurance datasets.
- Enhanced pipeline observability by implementing monitoring dashboards in GCP Stackdriver, tracking execution metrics, failures, and performance anomalies across ETL workflows.
- Structured cloud-native GCP data architecture using BigQuery, Cloud Storage, and Dataflow to ensure scalable, secure, and performant ETL processing for insurance operations.
- Developed automated reconciliation processes to validate data consistency between source systems, transformation outputs, and reporting dashboards, ensuring high data reliability.
- Introduced version control and CI/CD processes for DBT models using Git, enabling safe deployment and rollback across development, test, and production environments.
- Refined SQL queries and table designs to optimize storage, join performance, and data retrieval for claims and underwriting dashboards.
- Applied Python and Pandas for data enrichment, anomaly detection, and preprocessing, complementing DBT transformations and reducing manual data preparation tasks.
- Executed incremental ETL designs to reduce processing time and resource consumption for large insurance datasets, maintaining operational efficiency and data freshness.
- Guided cross-team collaboration between engineering, analytics, and business stakeholders to prioritize pipeline development and ensure alignment with strategic objectives.
- Documented data flows, transformations, and business logic for insurance datasets, supporting transparency, auditing, and knowledge sharing across technical teams.
- Laid the foundation for advanced reporting and analytics by delivering clean, structured datasets suitable for visualization, dashboarding, and predictive modeling initiatives.
- Fine-tuned GCP job orchestration and scheduling to balance latency, throughput, and cost for daily and monthly reporting pipelines.
- Enhanced data security and compliance through role-based access, audit logging, and encryption of sensitive insurance and customer data across GCP environments.

Environments: GCP, BigQuery, Cloud Storage, Dataflow, DBT, Control M, SQL, Git, Python, Pandas, Airflow, Tableau, Jupyter, Spark, Kubernetes, Docker, Linux, PySpark, Looker.

Lowe's: December 2015 – October 2017 (Location: Mooresville, NC)

Role: Data Engineer / Migration Specialist

Responsibilities:

- Created a centralized Data Integration Platform on Azure Cloud to consolidate enterprise data from multiple business applications. This improved overall data integration efficiency and ensured seamless information flow across retail systems.
- Built scalable data ingestion pipelines using Talend ETL and Azure Data Factory (ADF). These pipelines automated data collection from various operational sources, improving consistency and timeliness.

- Managed complex data migration projects transferring historical retail and customer data from legacy systems to modern Azure-based platforms. Ensured complete data accuracy and minimal downtime during migration.
- Designed and orchestrated complex ETL workflows using Apache Airflow. This automation improved workflow reliability and streamlined data processing across distributed environments.
- Configured Azure SQL Database for structured data management and optimized query execution. These optimizations significantly reduced response times for analytics and reporting queries.
- Set up Azure Data Lake Storage to handle large volumes of semi-structured and unstructured enterprise data. The solution provided scalable storage and improved data accessibility for analytics teams.
- Integrated diverse data sources including sales transactions, customer profiles, and product catalogs into a unified data model. This consolidation improved data visibility and supported business decision-making.
- Implemented strong data security and compliance measures using Azure policies and encryption techniques. Ensured adherence to corporate governance standards and secure handling of sensitive retail information.
- Executed migration from legacy retail management systems to modern Azure-based environments. Maintained full compliance with enterprise data governance and integration standards throughout the transition.
- Built machine learning models on Azure Databricks to forecast product demand and inventory needs. These predictive models improved planning accuracy and reduced stock shortages.
- Developed automated ETL pipelines integrating ADF with Talend for continuous data synchronization. This reduced manual effort and enhanced real-time data availability for business analytics.
- Created validation and transformation logic within Azure Data Factory for data quality management. These enhancements ensured clean and reliable data for downstream reporting systems.
- Collaborated with cross-functional teams to integrate Airflow DAGs with Azure Databricks jobs. This seamless orchestration enabled smooth execution of data science workflows and analytics pipelines.
- Designed metadata-driven frameworks for Talend ETL workflows, improving scalability and maintainability. This modular approach reduced deployment time and enhanced workflow monitoring.
- Implemented access control mechanisms using Azure Active Directory and role-based policies. This safeguarded critical enterprise data while ensuring authorized access for analytical and operational teams.

Environments: Azure Cloud, Talend ETL, Azure Data Factory (ADF), Azure Databricks, Azure Data Lake Storage, Azure SQL Database, Epic, FHIR

Tally Solutions: May 2014 – October 2015 (Location: Bangalore, India)

Role: ETL Developer

Responsibilities:

- Created a scalable ETL pipeline using AWS services, Informatica, and Talend to integrate data from multiple enterprise sources. This streamlined data movement from relational, NoSQL databases, and flat files into centralized storage.

- Configured AWS S3 as a secure storage layer for raw, staging, and transformed data. This setup enabled seamless ingestion workflows and supported consistent data accessibility for downstream analytics.
- Utilized Informatica PowerCenter to extract and consolidate data from MySQL, PostgreSQL, and MongoDB databases. The integration ensured accurate extraction and consistency across diverse business systems.
- Built Talend jobs to process and transform data from CSV and JSON files. These optimized workflows significantly reduced overall data processing time and improved throughput.
- Implemented advanced data transformations using Informatica for cleansing, validation, and quality checks. This improved data accuracy and ensured reliable datasets for reporting and analytics.
- Configured Talend Open Studio to perform complex data transformations and integrations. These configurations simplified ETL workflows and improved maintainability of data pipelines.
- Created AWS Glue jobs for automated cataloging and loading of transformed data into Redshift. This enhanced data organization and optimized the warehouse for high-performance analytical queries.
- Automated ETL scheduling using AWS Lambda and CloudWatch Events. This reduced manual intervention, improved job reliability, and ensured timely data updates.
- Monitored ETL job performance through AWS CloudWatch dashboards and alerts. Proactive monitoring helped identify failures early and minimize downtime in production pipelines.
- Built Tableau connections with AWS Redshift to visualize and analyze large datasets interactively. These dashboards provided business teams with real-time insights into key operational metrics.

Environments: AWS, Informatica, Talend, AWS S3, AWS Glue, AWS Redshift, Tableau, AWS Lambda

Education: JB Institute of Engineering and Technology- Computer Science (July 2010-May 2014)