

AMRUTHA SAI EDARA

Dallas, TX | +1 617-749-5883 | amruthasaiedara11@gmail.com | Tableau Public

PROFESSIONAL SUMMARY

Data Analytics Engineer with 4 years of experience transforming operational, financial, and risk data into reliable datasets, dashboards, and business insights. Experienced in SQL, Python, Tableau, Power BI, Snowflake, dbt, Airflow, and AWS, with a strong focus on ad hoc analysis, KPI reporting, data validation, and reporting automation. Proven ability to work with business and technical stakeholders to improve data quality and support informed decision-making.

TECHNICAL SKILLS

Programming Languages: Python, SQL, PySpark

ETL and Data Integration: SSIS, SnapLogic, ETL/ELT Development, Data Transformation, Incremental Loading, Batch Processing

Data Platform: Snowflake, Databricks, SQL, dbt, NoSQL, AWS S3, Glue, Lambda, Kinesis

Orchestration and Workflow Management: Apache Airflow (MWAA), Job Scheduling, Pipeline Monitoring / Alerting

Data Quality: Soda, dbt Tests, Data Profiling, Reconciliation, Completeness, Uniqueness, Referential Integrity, Schema Drift

Analytics and BI: Tableau, Power BI, Ad Hoc Analysis, KPI Development, Dashboarding, Statistical Analysis, Trend Analysis, Root-Cause Analysis

Orchestration and Devops: Apache Airflow, AWS MWAA, Jenkins, Git, CI/CD, Monitoring and Alerting

Machine Learning and Analytics Integration: Scikit-learn, MLflow, Feature Engineering, Model Validation, Precision, Recall, F1-Score, ROC-AUC, SVM, NLP, LLMs

Performance Optimization: Query Optimization, Pipeline Tuning, Snowflake Performance Tuning

Data Governance: Collibra, Data Lineage, Metadata Management, Source-to-Target Mapping, Data Contracts

CERTIFICATIONS

AWS Cloud Practitioner | SnowPro Core Certification | Databricks Data Engineer Associate | Google Data Analyst Certification | Data Visualizations with Advanced Excel by PWC

PROFESSIONAL EXPERIENCE

Cognizant

Feb 2026 – Present

Data Analyst

Dallas, TX

- Partnered with business, risk, compliance, and engineering teams to define KPI logic, reporting requirements, data-quality rules, and analytics-ready datasets supporting operational decisions.
- Conducted ad hoc SQL and Python analysis across 2M+ daily regulatory events to identify trends, anomalies, SLA risks, and process-improvement opportunities.
- Developed and maintained SSIS packages to extract, transform, and load data from SQL and operational source systems into reporting tables, implementing incremental loads, validation checks, logging, and error handling to improve data reliability and reduce manual processing.
- Developed and automated Tableau and Power BI dashboards, scorecards, and recurring reports, improving reporting availability and reducing data latency by 40%.
- Implemented dbt and Airflow validation checks covering reconciliation, completeness, uniqueness, freshness, schema drift, referential integrity, and business rules, reducing downstream data-quality issues by 25%.
- Automated reporting workflows, monitoring, and alerts through Airflow and CI/CD, reducing job failures by 35% and lowering manual operational effort.
- Developed analytics datasets for machine-learning models and operationalized model outputs through MLflow, SQL, and BI dashboards, enabling risk and compliance teams to monitor prediction trends, identify high-risk events, and perform targeted ad hoc analysis.

CBRE (through Populus Group)

Oct 2024 – Dec, 2025

Data Engineer

Richardson, TX

- Worked across two teams Enterprise Data Platform (EDP) and Client 360 developing enterprise data pipelines, dimensional models, data-quality frameworks, and analytics-ready datasets.
- On the EDP team, developed and maintained SnapLogic pipelines across the Landing, Unified, and Master layers, integrating data from multiple enterprise source systems into Snowflake.
 - Built SnapLogic ingestion pipelines to replace the legacy Lymbic Landing-to-Unified processing layer, automating source-data ingestion and transformation into Snowflake while supporting scalable, standardized data movement across enterprise systems.
 - Worked on and optimized end-to-end data pipelines by tuning SQL queries, implementing incremental loading, and improving transformation logic, reducing pipeline runtime by 30% and minimizing processing backlogs.
 - Developed and maintained dimensional data marts and microservice-ready datasets, incorporating new attributes, source-system changes, and evolving business rules across more than 50 business-critical attributes and KPIs.
 - Optimized Snowflake SQL, incremental loading, and transformation workflows, reducing pipeline processing time by 30% and dashboard load times by up to 35%.
 - Implemented Jenkins-based CI/CD pipelines to promote SnapLogic, SQL, and data-model changes across DEV, QA, and PROD, improving deployment consistency and production reliability.

- On the Client 360 team, developed advanced dbt models in Snowflake and created source-to-target mapping (STTM) documentation to map source attributes, transformation rules, business logic, and KPI definitions into standardized client, deal, revenue, finance, and operational datasets.
- Implemented more than 120 automated dbt and Soda validation checks, including uniqueness, not-null, referential-integrity, reconciliation, freshness, schema-drift, and custom business-rule tests, reducing downstream data discrepancies by 20%.
- Developed optimized Snowflake SQL queries and KPI datasets for Tableau and Power BI dashboards, enabling secure, standardized reporting for finance, operations, and regional leadership.
- Built a new Airflow DAG on AWS MWAA that coordinated Soda instances with project-specific SnapLogic and dbt workflows, using automated retries, alerts, and exception handling to reduce manual operational effort by 40% as a part of Proof of concept.
- Collaborated with business, application, finance, and IT teams to translate reporting requirements into data contracts, business logic, dbt transformations, data-quality rules, and dashboard-ready attributes.
- Leveraged Elise AI and Copilot to optimize Snowflake SQL and dbt transformations, troubleshoot SnapLogic and Airflow workflows, accelerate code development, and improve technical documentation while validating all AI-generated recommendations.

Northeastern University
Data Analyst (Part-time)

Sep 2023 – Apr 2024
Boston, MA

- Built Tableau and Power BI dashboards to support university reporting needs, consolidating data from multiple internal sources into KPI visualizations supporting academic and operational tracking.
- Performed data analysis on university datasets using SQL and Python to identify enrollment trends, validate data integrity, and generate reporting-ready datasets for internal academic use.

Accenture
Application Development Associate

Jun 2021 – Aug 2022
Hyderabad, India

- Built and maintained high-volume ETL pipelines using SQL and IBM DataStage to integrate operational data into centralized analytics platforms; improved pipeline reliability through logging, parameterization, and error handling.
- Optimized large-scale data workflows through SQL query tuning, reducing pipeline backlogs and improving performance of enterprise analytics processing jobs.
- Worked with JSON-based data from APIs and enterprise applications, parsing, validating, and transforming semi-structured data into SQL-based analytics tables to support reporting, reconciliation, and downstream business workflows.
- Implemented data governance and CI/CD frameworks using Collibra, Jenkins, and Git, improving data lineage, metadata management, and automated deployment pipelines for reporting dashboards.

PROJECTS

Heart Disease Prediction | Python, Scikit-Learn, Pandas, Matplotlib

- Built and evaluated multiple classification models including KNN, SVM, Decision Trees, and Random Forest using precision, recall, F1-score, and ROC-AUC metrics.
- Achieved 92.3% recall and 90.4% F1-score through rigorous EDA, feature engineering, hyperparameter tuning, and cross-validation, demonstrating strong application of statistical methods to real-world prediction problems.

Abstract Summarization using Transformer Mode | Python, TensorFlow, Hugging face, NLTK, LLM

- Developed a transformer-based model for efficient text summarization, Employed tokenization, batch processing, and data collation for sequence-to-sequence learning.
- Validated and fine-tuned model efficiency, achieving 89.9% summarization accuracy.

EDUCATION

Northeastern University
MS in Computer Engineering – Information Systems

Boston, MA
Sep 2022 – May 2024

Jawaharlal Nehru Technological University
BTech in Electronics and Communication Engineering

Hyderabad, India
Jun 2017 – Jul 2021