ANVESH B

Data Scientist | Machine Learning | AI

anveshbandari2025@gmail.com | +1 (551) 236-4699 | **LinkedIn:** https://www.linkedin.com/in/anvesh-b-6b4940300

Professional Summary:

Data Scientist with 4+ years of experience designing and deploying **Machine Learning, AI, and Generative AI** solutions across financial services and healthcare domains. Proficient in **Python, SQL, PySpark, TensorFlow, Scikit-learn, AWS SageMaker, NLP, LLMs, and MLOps** to build predictive models, develop scalable ML pipelines, and deliver actionable business insights. Experienced in **A/B testing, model deployment, inference pipelines, CI/CD, model monitoring, and production-grade ML systems** supporting enterprise operations. Hands-on expertise with **NLP, LLMs, RAG pipelines, prompt engineering, LangChain, REST APIs, MLflow, Docker, and Kubernetes**, delivering measurable business impact through intelligent automation and scalable AI solutions.

Technical Skills:

- **Programming & Databases**: Python, SQL, R, PySpark, PostgreSQL, MySQL, MongoDB, NoSQL
- **Machine Learning & AI:** Scikit-learn, TensorFlow, PyTorch, Keras, XGBoost, LightGBM, Random Forest, Gradient Boosting, Generative AI, Explainable AI, Reinforcement Learning, Model Evaluation, Experimentation, A/B Testing
- **Deep Learning, NLP & GenAI:** Neural Networks, CNN, RNN, LSTM, Transformer Models, BERT, GPT, LLMs, Hugging Face Transformers, Vector Databases, Fine-Tuning, spaCy, NLTK, LLM Evaluation, Prompt Engineering, RAG Pipelines, LangChain
- **Data Processing & Engineering:** Pandas, NumPy, SciPy, Data Analysis, Predictive Modeling, Feature Engineering, Data Preprocessing, ETL Pipelines, Apache Spark, Hadoop, Kafka, Hive, Databricks
- **Cloud Platforms:** AWS (SageMaker, S3, EC2, Lambda, EKS), Azure (ML Studio, Databricks), GCP
- **Visualization, MLOps & Tools:** Tableau, Power BI, Matplotlib, Seaborn, Plotly, Looker, MLflow, Airflow, Docker, Kubernetes, Jenkins, GitHub Actions, Terraform, REST APIs, CI/CD Pipelines, MLOps, Model Deployment, Model Monitoring, Flask, FastAPI, Git, Jupyter Notebook, Streamlit
- **Statistical Analysis & Business Intelligence:** Hypothesis Testing, Regression Analysis, Time Series Forecasting, Statistical Modeling, Data Storytelling, KPI Analysis

Work Experience:

Fidelity Investments | Data Scientist with ML | New York, NY | May 2025 – Present

- Developed fraud detection models using **XGBoost** and **Random Forest** on 50M+ financial transactions, achieving 94% accuracy while preventing $15M+ fraud annually and strengthening enterprise risk detection for financial systems
- Engineered real-time risk scoring systems using **Gradient Boosting** and **Neural Networks**, reducing false positives by 35% and improving transaction approval accuracy for financial transaction processing platforms
- Implemented customer segmentation using **K-means** and **RFM Analysis** on 100M+ customer profiles, increasing campaign engagement by 18% through targeted marketing, behavioral analytics, and retention optimization initiatives
- Built predictive churn models using **Logistic Regression** and **Deep Learning**, achieving 87% accuracy while reducing customer attrition by 22% through customer retention and engagement strategies
- Designed feature engineering, model deployment, and **MLflow-based Monitoring** pipelines using **PySpark** and **AWS SageMaker**, processing 5TB+ enterprise data and reducing deployment turnaround time by 40% for production ML systems
- Performed **A/B testing** and model validation on fraud detection models to support enterprise risk decisions
- Collaborated with fraud analysts and engineering teams to deploy ML models using **Docker, Kubernetes, and CI/CD Workflows**, enabling scalable payment security operations

Environment: Python, NumPy, Pandas, TensorFlow, XGBoost, Scikit-learn, AWS (SageMaker, S3, EKS), PySpark, SQL, Tableau, Docker, MLflow, LangChain, REST APIs, MLOps

The Cigna Group | Data Scientist with ML | Mumbai, India | Oct 2021 – Jul 2024

- Formulated healthcare cost forecasting and predictive modeling frameworks using **ARIMA, Prophet, and Machine Learning** techniques, achieving 85% forecasting accuracy and supporting $200M+ healthcare budget planning initiatives
- Constructed NLP-based claims processing systems using **BERT** and **Transformer Models**, extracting insights from 500K+ clinical documents monthly and reducing manual review effort by 65% through intelligent text processing
- Optimized patient risk stratification models using **Gradient Boosting** and **Collaborative Filtering**, reducing readmissions by 20% and increasing healthcare program enrollment by 35% using predictive analytics
- Delivered scalable data pipelines, model deployment, and **MLflow-based monitoring** processes using **Python, SQL, Flask, and Docker**, improving data accuracy from 82% to 96% and reducing manual processing effort
- Applied **Hugging Face Transformers**, NLP techniques, and model deployment workflows to improve clinical document analysis and accelerate healthcare insight generation

Environment: Python, Scikit-learn, TensorFlow, BERT, Transformers, NLP, Azure ML, SQL, Flask, Docker, Pandas, NumPy, Git, MLflow, Kubernetes, REST APIs, MLOps

Education:

- Master of Science in Data Science, New Jersey Institute of Technology (NJIT), USA
- Bachelor of Science in Computer Science, A.V. College of Arts, Science & Commerce, India