

Anyutha Reddy Jaidi Data Engineer

+1 (313) 558-2087 | anyuthajaidi@gmail.com | [LinkedIn](#)

PROFESSIONAL SUMMARY

Data Engineer with 3+ years of experience designing scalable, cloud-native data platforms and high-performance ETL/ELT pipelines across fintech and enterprise environments. Specialized in Python, SQL, Apache Spark, Databricks, and AWS, with a strong track record of building real-time ingestion systems, lakehouse architectures, and high-quality data models that power fraud detection, financial analytics, and executive reporting. Proven ability to improve pipeline reliability, optimize compute cost, automate orchestration with Airflow, and strengthen governance using Great Expectations, Collibra, and Apache Atlas. Adept at collaborating with cross-functional teams to deliver production-ready, highly available data solutions that boost analytics speed, accuracy, and business impact.

TECHNICAL SKILLS

Programming & Scripting: Python, SQL, PySpark, Bash, Data Pipeline Unit Testing (PyTest)

Data Engineering & Big Data: Apache Spark, Spark Streaming, Hadoop Ecosystem, Delta Lake, Data Lakehouse Architecture, Medallion Architecture, Unity Catalog, Parquet, Avro, JSON, ORC, EHR/EMR, Claims Data, FHIR/HL7 formats

ETL / ELT & Orchestration: Apache Airflow, Databricks Workflows, dbt, Data Contracts, Workflow Automation, SLA Management, CDC (Debezium/Kafka Connect), PHI Masking Pipelines & Healthcare ETL

Cloud Platforms: AWS (S3, Lambda, Redshift, Glue, IAM, CloudWatch), Azure (ADLS, Synapse), Serverless Computing

Data Warehousing: Snowflake, Dimensional Modeling, Star & Snowflake Schema Design, Query Optimization

Databases: PostgreSQL, MySQL, SQL Server, NoSQL (MongoDB), Data Lake Storage (S3, ADLS)

Data Quality & Governance: Great Expectations, Anomaly Detection, Data Validation Pipelines, Collibra, Apache Atlas, Lineage Tracking, Compliance (PCI, SOC2), HIPAA Compliance, PHI Masking

DevOps & Automation: Git, GitHub Actions, Jenkins, Terraform, CI/CD Pipelines, Docker, Kubernetes (basic), Logging & Observability (ELK, Prometheus, Grafana), Linux

PROFESSIONAL EXPERIENCE

Data Engineer | PayPal | Remote, USA

Jan 2025 – Present

- Engineer high-scale ETL/ELT pipelines using **Python, SQL**, and **Apache Spark** to ingest data from APIs, flat files, and SaaS systems, processing **2 TB/day** and increasing pipeline throughput, improving analytics for fraud, risk, and finance teams.
- Architect an **AWS** Lakehouse platform using **Databricks, Delta Lake**, and **S3**, implementing **Medallion architecture** and **Unity Catalog** to support real-time transaction monitoring and PCI-compliant data governance.
- Implemented automated data quality and observability using **Great Expectations**, custom anomaly rules, and **Databricks Workflows**, reducing downstream data errors by 30%+ and cutting manual validation hours by 25/week.
- Automate and optimize pipeline orchestration with **Apache Airflow**, redesigning 70+ DAGs, improving SLA adherence from 82% to 98%, and reducing manual intervention by 80% through lineage, alerting, and dependency management.
- Design high-performance **Snowflake** data models tailored for fraud analytics and KPI dashboards, improving query latency by 25–40% and accelerating delivery of executive insights across finance and operations.
- Enhanced pipeline reliability and deployments with **CI/CD, GitHub Actions, data contracts**, and **CloudWatch** monitoring, reducing failure recovery time by 40% and increasing deployment frequency by 2x.

Data Engineer | Accenture | India

Jun 2021 – Dec 2023

- Built secure data pipelines for a healthcare client using **Spark, Python**, and **Airflow** to ingest **EHR/EMR records** and claims data, applying **HIPAA-compliant** encryption, **PHI masking**, and quality checks that improved reporting accuracy for clinical and operational teams.
- Designed scalable warehouse schemas (Star/Snowflake) in **Snowflake** and **Redshift**, improving analytical performance by 50%, reducing compute costs, and supporting BI teams with faster self-service reporting.
- Modernized transformation pipelines by introducing **dbt**, converting 40+ ad-hoc SQL scripts into modular, version-controlled models with automated tests/documentation, cutting data inconsistencies by 35% across business teams.
- Optimized **Spark** performance through partitioning, caching, Z-ordering, and predicate pushdown, improving job runtime by 45% and lowering cloud compute spend by 20% across batch ETL workloads.
- Established enterprise data governance standards using **Collibra** and **Apache Atlas**, enabling end-to-end lineage, impact analysis, and compliance with internal audit requirements, increasing trust and transparency across analytics teams.
- Collaborated with engineering, QA, and product teams to troubleshoot critical data incidents, refine SLAs, and deliver production-hardened pipelines using **CI/CD, Git**, and **SQL**, boosting report availability and reducing data outages by 30%.

EDUCATION

Master of Science in Information Technology

Lindsey Wilson University

Jan 2024 – Aug 2025

Columbia, KY