

Aravind Kumar Reddy Yeddula (551)-260-8700 aravindkumarreddyyeddula@gmail.com LinkedIn	
SUMMARY	
Data Engineer with 4+ years of experience in designing and deploying scalable data pipelines, cloud-based solutions, and performance-optimized ETL workflows. Proficient in tools like Azure Data Factory, Databricks, PySpark, SQL, and Power BI, with hands-on expertise in Snowflake and Azure Synapse. Strong track record of improving data reliability, accelerating reporting processes, and enabling data-driven decision-making across business functions.	
SKILLS	
Programming Languages: Python, SQL, PySpark Data Manipulation: Pandas, NumPy Data Visualization: Power BI, Matplotlib, Seaborn Database Management: MySQL, PostgreSQL, Microsoft SQL Server Cloud Platforms: AWS S3, Azure Data Factory, Azure Synapse Analytics, Azure Databricks Data Warehousing: Snowflake, Redshift, BigQuery Data Engineering Tools: Apache Spark, Apache Kafka, Apache Airflow, dbt Workflow Orchestration: Airflow, Azure Data Factory DevOps & CI/CD: Docker, Kubernetes, Terraform, GitHub Actions Version Control & Collaboration: Git, Jira, Confluence Machine Learning: Scikit-learn, XGBoost, RapidMiner, TensorFlow, Keras, Classification, Regression, Model Tuning Soft Skills: Problem-Solving, Critical Thinking, Communication, Team Collaboration, Adaptability, Time Management	
PROFESSIONAL EXPERIENCE	
Data Engineer, CitiusTech	Dec 2024 - Present USA
- Designed and implemented scalable ETL pipelines using Azure Data Factory and Databricks, optimizing the integration of complex healthcare datasets and reducing reporting latency by 54%, directly improving the speed and reliability of decision-making for clinical and operational teams. - Developed PySpark workflows for large-scale data processing, improving performance by 37% and ensuring data consistency for real-time analytics, which resulted in more accurate insights that enabled faster response times in critical healthcare operations. - Engineered and maintained data models in Azure Synapse, enabling the creation of reliable, real-time reports that supported data-driven decisions across clinical and operational teams, enhancing overall healthcare service delivery. - Automated data validation and quality checks with dbt and Python, reducing manual efforts by 46%, enhancing data traceability, and ensuring compliance with stringent healthcare data regulations like HIPAA, thereby minimizing risk and ensuring trust in the data. - Delivered actionable insights through Power BI dashboards, simplifying complex healthcare data and empowering business and clinical teams to make informed, impactful decisions on care delivery, improving patient outcomes and operational efficiency.	
Data Engineer, Technovert Solutions Private Limited	Jan 2022 – Jun 2023 India
- Built and maintained scalable data pipelines using Azure Data Factory and Databricks, improving data processing efficiency and reducing pipeline failures, resulting in more reliable reporting and faster data delivery. - Developed dynamic PySpark scripts for transforming raw data into analytics-ready formats, increasing data accuracy by over 91% and cutting manual data correction time by 37%. - Consolidated data from SQL Server, APIs, and flat files into Snowflake and Azure Synapse, enabling unified access to enterprise data and decreasing report generation time. - Automated recurring ETL workflows with Apache Airflow, reducing manual workload by 64% and improving SLA adherence for data refresh cycles across business-critical dashboards. - Collaborated with analysts and business users to create actionable Power BI dashboards that improved visibility into key metrics, accelerating decision-making processes in sales and operations. - Optimized slow-running SQL queries and Spark jobs, leading to a 53% reduction in data load times and significantly improving performance for both scheduled and ad-hoc reports.	
Data Engineer, Wipro	Jun 2019 - Dec 2021 India
- Designed and maintained scalable ETL pipelines using Azure Data Factory and PySpark, automating the ingestion and transformation of data daily, which reduced manual effort and ensured timely data availability for analytics. - Developed modular data workflows using Azure Synapse Analytics, Databricks, and Snowflake, streamlining data integration across platforms and reducing the report delivery time by up to 52%, directly improving decision-making speed. - Created interactive Power BI dashboards to visualize key performance metrics, increasing business visibility and enabling leadership to make faster, data-backed strategic decisions. - Optimized SQL queries and procedures in PostgreSQL and SQL Server, reducing query runtime by over 61%, which significantly cut down on cloud compute costs and improved database performance. - Implemented Delta Lake features like schema enforcement and time travel, enhancing data consistency and reducing audit preparation time, while supporting regulatory compliance. - Participated in Agile development cycles, contributing to sprint planning and retrospectives, which helped boost team delivery velocity and maintain high sprint success rates. - Partnered with architects and analysts to translate business requirements into scalable data solutions, ensuring alignment with project goals and improving the reliability of downstream analytics. - Leveraged Git, Bitbucket, and Jira for version control and task tracking, improving development transparency and cutting resolution time for issues by 22%. - Established automated data validation frameworks and standardized documentation processes across ETL pipelines, resulting in a 38% drop in data quality issues and increased stakeholder confidence in analytics outcomes.	
EDUCATION	
Master of Science (MS) in Data Science GPA: 4.0/4.0	Sept 2023 - May 2025
Pace University, Seidenberg School of CSIS, New York	
Bachelors of Technology in Electronics and Commuincation Engineering GPA: 7.18/10	Jul 2015 - May 2019
SRM Institute of Science and Technology, Chennai, India	
PROJECTS	
Prediction of US Chronic Diseases Cases	Dec 2024
- Worked on a team of three to collect more than one million records from Center for Disease Control and Prevention (CDC) website which contains information on various chronic diseases, performing EDA and preparing a sample Power BI dashboard for analysis. - Incorporated ML algorithms such as Gradientboosting and Random Forest to predict the number of cases for particular diseases based on factors like age, location, and habits. - Hyper tuned the models to achieve better scores, resulting an an accuracy of 81% with Random Forest.	
Flight Delay Prediction Using Rapidminer	May 2024
- Used Hadoop Hive for querying and cleaning more than 30 million flight records, loading the transformed data into AWS S3 bucket. - Implemented machine learning classification approaches to predict flight delays with 98% accuracy using RapidMiner, developing a pipeline using Alter AI studio for algorithms. - Used Seaborn and Matplotlib libraries to build visuals for understanding the data.	
CERTIFICATIONS	
Udemy: Azure Data Factory	
Coursera : Machine Learning: Regression, Introduction to Python	
Great Lakes Institute of Management : Post Graduate Diploma in Data Science & Engineering	