

SAI ASHISH CHAVALI

Frisco Texas | +1 2147712417 | saiashish1512@gmail.com | [LinkedIn](#)

SUMMARY

Data Scientist / Data Analyst with **4+ years** of experience analyzing enterprise datasets and developing machine learning and AI-powered analytics solutions across marketing, product telemetry, and customer interaction platforms. Skilled in **exploratory data analysis**, predictive modeling, **statistical analysis**, and generative AI applications using **Python, SQL, Scikit-learn**, and **LangChain**. Builds data pipelines, **dashboards**, and AI-driven insight systems enabling product teams and business stakeholders to extract value from large operational datasets.

PROFESSIONAL EXPERIENCE

Epsilon | USA

Data Scientist | Jul 2025 - Present

- Designed retrieval-augmented generation workflows connecting analyst query tools with enterprise marketing knowledge bases, indexing documents across **4 internal repositories** through **Python, LangChain**, and OpenAI APIs with storage on **AWS S3**.
- Adapted transformer language models to campaign analytics datasets spanning **5 marketing platforms**, performing fine-tuning experiments in **PyTorch** and **Hugging Face Transformers** to generate structured insights from customer interaction logs.
- Established semantic search capabilities for enterprise knowledge assets by embedding campaign documentation and analytics reports across **6 internal content collections** using **Sentence Transformers** and **FAISS**.
- Deployed language model inference services through containerized **FastAPI** applications orchestrated with **Docker** on AWS compute infrastructure, exposing prediction endpoints consumed by **3 internal analytics services**.
- Prepared structured training datasets from CRM and campaign engagement records through **Python** and **SQL** pipelines consolidating marketing signals across **5 operational systems** for large language model experimentation.
- Evaluated generative model responses across **4 prompt configurations** through benchmarking experiments conducted in **Python** and **Jupyter Notebooks**, guiding prompt design decisions for internal AI assistant workflows.

Dell Technologies | India

Data Scientist | Mar 2022 – Jul 2023

- Architected machine learning workflows in **Python, Scikit-learn, and Pandas** across enterprise product telemetry datasets containing **millions of usage events**, enabling analytics teams to evaluate behavior across **5 enterprise product modules**.
- Trained classification and regression models with **Scikit-learn and XGBoost** on structured CRM and product telemetry datasets spanning **3 enterprise systems**, producing churn risk signals reviewed during account strategy planning.
- Engineered predictive features from customer interaction logs and CRM records through **Python, Pandas, and SQL** pipelines, producing training datasets with **40+ behavioral variables applied across 5 predictive modeling experiments**.
- Implemented data ingestion and transformation pipelines integrating **6 enterprise data sources**, producing standardized datasets used by **machine learning experiments** and product analytics dashboards.
- Conducted exploratory analysis through **Python, Pandas, NumPy, and Jupyter Notebooks** across large telemetry datasets, identifying usage patterns across **5 platform modules** discussed during product analytics and strategy reviews.
- Deployed trained prediction models through **Flask REST APIs and Docker containers**, enabling internal applications to retrieve inference outputs from **3 containerized services** supporting analytics and forecasting workflows.
- Benchmarked model performance through cross-validation experiments and hyperparameter tuning in **Scikit-learn**, comparing **4 candidate algorithms** to guide model selection decisions during internal machine learning review cycles.

Adani | India

Data Analyst | Jan 2021 – Feb 2022

- Analyzed logistics and operational datasets using **SQL**, querying **6 warehouse tables** to detect reporting inconsistencies and align records used in internal operations dashboards.
- Developed ETL data preparation pipelines in **Python and Pandas** to extract and standardize operational records from **4 business units**, producing structured datasets for analytics and reporting.
- Designed interactive **Tableau dashboards** tracking shipment volume, facility throughput, and service metrics across **8 logistics hubs**, enabling operations teams to monitor regional performance.
- Examined operational supply chain datasets spanning **6 warehouse tables** and **120K+ inventory records** using **Python** and **Excel**, tracing inventory flow patterns across distribution facilities supporting procurement planning discussions.
- Produced recurring **SQL reports** aggregating transaction activity from **5 operational systems**, supplying weekly operational metrics reviewed by analytics and operations leadership.
- Audited enterprise reporting datasets through **data validation queries and reconciliation checks**, maintaining consistency between ingestion tables and curated datasets supporting internal analytics.

TECHNICAL SKILLS

- **Programming:** Python, R, SQL, PySpark, Scala, Bash
- **Machine Learning & AI:** Supervised Learning, Unsupervised Learning, Regression, Classification, Clustering, Time Series Forecasting, Feature Engineering, Model Evaluation, Hyperparameter Tuning, Predictive Modeling, Natural Language Processing (NLP), Deep Learning, Neural Networks, Transformers, Large Language Models (LLMs), Generative AI, Prompt Engineering, RAG
- **ML & AI Libraries:** Scikit-learn, TensorFlow, PyTorch, Keras, XGBoost, LightGBM, CatBoost, Hugging Face Transformers, spaCy, NLTK, Sentence Transformers
- **Data Analysis:** Pandas, NumPy, SciPy, Statsmodels
- **Data Visualization:** Tableau, Power BI, Matplotlib, Seaborn, Plotly, Looker, Microsoft Excel
- **Big Data:** Apache Spark, Hadoop, Databricks, Hive
- **Data Engineering:** ETL Pipelines, Data Cleaning, Data Transformation, Data Modeling, Data Warehousing
- **Databases:** PostgreSQL, MySQL, SQL Server, Snowflake, BigQuery, MongoDB
- **Cloud Platforms:** AWS (S3, EC2, Lambda, SageMaker, Redshift), Azure (Azure Machine Learning, Data Factory, Blob Storage), Google Cloud Platform (BigQuery, Vertex AI)
- **MLOps & Deployment:** Docker, Kubernetes, MLflow, REST APIs, FastAPI, Flask, CI/CD
- **Workflow Orchestration:** Apache Airflow, Prefect
- **Version Control & Collaboration:** Git, GitHub, GitLab, Bitbucket, Jira, Agile
- **Data Governance:** Data Validation, Data Quality Monitoring, Data Integrity Checks

EDUCATION

Master of Science in Business Analytics and Artificial Intelligence, Data Science Concentration
University of Texas at Dallas – Jindal School of Management | Dallas, TX | May 2025

PROJECTS

TrueSense AI Fraud Detection Model

- Constructed a synthetic fraud detection dataset containing **530K transactions and 31 engineered features**, modeling transaction patterns and device behavior through **Python, Pandas, and SQL** for financial risk experiments.
- Implemented fraud detection models with **XGBoost** and **Random Forest**, tuning hyperparameters through randomized search and validating feature impact through **SHAP analysis** across large transaction datasets.

Multi-Label Tweet Emotion Classification

- Built transformer-based classification models with **BERT and RoBERTa** using **Hugging Face** and **PyTorch**, classifying **10K tweets across 11 emotional categories** for social media sentiment analysis.
- Prepared NLP training datasets through **spaCy preprocessing, tokenization pipelines**, and dataset balancing, enabling consistent model training experiments tracked through **Weights & Biases**.

Product Usage Analytics Dashboard

- Analyzed product interaction datasets containing **120K user activity records** through SQL and Python, identifying engagement patterns and behavioral cohorts across **6 application modules**.
- Created interactive **Tableau dashboards** presenting feature adoption metrics, retention patterns, and usage indicators across **6 application components** for product analytics review.

CERTIFICATIONS

- AWS Certified Cloud Practitioner
- Python for Data Science
- Google Analytics GA4 Certification
- Machine Learning
- NVIDIA Large Language Models (LLM)
- Data Science Prerequisites Certificate – 360DigiTMG
- IBM Generative AI Applications for Beginners
- Java Spring Boot Training