

Ashritha Kotla

ashritha0704@gmail.com | (469) 956-7184 | USA | [LinkedIn](#)

Summary

A highly skilled Data Engineer with 4 expertise in designing, optimizing, and managing data pipelines, cloud architectures, and business intelligence solutions. Proficient in AWS, Azure, Apache Spark, SQL, and Python. Experienced in implementing scalable ETL processes, optimizing query performance, and enhancing data processing efficiency. Strong ability to collaborate with cross-functional teams to deliver high-quality, actionable insights. Committed to maintaining data integrity, governance, and compliance while driving innovation in data-driven decision-making across various industries.

Technical Skills

- Programming Languages:** Python, SQL
- Big Data & Cloud Platforms:** AWS (Glue, Redshift, Lambda), Azure (Data Factory, Data Lake Storage Gen2, Synapse Analytics), GCP (Dataflow, BigQuery, Cloud Storage)
- Data Engineering Tools:** Apache Spark, Kafka, PySpark, Apache Beam
- ETL & Data Integration:** Apache Spark, SQL Server Integration Services (SSIS)
- Data Wrangling & Transformation:** Python (pandas, pyarrow), SQL, Apache Spark
- Business Intelligence & Reporting:** Tableau, Power BI, Amazon QuickSight
- Metadata & Data Governance:** Azure Purview, AWS Glue Data Catalog
- DevOps & Automation:** Apache Kafka, GCP Cloud Functions, GCP Shell
- Machine Learning & AI:** Vertex AI (Gemini Pro), Google Document AI
- Query Optimization:** SQL (window functions, indexing, partitioning)

Professional Experience

- Data Engineer, KPMG**

01/2024 – Present | Remote, USA

 - Worked on Data Warehouse Optimization for Business Intelligence, conducting detailed requirement-gathering sessions through Agile sprints. Collaborated with data architects, BI analysts, and cloud engineers to optimize data lakes and manage multi-source integrations (Parquet, CSV).
 - Migrated legacy ETL workflows to scalable data pipelines using AWS Glue, Apache Spark, and Amazon Redshift Spectrum. Improved data processing efficiency by 55%, utilizing dynamic partitioning and auto-scaling for handling large-scale datasets in near-real time.
 - Built data wrangling scripts with Python, using pandas for transformation and pyarrow for optimized Parquet file handling. Reduced processing time by 40% via multi-threaded transformations and efficient memory management in large-scale data environments.
 - Optimized SQL queries in AWS Redshift with window functions (ROW_NUMBER (), RANK ()) for deduplication and aggregations. Enhanced performance by 60% through proper indexing, distribution, and partitioning strategies on multi-terabyte datasets.
 - Collaborated with BI teams to redesign AWS QuickSight dashboards and optimize OLAP cubes. Improved report performance by 50%, using pre-aggregated metrics, multi-dimensional modeling, and efficient query execution to aid real-time analytics for business users.
 - Implemented AWS Glue Data Catalog for metadata management and lineage tracking, automating schema evolution and versioning. Integrated AWS Lambda for automation, ensuring 100% auditability and compliance with governance policies across all data flows.
- Associate Data Engineer, Invesco India Pvt. Ltd.**

01/2020 – 07/2022 | Hyderabad, India

 - Initiated Performance Data Aggregation and Calculation Pipeline through Agile-led stakeholder sessions, aligning goals across data, compliance, and reporting. Improved project clarity by 40% via clearly documented KPIs, timelines, and sprint-level expectations.
 - Built scalable pipelines with Apache Spark, Kafka, and Azure Data Factory, improving ETL efficiency by 60%. Enabled processing of 4x more data from sources like Morningstar and State Street with zero latency backlog.
 - Used Python and pandas to clean, transform, and normalize data, boosting calculation accuracy by 35% across currencies and timeframes. Reduced exception handling time by 50% through dynamic schema validation scripts.
 - Created advanced SQL queries in Azure Synapse Analytics using WITH clauses and window functions, improving query performance by 45% and ensuring consistent composite aggregation over datasets exceeding 5 billion records.
 - Collaborated with data science, BI, and compliance teams, implementing PySpark tests and quality gates that reduced data issues in reporting pipelines by 70%, enhancing overall reliability of monthly return calculations.
 - Delivered fund performance dashboards using Tableau and Power BI, increasing stakeholder visibility into trends by 50%. Drove 30% faster decision-making for investment teams through enriched, real-time composite analytics.
 - Implemented data lineage using Azure Purview, increasing traceability by 90%. Automated metadata management in Azure Data Lake Storage Gen2, aligning 100% with internal governance and external audit requirements.

Education

The University of Texas at Dallas	TX, USA
Master of Science in Business Analytics	08/2022 – 08/2024
Chaitanya Bharathi Institute of Technology	Hyderabad, India
Bachelor of Engineering in Information Technology	07/2017 – 08/2021

Certificates

- Microsoft Certified: Azure Data Fundamentals**
- Azure Data Engineer Associate**
- Google Cloud Professional Data Engineer**