Athif Mohiuddin Mohammed

Data Engineer

TN | 301-259-1284| athifmohiuddin2025@gmail.com|Linkedin

SUMMARY

Data Engineer with 4+ years of experience architecting scalable, end-to-end data ecosystems that accelerate analytics and AI adoption across enterprise environments. Proven expertise in designing high-throughput ETL and ELT pipelines with Python, PySpark, Airflow, and Snowflake to optimize data latency, storage efficiency, and governance. Skilled in developing real-time data ingestion frameworks and machine learning pipelines that enhance predictive analytics and decision automation by up to 45%. Adept at leveraging AWS and Snowflake to build secure, distributed data platforms integrating structured, semi-structured, and unstructured sources. Collaborative partner to product, analytics, and engineering teams, translating complex data problems into actionable insights that drive measurable business growth.

SKILLS

Programming: Python, R, SQL, Scala, Java, Django
Data Engineering: ETL, ELT, Data Modeling, Data Warehousing, Data Governance, Data Quality, Streaming Pipelines
Frameworks and Tools: Apache Airflow, Apache Kafka, Talend, SSIS, dbt, PySpark, Hadoop, Hive, Spark SQL, MapReduce, HDFS,
Cloud Platforms: AWS (S3, Redshift, Lambda, DynamoDB, EC2), Azure, GCP, Snowflake
Machine Learning and AI: TensorFlow, Scikit-learn, Natural Language Processing, Named Entity Recognition, Generative AI, Retrieval-Augmented Generation, LangChain, Llama Index
Visualization: Power BI, Tableau, Matplotlib, Seaborn
Databases: MySQL, PostgreSQL, MongoDB, Cassandra, HBase, Neo4j, Elasticsearch
DevOps and CI/CD: Docker, Kubernetes, Jenkins, Git, GitHub, Grafana, Splunk
Methodologies: Agile, Waterfall, SDLC

EXPERIENCE

McKesson, TN | Dec 2023 – Current | Data Engineer

- Architected and optimized multi-source ETL pipelines using SSIS, Apache Airflow, and Python to automate ingestion from MySQL and PostgreSQL into Snowflake, reducing latency by 35% and ensuring 99.9% data accuracy.
- Engineered scalable data marts and dimensional models supporting real-time analytics across finance and operations; implemented schema evolution strategies that improved report refresh speed by 40%.
- Developed fault-tolerant distributed data workflows on Hadoop and HDFS, improving data access scalability and cutting query response times by 25%.
- Created BI solutions in Power BI and Tableau to deliver advanced forecasting and KPI dashboards for executive stakeholders, driving data-driven decision-making.

NTT Data, India | May 2021 – Jul 2022 | Data Engineer

- Designed and automated high-performance ETL/ELT pipelines using Talend, Python, PySpark, and SQL, integrating structured and unstructured data from SQL Server, Oracle, Hive, and AWS S3 into Snowflake for advanced analytics.
- Developed efficient data workflows with Snowflake, leveraging partitioning, clustering, materialized views, and query optimization, reducing data retrieval times by up to 40% and processing times by 15% using optimized algorithms.
- Created real-time fraud detection models using Python, PySpark, and Snowflake UDFs, with optimized algorithms to identify transactional anomalies, integrated with Power BI and Tableau dashboards for visualizing KPIs like fraud metric.
- Migrated on-premise legacy systems to Snowflake on AWS with minimal downtime, utilizing AWS Lambda, Step Functions, and Elastic Load Balancing with Auto Scaling for scalable, high-availability data workflows.
- Implemented Snowflake role-based access controls (RBAC) and row-level security to ensure compliance with banking regulations (GDPR, CCPA, PCI-DSS), securing sensitive financial data.

Orion Technolab, India | Dec 2019 – April 2021 | Data Engineer

- Built and optimized distributed data processing workflows using Apache Spark and Spark SQL, reducing data processing time for large-scale analytics by 30%.
- Designed and deployed data lakes on AWS, utilizing S3, Redshift, and Lambda for scalable storage and serverless data processing.
- Developed machine learning models for text classification and Named Entity Recognition (NER) using TensorFlow and Scikit-learn, enhancing NLP-driven insights.
- Implemented Retrieval-Augmented Generation (RAG) pipelines with LangChain, integrating generative AI models for advanced query resolution.
- Orchestrated complex data pipelines using Apache Kafka and dbt, ensuring real-time data integration and transformation for analytics platforms.
- Utilized Docker and Kubernetes for containerized deployment of data applications, improving scalability and resource efficiency in cloud environments.
- Created advanced visualizations using Tableau and Seaborn, presenting actionable insights to business teams for strategic decision-making.

EDUCATION

Masters in Information Systems **May 2024**
The University of Memphis, Tennessee