

BESH VIGNESH REDDY MATTA

Dallas, TX | (940) 268 5588 | beshtreddy.m@gmail.com | [LinkedIn](#)

SUMMARY

Data Analytics Professional with over 4 years of experience building reliable data pipelines, ETL/ELT processes, and analytics solutions using Python, SQL, PySpark, Databricks, Azure Data Factory, and Git. Experienced in medallion architectures, data quality frameworks, REST API integration, and automated financial reporting, with a record of improving data reliability and operational decision-making across revenue-focused environments.

SKILLS

- **Data Engineering:** Python, SQL, Databricks, ETL/ELT, Snowflake, Data Integration, Data Pipelines, Data Governance
- **Business Intelligence & Analytics:** Power BI, Cognos, Excel (Power Query, Pivot Tables)
- **Cloud Platforms:** Microsoft Azure (Azure Data Factory), Amazon Web Services (EC2, S3, Redshift, Glue)
- **AI & Automation:** LLMs(Llama 3.2), RAG, LangChain, Microsoft Power Automate, n8n
- **Tools & DevOps:** Git, Docker, REST APIs, Terraform, Azure DevOps
- **Certifications:** Databricks Certified Data Engineer Associate, Azure Databricks Platform Architect

PROFESSIONAL EXPERIENCE

UT Dallas

Jan 2026 - May 2026

AI Data Intern

Richardson, TX

- Automated 7 recurring operations reports using SQL, scheduled queries, and parameterized templates, reclaiming 20 hours per month across the team, eliminating errors from manual refreshes, and improving the reliability of recurring data workflows.
- Implemented feature and prediction drift monitoring with automated retrain triggers, that caught two upstream data regressions before they degraded production capture rate.
- Collaborated on an AI agent solution to analyze and categorize student inquiries, enabling faster information processing and supporting university outreach initiatives.

Hilton Grand Vacations Inc.

Jun 2025 - Sep 2025

Reporting & Analytics Intern

Orlando, FL

- Developed office-level financial dashboards in Power BI for 20+ operational units, automating budget validation and reconciliation, reducing budget variance from 9% to under 3.5%, and improving yearly financial reviews used by regional operators.
- Deployed Microsoft Power Automate workflows for budget approval processes, eliminating manual bottlenecks and cutting Cognos report delivery from 3+ weeks to 2–3 days, enabling more timely financial decisions through automated data integration.
- Integrated a conversational AI commentary layer using Databricks Mosaic AI in Power BI to automatically generate executive summaries from resort KPIs, saving 4–5 hours per reporting cycle.

Reliance Jio Platforms

May 2020 - Jul 2024

Analytics Engineer

Remote

- Processed 15M+ daily transactions across user, subscription, and revenue data using Python and PySpark on Databricks, delivering low-latency analytics through scalable ETL data pipelines for an internal SaaS platform serving ecommerce and cosmetics business units.
- Analyzed campaign performance and subscription engagement through cohort and conversion analytics, identifying key retention drivers and reporting engagement strategies that increased 90-day retention by 11 percentage points.
- Maintained a medallion architecture on Delta Lake with a config-driven data quality and governance framework featuring 40+ checks, reconciliation identities, and severity-based actions, reducing bad-data escalations from finance from monthly to zero across two consecutive quarters.
- Collaborated with engineering and dependency teams to define data contracts, maintain Git-based version control for ETL changes, ensure upstream data reliability, and coordinate schema updates that improved pipeline stability and reduced downstream data errors.

PROJECTS

AuditFlowX - Hierarchical Data Validation Framework

Jan 2026 - Mar 2026

- Built an AI-driven data validation framework using RAG and Llama 3.2 to detect hierarchy mismatches, duplicate records, missing members, row-level discrepancies, and rollout breaks, translating validation failures into plain-language diagnostic reports and eliminating manual formula auditing.

Distributed Airline Delay Analytics

Feb 2025 - May 2025

- Designed and implemented an automated ETL pipeline using ADF and Databricks in AWS environment to ingest 50+ million daily flight operations records, building a Lakeflow architecture that enabled reliable downstream reporting in Redshift.
- Optimized Spark jobs by identifying skewed joins and applying broadcast joins and partition pruning, reducing pipeline runtime from 2 hours to under 50 minutes and cutting DBU costs by 45%, without increasing cluster size.

EDUCATION

The University of Texas at Dallas

May 2026

Master of Science, Business Analytics and AI

- **Coursework:** Big Data Processing, Applied ML, Gen AI, Predictive Analytics