

JAIRAJ VENKATA KAPEESH JAGARAPU

Saint Louis, MO | +1 314-249-8607 | jairajvkapeesh@gmail.com | [linkedin.com/in/jairajjagarapu](https://www.linkedin.com/in/jairajjagarapu)

PROFESSIONAL SUMMARY

Results-driven Data Engineer with 4+ years of experience building scalable batch and near-real-time ETL/ELT pipelines in financial services and banking environments, processing 500+ GB of high-volume transactional data daily. Strong expertise in Python, SQL, PySpark, Apache Spark, Apache Kafka, and modern cloud data platforms, including Snowflake, Redshift, Azure Databricks, AWS, and Azure to deliver enterprise-grade data solutions. Proven track record in dimensional data modeling, incremental processing, performance optimization, and implementing robust data governance frameworks that ensure reliable, SLA-driven analytics. Passionate about transforming complex data ecosystems into high-impact, business-ready insights and eager to contribute scalable, production-ready data infrastructure that drives strategic decision-making.

TECHNICAL SKILLS

Programming & Query Languages: Python, SQL, PySpark

Big Data & Streaming Technologies: Apache Spark, Apache Kafka, Apache Flink, Spark Structured Streaming

Cloud & Data Platforms: Snowflake, Amazon Redshift, Azure Databricks, Azure Data Factory, Azure Data Lake Storage (ADLS), AWS S3, AWS Glue, Delta Lake, Azure Synapse Analytics

Data Engineering & Analytics: ETL/ELT Pipeline Development, Batch & Incremental Processing, Change Data Capture (CDC), Dimensional Modeling, Star Schema, Snowflake Schema, SCD Type 1/2, Data Quality Validation, Medallion Architecture, Data Lakehouse Architecture

Data Governance: Data Quality Validation, Data Lineage, Data Catalog, RBAC

Data Transformation & BI Tools: dbt (Data Build Tool), Fivetran, Tableau, Looker, Power BI

Orchestration & DevOps: Apache Airflow, Terraform (IaC), Git, Docker, Kubernetes, GitHub Actions, CI/CD Pipelines, Azure DevOps

Data Formats & Protocols: Parquet, ORC, Avro, JSON, REST API Integration

PROFESSIONAL EXPERIENCE

Data Engineer

Apr 2025 - Present

Edward Jones | Saint Louis, MO

- Architected and deployed scalable batch and real-time ETL/ELT pipelines using PySpark, Apache Kafka, Azure Databricks, and Snowflake, processing 500GB+ of financial data daily to support analytics and regulatory reporting.
- Designed event-driven ingestion architecture integrating Kafka streams with batch workflows, achieving sub-5-minute end-to-end latency for time-sensitive transactional metrics.
- Engineered 20+ modular dbt transformation models, consolidating fragmented business logic into a single source of truth — reducing ad-hoc SQL requests from business teams by ~40% and cutting analyst data prep time from ~4 hours to under 20 minutes per report cycle.
- Engineered dimensional star schema models in Snowflake and Redshift for 5+ finance and operations dashboards, reducing average Tableau and Looker query execution time from 45 seconds to under 8 seconds.
- Implemented incremental processing, partition pruning, and automated reconciliation for 500GB+ daily financial data loads, cutting refresh latency by 25% while sustaining 95%+ SLA compliance across 10+ downstream reporting pipelines.
- Optimized PySpark and Snowflake workloads through advanced performance tuning and query clustering, reducing pipeline execution time by 30% and eliminating 3+ recurring SLA breach incidents per month.

Associate Data Engineer

Feb 2021 – Sep 2023

KreditBee | Bengaluru, India

- Architected end-to-end batch and streaming pipelines across Azure Data Factory, Airflow, Kafka, and Snowflake to ingest and transform 200GB+ of daily financial data, enabling real-time loan origination decisioning and regulatory reporting for 3M+ active customers.
- Developed high-performance PySpark jobs in Azure Databricks to process 5M+ financial records daily, enabling credit risk modeling and customer segmentation initiatives.
- Integrated Kafka event streams with Databricks Structured Streaming, reducing critical KPI data latency from 4 hours to under 15 minutes.
- Designed relational and dimensional data models standardizing logic across 15+ reporting tables, reducing average query complexity by 35% and enabling 3 business teams to independently access analytics without engineering support.
- Implemented automated data validation, reconciliation checks, and error-handling frameworks across 10+ pipelines, decreasing reporting discrepancies by 25% and reducing manual data correction effort from 6 hours to under 1 hour per week.
- Tuned Snowflake and Azure Synapse workloads using clustering, indexing, and materialized views, improving performance by 30% and reducing compute costs by \$15K annually.

Data Engineer Intern

Apr 2020 - Jan 2021

KreditBee | Bengaluru, India

- Built Spark-based batch pipelines in Azure Databricks using PySpark to process daily customer transaction and loan data for 4+ enterprise BI reports, reducing manual data preparation by 60%.
- Optimized SQL queries in Azure Synapse Analytics through indexing and query refactoring, improving execution time by 15–20% for high-frequency analytical workloads.
- Automated ingestion from 10+ third-party sources using Python, REST APIs, and Azure Functions, reducing manual data entry by 80%.
- Designed and deployed 25+ data validation rules and exception-handling workflows, reducing pipeline failure incidents by 15% and cutting manual QA intervention by ~3 hours per week.
- Created comprehensive data dictionaries and pipeline architecture docs for 10+ data sources, reducing onboarding time for new engineers by an estimated 30% and improving data governance audit readiness.

EDUCATION

- M.S. in Information Technology Management**
Webster University - Saint Louis, MO

2023 - 2025

CERTIFICATIONS

- Data Engineering with Google Cloud Specialization - **Coursera**
- Data Warehousing for Business Intelligence - **Coursera**
- Apache Spark Programming with Databricks - **Coursera**
- Google Cloud Data Engineering, Big Data, and Machine Learning Fundamentals - **Coursera**
- Data Processing with Python and SQL - **Coursera**