

SravanKumar Gatla

+ 1 (913) 608-4237 | sravankumargatla.work@gmail.com | [linkedin.com/in/sravankumargatla/](https://www.linkedin.com/in/sravankumargatla/)

Professional Summary

Data Engineer with 7 years of experience architecting scalable **AWS** and **Azure** pipelines using **Spark**, **Databricks**, and **Snowflake**. Expert in **ETL/ELT** orchestration with **Azure Data Factory**, **Synapse**, **Airflow**, **Kafka**, and **NiFi**, featuring **Terraform**-driven **IaC** automation. Proficient in **Python** and **SQL** for distributed processing across **Redshift**, **Oracle**, and **NoSQL** environments. Specialized in **GenAI** and **RAG** implementation using **LangChain** and **OpenAI APIs**, delivering **Power BI** insights and managing production via **CI/CD**, **Docker**, and **Kubernetes**.

Technical skills:

Programming Languages: J2EE, SQL, Python, Unix Shell Scripting. | **AI/ML & Generative AI:** TensorFlow, PyTorch, Scikit-learn, Hugging Face, LangChain, OpenAI API, NLP, Predictive Modeling | **Databases:** MS-SQL SERVER, Oracle, MySQL, PostgreSQL, DB2, Cassandra, MongoDB. | **Big Data Technologies:** Apache Spark, Apache Flink, Apache Hive, Apache HBase, Generative AI | **Data Streaming, Data Warehousing, Data Visualization:** Apache Kafka, Apache Storm, Apache Flink, Apache Beam, AWS Firehose, Tableau, Power BI, QlikSense, Looker, AWS QuickSight, Qlik, MS Excel, AWS Redshift, Azure Synapse, Snowflake | **Cloud Distributions:** AWS (EMR, EC2, EBS, RDS, S3, Glue, Lambda, Kinesis, Firehose, SQS, DynamoDB, Redshift, ECS, IAM, KMS) Azure (Databricks, Data Lake, Blob Storage, Data Factory, SQL DB, SQL DWH, Cosmos DB, Fabric, Azure DevOps, Active Directory). | **ETL & Data Integration:** Apache NiFi, Talend, Informatica, AWS Glue, Azure Data Factory, Apache flink | **DevOps, CI/CD, Containerization, Orchestration, Monitoring tools & Version Control :** Jenkins, GitLab CI, DevOps, GitHub, AWS CloudFormation, Terraform, Docker, Kubernetes, Jenkins, Airflow. Splunk, Grafana, Git, SVN.

Professional Experience:

Data Engineer | MCKESSON, New York, NY

Mar 2023- Current

- Architected **AI-driven** data pipelines using **Python**, **Azure Databricks**, **Azure Synapse**, and **Microsoft Fabric** alongside **LangChain (RAG)** and **OpenAI APIs** to build vector-indexed knowledge bases.
- Engineered high-scale **ETL/ELT** workflows using **Python**, **PySpark (Spark SQL/Streaming)**, **Kafka**, and **NiFi** to ingest 1.5TB of RDBMS data (**SQL Server**, **Oracle**, **SQL**) into **Azure Data Lake Storage (ADLS)**, reducing query latency by 35%.
- Automated cloud orchestration and (**IaC**) using **Python**, **Terraform**, **Azure Functions**, and **Azure Logic Apps**; achieved a 30% reduction in **Azure** costs through **CI/CD (Jenkins, Azure DevOps, Git)**.
- Developed 50+ **Power BI** dashboards and **Azure Databricks SQL** reports using advanced **SQL** queries and **DAX** to provide automated, executive-ready insights for cross-functional stakeholders.
- Optimized **Snowflake** and **SQL** performance for 1PB of monthly data using advanced partitioning and clustering, increasing analytics throughput by 20% while ensuring 99.8% data integrity.
- Managed complex workflow orchestration and metadata governance using **Python**, **Apache Airflow**, and **Azure Unity Catalog**, boosting **SQL** pipeline monitoring efficiency by 40%.

Data Engineer | American Airlines, Charlotte, NC

Jan 2022- Dec 2022

- Streamlined **AWS** data operations by deploying **Python**-based **ETL** pipelines with **AWS Glue**, **AWS EMR**, and **AWS Lambda**, utilizing **AWS S3** as a centralized Data Lake for airline operations data.
- Enhanced real-time data processing capabilities using **Python**, **AWS Kinesis**, and **Spark Streaming**, enabling faster processing of flight telemetry and optimized **SQL** query execution.
- Developed complex **SQL** scripts for data transformation in **AWS Redshift** and maintained infrastructure using **AWS CloudFormation**, **Python**, and **AWS CloudWatch**.
- Orchestrated **SQL** data migrations and automated recurring data quality checks using **Python** scripts to ensure zero-loss transfers between on-premise and **AWS** environments.

Data Engineer/ Software Engineer | ACCENTURE Mumbai, IN (Client: Prudential Financial, NJ)

Feb 2020 - Jun 2021

- Architected scalable data architectures and **Python/Scala** validation frameworks using **Python**, **Spark**, **AWS Glue**, and **SQL** to process large-scale insurance datasets for **SQL**-based **OLAP** and **OLTP** environments, ensuring deep data profiling and cleansing
- Migrated legacy on-premise **SQL** databases (**MySQL**, **PostgreSQL**, **Cassandra**, **MongoDB**) to **Snowflake** and **AWS Redshift**, improving overall data accessibility.
- Automated deployment cycles and **SQL** schema migrations using **Python**, **Docker**, **Kubernetes**, and **GitLab CI/CD** on **AWS EC2** and **EKS**, reducing manual intervention by 45%.

Data Analyst at OURANOS TECHNOLOGIES PVT LTD, Hyderabad, INDIA

Jul 2018- Nov 2019

- Preprocessed insurance datasets using **Oracle PL/SQL** and extracted insights from **SQL Server** and **MySQL** to support cross-functional business reporting and data integration.
- Supported **Azure Databricks** data exploration and **Power BI** dashboard design by configuring **JDBC/ODBC** connectivity, utilizing **Python** for data validation, and developing **DAX**-driven automated refreshes for reporting accuracy.

Certifications : Azure Data Engineer Associate | Snowflake (Data Lake, Data Sharing, Data Engineering, Data Warehouse, Data Applications)

Education: Master's in data science - New Jersey Institute of Technology - Dec 2022.