

Devika Manoj Vengurlekar

Harrison, New Jersey - Open to Relocate • Email: devvg1612@gmail.com • Phone: (973) 393-7814
[LinkedIn](#) • [GitHub](#) • [Portfolio](#)

Profile Summary

- Data Engineer with **2+ years of experience** building, automating, and optimizing **data pipelines and analytical systems** across fintech domains
- Skilled in **Python, SQL, Airflow, Spark**, and cloud platforms including **AWS and Azure**
- Adept at designing **scalable ETL/ELT workflows**, implementing **CI/CD automation**, and integrating data pipelines into production systems
- Hands-on expertise in **Azure Databricks, Data Factory, Synapse Analytics, and Delta Lake** for enterprise data warehousing
- Experienced collaborating with **cross-functional teams** to deliver reliable, data-driven insights that enhance business outcomes

Technical Skills

- **Programming & Databases:** Python, SQL, MySQL, SQL Server, Oracle, NoSQL, MS Advanced Excel
- **Cloud & Infrastructure:** AWS (S3, EC2, Redshift), Microsoft Azure, GCP, Snowflake, Terraform, Docker, Serverless & Cloud Architecture, VMware
- **Data Engineering:** ETL/ELT, Data Acquisition, Data Cleansing, Data Mining, Data Wrangling, Hadoop, Spark, Airflow
- **Modeling & Analytics:** SVM, KSVM, RPART, Conceptual & Logical Data Models, Statistical Concepts, A/B Testing
- **Visualization:** Power BI, Tableau, VBA, Microsoft PowerPoint
- **DevOps & Tools:** GitLab CI/CD, Jenkins, Git, GitHub, Linux, UNIX, NGINX, Visual Studio Code

Work Experience

Cloud Engineer

Jul 2025 – Present

Forman Tech Group LLC

New Jersey, Remote

- Managed a **seamless migration from hybrid AWS/Azure infrastructure** to a consolidated Microsoft Azure environment, reducing **cloud resource overhead by 20%** with **zero downtime** during transition
- Engineered **high-performance data processing engines using Azure Databricks**, utilizing collaborative workspaces and **PySpark/Scala notebooks** to process multi-terabyte raw datasets for data science teams
- Architected robust **ETL/ELT pipelines in Azure Data Factory (ADF)**, leveraging Mapping Data Flows to orchestrate complex transformations that **increased data ingestion throughput by 35%**
- Designed enterprise-grade **data warehouses using Azure Synapse Analytics**, implementing optimized **star-schema dimensional modeling** to establish a single source of truth for corporate reporting
- Leveraged **Azure Databricks Delta Lake architectures** to enforce **ACID transactions** on the corporate data lake, **boosting downstream query performance by 45%**
- Optimized **Synapse dedicated SQL pools** and fine-tuned database indexing, **slashing daily query execution runtimes by 40%** and significantly reducing operational cloud compute costs
- Spearheaded a highly sensitive, **end-to-end data engineering project** for a private client, implementing rigorous **data governance, RBAC, and Azure Key Vault integration** to meet strict regulatory compliance
- Pioneered **modernization of legacy data pipelines** by transitioning to **Microsoft Fabric and Fabric Data Factory**, targeting a **30% reduction in data latency** via unified OneLake integration

Web Developer

Aug 2024 – Jun 2025

Rebecca Everlene Trust Company

New Jersey, Remote

- Designed and implemented **responsive pages using WordPress Elementor** and GitHub-hosted media, enhancing **performance, accessibility, and stability by 20%** during MedKids staging reviews
- Provisioned **static website hosting via Amazon S3**, configured **Route 53** for domain routing, and managed **IAM role-based access control**, enabling **99.99% staging uptime**
- Monitored website health using **Amazon CloudWatch**, troubleshooting DNS resolution and S3 permission issues to ensure **error-free deployment** and faster page loads
- Collaborated across teams using **Slack and Monday.com** to bridge front-end development with AWS cloud infrastructure, supporting testing and continuous optimization

Graduate Teaching Assistant

Jan 2023 – May 2024

New Jersey Institute of Technology

New Jersey

- Mentored **40 students** in mastering **Java fundamentals** through hands-on labs, reinforcing key programming concepts
- Assessed **20+ assignments and labs**, delivering tailored feedback that drove notable improvements in technical deliverables and academic performance

Projects

Metadata-Driven Azure Data Lakehouse | [GitHub](#)

- Built an end-to-end **Azure medallion (Bronze-Silver-Gold) pipeline**, ingesting retail data via **Azure Data Factory** into **ADLS Gen2**.
- Engineered a dynamic, **JSON-parameterized ADF pipeline** to migrate 10 CSVs through a single Copy activity, cutting migration time **~70%** (30-40 min → 7-10 min).
- Transformed and cleansed **Silver-layer** data using **PySpark** on **Azure Databricks** with Service Principal-based secure access.
- Modeled **Gold-layer** analytics tables in **Synapse serverless SQL** using **OPENROWSET()** and **External Tables**, avoiding any duplication of data in storage.
- Delivered **Power BI dashboards** on curated sales/customer data for business reporting.

Real-Time Streaming Data Pipeline | [GitHub](#)

- Designed a real-time event-driven pipeline streaming live **WebApp** data through **Azure Event Hubs**, orchestrated via **Azure Data Factory** and consolidated multiple event streams into a single processing flow.
- Built a **metadata-driven streaming framework** using **PySpark Structured Streaming** on **Azure Databricks** to process high-velocity event data in near real-time, dynamically handling multiple sources without hardcoded pipeline logic.
- Applied **SCD Type 1 and Type 2** techniques in Azure Databricks to manage dimension table updates, overwriting non-critical fields while preserving historical accuracy for key attributes.
- Modeled a **Star Schema** with **fact and dimension tables** to optimize query performance for downstream analytics and reporting.

Production-Ready Azure Data Pipeline | [GitHub](#)

- Architected a production-grade **ADF pipeline** using **Self-Hosted Integration Runtime** to migrate on-prem data to Azure, with **Git-based CI/CD** for **version-controlled deployments**.
- Built **automated Logic Apps + Azure Monitor alerting**, capturing pipeline name, status, run ID, and execution time on every run, which eliminated the need for manual **log checks**.
- Secured all credentials via **Azure Key Vault + Databricks Secret Scope** integration, removing hardcoded secrets across the pipeline.
- Built **Synapse SQL-driven Power BI dashboards** surfacing key sales/operations metrics for stakeholder-style reporting.

Education

Master of Science in Computer Science | New Jersey Institute of Technology | Newark, NJ

Sep 2022 – May 2024

Bachelor of Engineering in Computer Engineering | Xavier Institute of Engineering | Mumbai, MH

Aug 2017 – Jun 2021

Certifications

- Microsoft Fabric Data Engineer Associate (DP-700)
- Databricks Certified Data Engineer Associate