

Dhanush Chavadi

Data Engineer

Corpus Christi, TX | (361)737-6047 | dhanushchavadi1@outlook.com | LinkedIn

SUMMARY

Data Engineer with 5 years of experience designing scalable ETL/ELT pipelines, cloud-native data platforms, and modern lakehouse solutions using Python, SQL, Spark, PySpark, Databricks, Airflow, and Kafka. Experienced across AWS, Azure, and Google Cloud Platform with expertise in Snowflake, Redshift, BigQuery, and data modeling. Skilled in building high-performance data pipelines, implementing CI/CD and data quality frameworks, and integrating Generative AI technologies including LangChain, OpenAI API, and RAG into enterprise data platforms.

TECHNICAL SKILLS

Programming Languages: Python, SQL, Scala, Java, Bash

Big Data Technologies: Apache Spark, PySpark, Hadoop, Databricks, Delta Lake, Apache Iceberg

Cloud Platforms: AWS (S3, Glue, EMR, Redshift, Lambda, Athena, IAM, CloudWatch), Azure (Azure Data Factory, ADLS Gen2, Azure Databricks, Synapse Analytics), Google Cloud Platform (BigQuery, Dataflow, Pub/Sub, Cloud Storage)

Data Warehousing: Snowflake, Amazon Redshift, Google BigQuery, Azure Synapse Analytics, PostgreSQL, MySQL

ETL & Orchestration: Apache Airflow, dbt, AWS Glue, Azure Data Factory, Apache NiFi, Prefect, ETL/ELT Pipelines

Streaming & Messaging: Apache Kafka, Kafka Connect, Spark Structured Streaming, AWS Kinesis

Data Modeling: Dimensional Modeling, Star Schema, Snowflake Schema, Data Vault 2.0, Slowly Changing Dimensions (SCD), Lakehouse Architecture

DevOps & Infrastructure: Docker, Kubernetes, Terraform, Git, GitHub, GitHub Actions, Jenkins, CI/CD, Linux

Data Quality & Governance: Great Expectations, Data Validation, Data Quality Monitoring, Data Lineage, Data Catalog, Metadata Management

AI & Modern Data Engineering: LangChain, OpenAI API, Retrieval-Augmented Generation (RAG), Vector Databases, AI Agents, MCP (Model Context Protocol), LLM Applications Prompt Engineering

APIs & Data Formats: REST APIs, GraphQL, JSON, Parquet, Avro, ORC, XML

Business Intelligence: Power BI, Tableau, Amazon QuickSight

EXPERIENCE

Data Engineer | McKinsey & Company

Sep 2024 – Present

- Designed scalable ETL/ELT pipelines using Python, SQL, PySpark, Apache Spark, and Apache Airflow to ingest and transform enterprise data from 20+ source systems for analytics and reporting.
- Engineered cloud-native data processing workflows on AWS (S3, Glue, EMR, Redshift, and Athena), reducing pipeline execution time by 38% while improving data availability.
- Developed real-time streaming pipelines using Apache Kafka, Spark Structured Streaming, and AWS Kinesis to process over 5 million events daily, enabling near real-time operational reporting.
- Implemented Delta Lake and Apache Iceberg-based lakehouse architecture, improving data consistency, ACID compliance, and reducing storage redundancy across enterprise datasets.
- Designed dimensional data models using Star Schema, Snowflake Schema, and SCD Type 2 to support Power BI dashboards and enterprise reporting across multiple business domains.
- Automated data transformation and orchestration using dbt and Apache Airflow, reducing manual data processing efforts by 45% across multiple workflows.
- Integrated LangChain, OpenAI API, RAG, and Vector Databases to build AI-powered metadata search and document retrieval solutions, improving data discovery for analytics teams.
- Collaborated with cross-functional Agile teams and implemented CI/CD using Docker, Terraform, GitHub Actions, and Jenkins, accelerating deployment cycles by 40%.

Environment: Python, SQL, PySpark, Apache Spark, Databricks, Apache Airflow, dbt, AWS (S3, Glue, EMR, Redshift, Lambda, Athena), Apache Kafka, Spark Structured Streaming, AWS Kinesis, Delta Lake, Apache Iceberg, Snowflake, Terraform, Docker, GitHub Actions, Jenkins, LangChain, OpenAI API, RAG, Power BI.

Data Engineer | Avenir Technologies

Jan 2020 – Dec 2022

- Built scalable ETL pipelines using Python, SQL, Apache Spark, and Azure Data Factory to integrate data from enterprise applications into centralized data platforms.
- Developed cloud-native data pipelines using Azure Data Factory, Azure Databricks, ADLS Gen2, and Azure Synapse Analytics to support enterprise reporting and analytical workloads.

- Built reusable transformation models using PySpark and dbt, standardizing business logic and improving pipeline maintainability across multiple projects.
- Designed event-driven ingestion pipelines with Apache Kafka to process streaming data for operational analytics and near real-time reporting.
- Optimized data warehouse solutions using Snowflake, PostgreSQL, and Azure Synapse Analytics, improving complex query performance by 32%.
- Implemented automated data quality validation with Great Expectations and metadata management frameworks, increasing data accuracy and governance across analytical datasets.
- Developed REST API-based ingestion frameworks using JSON, Parquet, and Avro, reducing data onboarding time by 35% for new source systems.
- Supported infrastructure automation and CI/CD using Docker, Terraform, Git, and Jenkins, decreasing deployment effort by 30% while improving release consistency.

Environment: Python, SQL, PySpark, Apache Spark, Databricks, Airflow, dbt, AWS (S3, Glue, EMR, Redshift), Kafka, Delta Lake, Apache Iceberg, Snowflake, Docker, Terraform, Jenkins, GitHub Actions, Power BI.

EDUCATION

Master of Science in Computer Science | Texas A&M University–Corpus Christi