DHRUV AMIN

Old Bridge NJ | dhruvamin2010@gmail.com | (732) 397-8728

PROFESSIONAL SUMMARY

- Azure & Snowflake Data Engineer with 4 years of experience in data engineering, ETL development, cloud data solutions, and performance optimization.
- Proficient in Snowflake, Azure Data Factory (ADF), Synapse Analytics, Databricks, SQL, and Python, specializing in data ingestion, transformation, and processing.
- Experience designing scalable and secure cloud data solutions, optimizing data pipelines for structured and semi-structured data.
- Expertise in ETL/ELT development, integrating Azure Data Lake, Data Bricks, and Power BI to support data analytics and reporting.
- Strong background in performance tuning, cost optimization, and data modeling using Snowflake and Azure Synapse.
- Hands-on experience with CI/CD deployment using DevOps, Terraform, and ARM templates for cloud-based data solutions.
- Proficient in handling large-scale data warehousing solutions, working with real-time and batch processing workflows.
- Deep understanding of data governance, security, RBAC, and best practices in cloud data management.
- Strong domain expertise in healthcare data processing, payer/provider claims, and HIPAA compliance.

TECHNICAL SKILLS

Cloud Platforms	Microsoft Azure (ADF, Synapse, Data Lake, Functions), AWS (Redshift, Lambda)
Data Warehousing	Snowflake, Azure Synapse Analytics, Redshift, BigQuery
ETL & Data Pipelines	Azure Data Factory (ADF), SSIS, DBT, Informatica, Matillion
Programming	Python, SQL, Scala, Java
Big Data Processing	Databricks, Apache Spark, Delta Lake
Database	SQL Server, PostgreSQL, Oracle, MySQL
DevOps & CI/CD	Azure DevOps, Terraform, GitHub Actions, Jenkins
BI & Visualization	Power BI, Tableau, Looker, SSRS

PROFESSIONAL EXPERIENCE

Azure Data Engineer | Emblem Health (New York, NY)
March 2024 – Present
Leading the Azure-based data transformation project for healthcare claims processing to improve data availability, analytics, and reporting capabilities for providers and payers.

- Designed and implemented Azure Data Lake architecture for managing healthcare claims and patient records.
- Developed complex ETL pipelines using ADF and Databricks to process payer claims, enrollment, and provider network data.

- Built Snowflake data warehouse models using Star and Snowflake Schema for optimized querying.
- Implemented incremental and full load strategies for handling large-scale data ingestion.
- Developed Python scripts for automating data transformations and quality checks.
- Optimized ADF pipeline performance by leveraging parallel processing, partitioning, and metadata-driven frameworks.
- Worked on CI/CD automation using Azure DevOps, integrating Terraform for infrastructure as code (IaC).
- Ensured HIPAA compliance and implemented Role-Based Access Control (RBAC) for secure data access.
- Built Power BI dashboards that improved claims processing insights by 30%.
- Worked closely with business analysts and stakeholders to define data reporting strategies and KPIs.

Environment: Azure Data Factory, Azure Synapse Analytics, Databricks, Snowflake, Power BI, Python, SQL Server, Azure DevOps, Terraform.

Snowflake Developer | AmeriHealth New Jersey (Cranbury, NJ)
Oct 2022 – Feb 2024
Developed a centralized Snowflake data warehouse to consolidate payer claims, risk management, and fraud detection insights for improved healthcare provider risk assessment.

Key Responsibilities:

- Designed Snowflake architecture, including Virtual Warehouses, Clustering, and Micro-partitioning for efficient querying.
- Developed ETL/ELT processes using Snowflake SQL, Streams, and Tasks for real-time claims data ingestion.
- Integrated Azure Data Lake with Snowflake for structured and semi-structured data processing (JSON, Parquet, XML).
- Improved query performance by 40% through clustering strategies, caching, and query optimization.
- Built automated alerting for data anomalies and fraud detection using Python and Snowflake UDFs.
- Developed dynamic role-based access control (RBAC) policies to enhance data security.
- Designed data pipelines for real-time processing of provider transactions using Kafka and Snowflake Streams.
- Optimized storage and compute costs by implementing auto-suspend & auto-clustering mechanisms.
- Created comprehensive data models in dbt (data build tool) to enable BI and analytics.
- Provided Snowflake query tuning and troubleshooting for business users and data analysts.

Environment: Snowflake, Azure Data Lake, ADF, dbt, Python, SQL, Power BI, AWS S3, Kafka.

Intern Data Engineer | ConnectiCare (Farmington, CT)
Aug 2021 – Sep 2022

Developed data pipelines for payer claims data transformation, improving provider claims validation and fraud detection.

- Designed data ingestion pipelines for processing dental insurance claims using Azure Data Factory.
- Built Snowflake stored procedures for automating data transformation and deduplication.
- Integrated Azure Functions for real-time claim validation workflows.
- Developed Power BI reports and dashboards for claims analytics and fraud detection.
- Managed data migration from legacy SQL Server databases to Snowflake.
- Implemented data masking and encryption to ensure PHI data security.
- Improved query performance using materialized views, clustering, and indexing strategies.
- Automated Azure Blob Storage ingestion to Snowflake using external stages and COPY INTO commands.
- Created data quality checks and validation rules to improve data accuracy.
- Provided technical guidance to business analysts and stakeholders on data strategies.

Environment: Snowflake, ADF, Azure Functions, SQL Server, Power BI, Python, dbt.

EDUCATION

- Bachelor's Degree in Information Technology (NJIT)