

# Dev Mithunisvar Premraj

(857) 423-5347 | devmithunisvar@gmail.com | [Linkedin](#) | [Github](#) | [Tableau Portfolio](#)

## Education

**Northeastern University, Boston, MA | Master of Science in Information Systems** **Dec 2025**

Relevant Courses: Database Management, Business Intelligence & Analytics, Data Governance & Quality

**Anna University, Chennai, India | Bachelor of Engineering in Electrical Engineering** **Apr 2021**

Publication: Predicting Power Consumption using Machine Learning models at Springer (ICICCIS 2022)

## Professional Experience

**Role: Data Analyst | Boehringer Ingelheim | Ridgefield, USA** **Feb 2025 – Jun 2025**

- Developed and maintained 15+ **Tableau dashboards** to consolidate 50+ **KPIs** across therapeutic areas, reducing reporting time by 25% and providing executives with a single source of truth.
- Built data pipelines using **DBT, Python, and SQL** to clean and standardize 200+ data sources in **Snowflake**, implementing masking and access controls to deliver secure and governed reporting layers for analysts.
- Enhanced dashboard performance by optimizing **SQL queries**, resulting in a 30% reduction in server resource consumption and 60% decrease in load times for end-users.
- Engineered scalable data pipelines on Databricks using **PySpark** to ingest and process over 10 million HCP engagement and patient encounter records bi-weekly, reducing data latency by 40% and improving downstream reporting.

**Role: Data Engineer | LTIMindtree | Chennai, India** **Oct 2021 – Dec 2022**

- Led collaboration efforts to refine reports, develop new KPIs, and update data models to meet evolving business needs, delivering 35+ Tableau dashboards and 10 ad-hoc reports that improved visibility and informed decision-making.
- Evaluated existing data models against business requirements and designed ETL pipelines in **AWS Glue** to enrich datasets, expanding data coverage by 40% and enabling 10+ new analytical use cases.
- Redesigned third-party data validation pipelines with **SQL, Python, and Alteryx** based workflows to automate quality checks, reducing manual effort by nearly 20 hours each week.
- Developed predictive optimization solution in **Python and SQL** for 150+ SKUs, by applying **Regression and clustering** models to forecast demand, reducing stockouts by 15% and costs by \$350K.
- Improved query performance by 60% by optimizing **SQL/PySpark** scripts, enhanced scalability for datasets by refining normalization and distribution strategies, and documented solutions for maintainability and knowledge sharing.
- Built **Python and SQL** validation scripts to test KPIs and ETL transformations before production release, performing root cause analysis to identify and fix data issues early, reducing reporting errors by 30%.

## Technical Skills & Certifications

**Programming Languages:** Python, SQL, Java, R, Scala, JavaScript

**Databases / Warehouses:** Snowflake, Redshift, BigQuery

**Distributed Systems:** Kafka, PySpark, EMR, Kinesis, Docker, Git

**Big Data / ETL Tools:** Airflow, Alteryx, Azure Data Factory, DBT, Glue, Lambda

**Data Visualization:** Tableau, Power BI, QuickSight, ThoughtSpot, Excel

**Machine Learning:** Predictive Modeling, Supervised Learning, Unsupervised Learning, NLP, LLMs

**Certifications:** Tableau Certified Desktop Specialist, AWS Certified Data Engineer Associate, Snowflake SnowPro Core Certification, Databricks Certified Data Engineer Professional

## Projects

### Food Safety Analytics Warehouse & Agentic BI Assistant

- Designed a **STAR schema** using **ER Studio**, integrated over 1M+ records into Fact and Dimension tables in **Azure** through **Talend** ETL jobs and CDC pipelines, used Alteryx and Python for Data Profiling and cleaning.
- Developed PowerBI dashboards, and **Agentic AI** chatbots to aid analysts in conducting complex querying and ad-hoc visualizations based on requirements.

### AI healthcare System

- Engineered a medical recommendation app using **FastAPI** and **Streamlit**, integrating **OpenAI LLM** with **RAG** to match user symptoms and history with medical content, delivering tailored medication suggestions.
- Built containerized **CI/CD** and ETL pipelines using **Git, Docker**, and **Airflow** with **Lambda** for scraping 12,000+ medical pages; applied Pydantic and pytest for validation before loading data into Snowflake and **Pinecone**.