

Name: Nishwitha Gopu

Contact Number:408-766-4661

Email: nishwitha06@gmail.com

LinkedIn: [linkedin.com/in/nishwitha-g-aa8989337](https://www.linkedin.com/in/nishwitha-g-aa8989337)

PROFESSIONAL SUMMARY: -

Highly detail-oriented and performance-driven Data Engineer with 5+ years of experience designing and optimizing scalable data pipelines using Palantir Foundry, Spark, SQL, and Python. Skilled in ontology-based data modelling, semantic layer construction, and delivering production-ready pipelines that enable enterprise analytics and machine learning capabilities. Demonstrated ability to modernize legacy data frameworks within collaborative platforms like Foundry, while upholding strict data quality and governance practices.

- Engineered and supported scalable pipelines in Palantir Foundry Code Workbooks to ingest, cleanse, transform, and distribute large-scale datasets for analytics and operational reporting.
- Modelled semantic data layers and developed Ontologies in Foundry's Ontology Manager to represent key business domains such as Customers, Orders, and Inventory, facilitating reusable, trusted data access across departments.
- Utilized PySpark to implement advanced business logic, including nested aggregations, multi-level joins, and derived metric calculations to power data-driven applications.
- Maintained data privacy and compliance by implementing row-level security and role-based access controls within Foundry.
- Used Dataset and Ontology APIs via SDK to validate schema changes, object relationships, and metadata consistency across environments.
- Collaborated with analysts and data scientists to co-create feature-enriched datasets for machine learning, segmentation models, and time-series forecasting pipelines.
- Leveraged Palantir Foundry Python SDK and REST APIs to programmatically trigger pipelines, manage dataset versions, and automate refresh and backfill workflows.
- Applied optimization techniques including coarsening, partitioning, and parallel processing to reduce execution time and enhance system efficiency.
- Experienced in using Palantir Foundry Python SDK and REST APIs to programmatically access datasets, trigger pipeline executions, validate ontology objects, and automate operational data workflows alongside Code Workbooks.
- Authored detailed documentation covering pipeline architecture, transformation logic, and ontology best practices to accelerate onboarding and promote team knowledge sharing.
- Worked in agile delivery teams, participating in stand-ups, sprint planning, and retrospectives to deliver timely, high-quality data solutions.
- Programmatically accessed datasets and metadata using Foundry Dataset APIs (SDK) to perform reconciliation against source systems.

- Developed Slate dashboards and personalized views to support key metrics tracking, reporting, and business operations.
- Strong experience in data modeling and structuring large, complex datasets to improve performance, flexibility, and extensibility for downstream analytics and application consumption.
- Proven expertise in developing Palantir Foundry data pipelines using Python, PySpark, and SQL to ingest, transform, and curate large-scale structured and semi-structured datasets.
- Hands-on experience leveraging Palantir Foundry ontologies to model complex business entities, relationships, and metrics, enabling governed, analytics-ready datasets.
- Designed and implemented ingestion pipelines to integrate data from diverse internal and external sources, ensuring reliable, high-performance data flow and consistency across systems.
- Conducted internal knowledge transfer sessions to train new team members on Foundry tooling, performance best practices, and data modelling standards.
- Contributed to process improvements by identifying gaps, recommending enhancements, and implementing solutions to improve developer productivity and scalability.

TECHNICAL SKILLS

Data Engineering: Palantir Foundry, PySpark, Spark SQL, ETL Pipelines

Cloud & Big Data: Azure Data Factory, Azure Databricks, ADLS Gen2, AWS S3

Programming: Python, SQL, Scala

Databases: Snowflake, Azure Synapse, SQL Server, Oracle, MySQL

Streaming: Event Hubs

Tools: Git, Azure DevOps, Docker, Kubernetes

Data Modeling: Dimensional Modeling, Ontology Modeling

Visualization: Foundry Slate, Tableau

CERTIFICATIONS:

Certified **Palantir Foundry Building Your First Pipeline** Issued by Palantir Technologies:

(<http://verify.skilljar.com/c/t7q55wg64f22>)

Certified **Transforming Your Data with Code Repositories** Issued by Palantir Technologies:

(<http://verify.skilljar.com/c/sztcas8hd2g2>)

Certified **Creating Your First Data Connection** Issued by Palantir Technologies:

(<http://verify.skilljar.com/c/f94fex2cxoa5>)

Certified **Data Analysis in Contour** Issued by Palantir Technologies :
(<http://verify.skilljar.com/c/jq49uivzf26u>)

Certified **Foundry & AIP Builder Foundations** Issued by Palantir Technologies :
(<http://verify.skilljar.com/c/p4efubv8hvui>)

PROFESSIONAL SUMMARY:

IQVIA, KOP, PA

Data Engineer - Palantir Foundry -

July 2024 – Till Date

Roles & Responsibilities:

- Maintain Palantir Foundry data pipelines using Code Workbooks and PySpark for ingesting and transforming enterprise datasets.
- Implement incremental processing logic and deterministic primary keys to support reliable data ingestion and deduplication.
- Design Ontology models (Objects, Links, Actions) representing business entities such as customers, orders, and inventory.
- Built end-to-end data pipelines and analytics workflows in Palantir Foundry using Pipeline Builder for large-scale data processing and full-stack application development.
- Engineered governed ETL/ELT processes using Python and PySpark in Foundry's Code Repositories and Code Workbook, ensuring high data quality, reliability, and performance for production applications.
- Modeled and maintained Ontology models (OSDK) to support cross-domain data integration and power Foundry applications, business workflows, and real-time analytics.
- Developed scalable backend services and application APIs using Python and TypeScript with Foundry Functions, enabling integration with external client systems, leveraging Foundry SDKs.
- Built production-grade front-end applications using React and TypeScript in Foundry Workshop, Slate, and Quiver, tightly integrated with ontology data models.
- Collaborated with business stakeholders, data analysts, and data scientists to translate technical and user requirements into scalable Foundry solutions and Ontology business operational layer.
- Designed and deployed AI features using LLMs, RAG pipelines, and prompt engineering, integrating them into Foundry applications for intelligent automation.
- Leveraged Palantir Foundry APIs and SDKs (Python/TypeScript) to build custom integrations and backend services for automation and application coordination.
- Integrated Foundry applications with external client systems using REST APIs and SDKs to support enterprise system interoperability.

- Integrate data from REST APIs, SFTP, AWS S3, Azure Blob Storage, SQL Server, and Oracle into curated bronze, silver, and gold data layers.
- Build automation using Foundry SDK and APIs to monitor pipelines, enforce SLAs, and trigger retries and backfills.
- Develop Slate dashboards and operational reporting layers for business KPI monitoring.
- Implement row-level security and RBAC policies to ensure governance and compliance.
- Optimize Spark pipelines using partitioning, caching, and parallel processing to reduce processing time.
- Collaborate with analytics and data science teams to develop feature datasets for machine learning and forecasting models.

Paypal, San Jose, CA

May 2021– June 2024

Data Engineer with Palantir -

Roles & Responsibilities:

- Built high-volume ETL pipelines using Azure Data Factory and Azure Databricks (PySpark) to process multi-terabyte payment transaction datasets.
- Developed automated financial reconciliation pipelines matching merchant settlements with bank transaction Reports.
- Implemented real-time event streaming pipelines using Azure Event Hubs to track payment lifecycle events.
- Designed dimensional data models (Star Schema, SCD Type-2) for transaction analytics and historical reporting.
- Designed and implemented end-to-end data pipelines on AWS and Palantir Foundry to ingest, transform, and curate pharma data including clinical trials, drug manufacturing, supply chain, and commercial datasets.
- Built and maintained Palantir Ontology models to represent core pharma entities such as Drug, Batch, Trial, Patient, Site, Shipment, and Adverse Event, enabling business-friendly data access and decision-making.
- Leveraged Foundry pipelines, transforms, and Code Repositories (Python / PySpark) to standardize data from multiple source systems (LIMS, CTMS, ERP, CRM).
- Integrated AWS services (S3, Glue, Lambda, Redshift) with Foundry to support scalable analytics and regulatory workloads.
- Enabled real-time and near-real-time insights using Foundry operational workflows for drug supply tracking, inventory optimization, and demand forecasting.
- Implemented data quality checks, lineage, and audit trails using Foundry's governance features to support GxP, FDA, and 21 CFR Part 11 compliance.
- Collaborated with clinical, regulatory, and supply chain stakeholders to translate business requirements into ontology-driven analytical use cases.

- Designed and optimized Snowflake data warehouses using star and snowflake schemas to support fast, scalable analytics and BI reporting.
- Built real-time streaming data pipelines using Apache Kafka and Kafka Streams, enabling low-latency event processing and improved system responsiveness.
- Automated ETL workflows with Matillion and Informatica, streamlining data ingestion, transformation, and integration across diverse platforms.
- Leveraged Amazon Athena for ad hoc querying of structured and semi-structured datasets in S3, accelerating access to business-critical insights.
- Developed data validation and anomaly detection frameworks improving financial data accuracy.
- Secured sensitive financial datasets using Azure Key Vault and RBAC aligned with PCI-DSS compliance.
- Improved query performance through partitioning, compression, and optimization strategies, reducing processing time by 40%.
- Containerized data frameworks using Docker and Kubernetes (AKS) for scalable deployment.