

Certified **data scientist** with hands-on experience in ML, infrastructure, and analytics, seeking a role to drive data-informed solutions.

EDUCATION

• University of San Francisco <i>Master of Science – Data Science</i>	San Francisco, CA <i>July 2024 – June 2025</i>
• University of San Francisco <i>Bachelor of Science – Data Science</i> Minor: Business Analytics Achievements: Dean’s List for 4 semesters	San Francisco, CA <i>August 2020 – May 2024</i>

EXPERIENCE

• Qventus <i>Data Scientist</i>	Remote <i>Mar 2025 – Present</i>
• Development & Deployment of a Predictive Classification Model: Built and deployed a binary classifier using XGBoost on Amazon SageMaker, trained on 1.2M+ surgery scheduling records. The model reduced last-minute cancellations by 20% through automated alerts integrated into the hospital’s scheduling system.	
• Time-Aware Feature Engineering & Leakage Prevention: Developed 20+ patient-, surgeon-, and system-level temporal and categorical features using SQL in Snowflake; implemented rolling statistics and applied time-aware data splits to prevent leakage, leading to an 8% improvement in model AUC.	
• Data Preparation & Pipeline Engineering: Processed 580K+ multi-source records using PySpark SQL; engineered a modular pipeline with imputation, feature transformation, and threshold-based filtering. Addressed class imbalance using SMOTE with parameter tuning, optimizing for AUCPR and F_β , resulting in a 12% gain in minority-class precision.	
• Model Optimization & Deployment: Optimized training using the Adam optimizer and deployed real-time inference endpoints, enabling continuous online evaluation and integration into hospital scheduling workflows.	
• Model Monitoring & Maintenance: Detected data drift using KS test and PSI, enabling proactive retraining. Sustained model performance through iterative updates based on feedback and evolving patterns.	
• Business Impact Analysis: Quantified model impact, reducing last-minute surgical cancellations by 20% and improving scheduling efficiency.	
• Klarity Health <i>Data Scientist</i>	Hybrid <i>Oct 2024 – Mar 2025</i>
• Booking Probability Model for Targeted Outreach: Established and deployed a logistic regression model to predict booking likelihood, achieving a 96% capture rate and boosting marketing efficiency through precise targeting.	
• Provider Engagement Scoring: Built a rule-based scoring model using cohort segmentation, activity weighting, pivot tables, and advanced Excel functions to identify high-potential inactive providers and improve outreach targeting.	
• Geospatial Demand Forecasting: Applied K-Means clustering and outlier analysis on zip code-level socio-economic indicators to uncover top 20% demand regions, supporting market expansion and regional campaign planning.	
• Funnel Optimization through Conversion Tracking: Monitored weekly metrics in Looker Studio and ran A/B tests on landing pages, uncovering key drivers that increased booking conversion from 0.4% to 1.25% over 8 months.	
• Root Cause Analysis of Booking Declines: Investigated booking drop-offs via time-series analysis and cross-state comparisons; identified spam triggers and informed mitigation efforts that reduced related losses by 77%.	
• Visibility Impact Modeling: Modeled listing position effects on user engagement using gradient boosting with interaction terms, revealing a 40% lift from increased provider visibility and informing ranking policy updates.	

PROJECTS

EduGgenie - AI-Powered Personal Scheduler & Tutor	<i>Apr 2025</i>
• Conversational RAG Assistant with Gemini API: Built a Streamlit-based RAG chatbot using Gemini API to answer study queries from uploaded PDFs and calendars, reducing planning time by 30%.	
• Secure File & Calendar Integration: Supported PDF uploads and iCal URL sync via GCP Storage API, enforcing encryption and HTTPS for secure file handling.	
• Monitoring Without Traditional Model Training: Tuned RAG pipeline via prompt evaluation and logged all queries, sync attempts, and GCS errors using Google Cloud Logging.	
• Lightweight Cloud Deployment with CI/CD: Hosted on Streamlit Cloud with GitHub Actions for auto-deployments, enabling continuous delivery without custom orchestration.	

SKILLS SUMMARY

- **Machine Learning:** Supervised/unsupervised learning, deep learning (CNNs, RNNs, Transformers), NLP, feature engineering, A/B testing, model evaluation (ROC, AUC, F1), model deployment (Docker, REST APIs), and pipeline automation
- **Languages:** Python, SQL, Java, R, JavaScript, NoSQL
- **Frameworks:** Scikit-learn, PyTorch, TensorFlow, Spark, Hugging Face Transformers, XGBoost, LightGBM, MLflow
- **Tools:** Airflow, GitHub, Pandas, NumPy, Matplotlib, Seaborn, Tableau, Looker, d3.js, PostgreSQL, MySQL, MongoDB
- **Cloud Tools:** MLflow, Kubernetes, Cloud Run, Cloud Storage, Cloud SQL
- **Platforms:** AWS (S3, SageMaker, Lambda), GCP (BigQuery, Cloud Run), Snowflake, Kubernetes, Linux, Windows

PROFESSIONAL CERTIFICATIONS

AWS Certified Cloud Practitioner - Nov 2024 \ Google Data Analytics Professional Certificate - April 2024