

JAGADISH BACHINA

Data Engineer

New Jersey | +1-217(761)-8543 | [E-mail](#) | [LinkedIn](#)

Data Engineer with 3+ years of experience building cloud-based data platforms, API-driven ingestion services, distributed data workflows, and reliable database pipelines across AWS, GCP, Azure, and Snowflake environments. Strong hands-on experience with Python, SQL, PySpark, Kafka, Airflow, Databricks, BigQuery, Synapse, and cloud APIs for scalable batch and near real-time processing. Skilled in designing maintainable data services, optimizing databases and Spark workloads, improving pipeline reliability through validation and monitoring, and supporting highly available platform infrastructure for analytics and operational systems.

EDUCATION

Master of Science in Management Information Systems

University of Illinois Springfield, Springfield, IL

Jan 2024 - May 2025

Bachelor of Engineering in Electronics and Communication Engineering

Anna University, India

Aug 2017 - Aug 2021

TECHNICAL SKILLS

- **Programming:** Python, SQL, PySpark, Bash
- **Cloud Platforms:** AWS, Google Cloud Platform, Microsoft Azure
- **Cloud APIs & Services:** AWS Glue, Lambda, S3, GCP Cloud APIs, GCS, Cloud Composer, Azure Data Factory, Azure Databricks, ADLS, Azure Event Hubs
- **Databases & Warehousing:** Snowflake, BigQuery, Azure Synapse Analytics, Azure SQL, Data Modeling
- **Distributed Data Systems:** Databricks, Kafka, Azure Event Hubs, Batch Processing, Near Real-Time Processing
- **API-Driven Services:** REST API Ingestion, Event-Driven Processing, Python Automation, Workflow Orchestration
- **Orchestration & Reliability:** Apache Airflow, Control-M, Prometheus, Grafana, Data Validation, Monitoring, Alerting
- **Engineering Practices:** Git, Docker, Kubernetes, Performance Tuning, Structured and Unstructured Data Processing

PROFESSIONAL EXPERIENCE

Data Engineer | Geak Minds Tech

Dec 2022 - Oct 2023

- Built Python and PySpark ingestion services to collect, normalize, and curate gNMI/SNMP telemetry data for compliance, engineering, and operational analytics teams.
- Designed API-based ingestion workflows that pulled external network data into centralized datasets, improving downstream availability for analytics and monitoring use cases.
- Automated distributed systems workflows with Apache Airflow and Kafka-integrated processing patterns, reducing manual intervention by 75% across recurring platform jobs.
- Developed Azure Data Factory and Databricks workflows to transform multi-source data into curated, database-ready datasets for reporting and operational analysis.
- Improved reliability of production pipelines by implementing validation checks, Prometheus/Grafana alerts, and failure visibility for latency and data-quality issues.
- Optimized PySpark and SQL transformations for high-volume telemetry workloads, improving processing efficiency across cloud-based analytical workflows.
- Supported real time telemetry processing with Apache Kafka and Event Hubs distributed workflows, enabling faster operational insight for engineering stakeholders.
- Partnered with cross-functional teams to convert requirements into production-ready pipelines, monitoring improvements, and maintainable platform data workflows.

- Developed PySpark jobs in Azure Databricks to merge multi-source data and publish metadata tables for analytics, reporting, and downstream service consumption.
- Implemented ingestion pipelines for JSON, CSV, TXT, and Avro data using Azure Data Factory and Databricks, improving readiness across multiple source systems.
- Built ETL pipelines to move data from ADLS and external systems into Azure Synapse Analytics and Azure SQL for curated database delivery.
- Built Python automation that interacted with cloud APIs across GCP and Azure services, improving pipeline execution, event-driven processing, and operational reliability.
- Used Kafka, Azure Event Hubs, and cloud messaging patterns to support streaming and near real-time ingestion across distributed data workflows.
- Developed Snowflake transformation workflows with SQL and dbt, improving maintainability, data quality, and availability of curated warehouse datasets.
- Applied validation and transformation controls to improve consistency, accuracy, and reliability of analytics-ready outputs across database environments

INDEPENDENT PROJECTS

End-to-End E-Commerce Data Engineering Pipeline

Azure, Databricks, Spark, Kubernetes, Docker

Oct 2025 - Nov 2025

- Built Python-based ingestion pipelines to load large-volume data into Azure Data Lake for downstream transformation and reporting.
- Used Azure Data Factory and Databricks workflows to standardize raw data into curated datasets for analytics consumption.
- Designed star schema models in Azure Synapse to support reporting and strengthen warehouse data modeling.
- Improved Spark job performance by tuning transformation logic and execution flow for faster large-volume processing.

Automated Cloud Data Warehouse Pipeline

GCP, Airflow, BigQuery

Dec 2025 - Feb 2026

- Built Cloud Composer workflows to move structured and unstructured data into GCS and BigQuery for scalable analytics processing.
- Created reusable Airflow DAGs for ingestion and transformation tasks, improving maintainability across GCP workloads.
- Improved BigQuery query performance through partitioning and clustering to support reporting use cases.
- Added monitoring and alerting for ingestion latency and job failures, improving reliability across GCP data workflows.

Network-Centric Fraud & Anomaly Detection Pipeline

AWS Glue, Lambda, S3

Jan 2025 - Apr 2025

- Built automated ETL workflows with AWS Glue and Lambda to ingest and transform large datasets into S3 for downstream analysis.
- Supported event-driven processing with Lambda-based workflows to handle time-sensitive data more efficiently.
- Used Python to examine high-velocity data patterns and support anomaly-focused analytical workflows.
- Created Tableau dashboards to visualize trends and detection metrics for technical and business stakeholders.