

Jaswanth Kumar Malla

Data Engineer

USA | +1 (571) 663-9498 | jaswanth10499@gmail.com | LinkedIn

SUMMARY

Experienced Data Engineer with around 4 years of experience designing and supporting scalable data pipelines, ETL processes, real-time data integration, and cloud-based data platforms across financial data environments. Leveraged expertise in Python, SQL, Databricks, Snowflake, Apache Spark, Kafka, Airflow, Azure, AWS, data governance, and AI-driven data engineering to deliver trusted, high-quality datasets for analytics and reporting. Collaborated effectively with cross-functional teams, business stakeholders, analysts, and engineers to optimize data workflows, improve data quality, ensure compliance, and support data-driven decision-making initiatives.

TECHNICAL SKILLS

Programming & Query Languages: Python, SQL, PySpark, Shell Scripting, R

Data Engineering & ETL: ETL Development, Data Pipeline Development, Data Integration, Data Transformation, Data Cleansing, Data Validation, Data Migration, CDC Pipelines, Source-to-Target Mapping, Workflow Scheduling

Big Data & Streaming Technologies: Apache Spark, Apache Kafka, Databricks Lakehouse, Delta Lake, Real-Time Data Processing, Batch Processing

Cloud & Data Platforms: Microsoft Azure, Azure Data Factory, Azure Data Lake Storage, AWS S3, Snowflake, SQL Server, GCP

Data Warehousing & Modeling: Data Warehousing, Data Modeling, Dimensional Modeling, Data Lakehouse Architecture, Master Data Management, Data Reconciliation

Data Governance & Quality: Data Quality Management, Data Governance, Microsoft Purview, Unity Catalog, Data Lineage, Metadata Management, Audit Compliance, Data Security

AI-Driven Data Engineering: OpenAI GPT APIs, LangChain, Retrieval-Augmented Generation (RAG), Pinecone Vector Database, AI-Ready Data Platforms, Feature Engineering, MLflow

DevOps, Monitoring & Visualization: Apache Airflow, Git, Azure DevOps, CI/CD Pipelines, Datadog, Grafana, Power BI, Microsoft Excel, Linux, Production Support, Performance Optimization

EXPERIENCE

Data Engineer | Charles Schwab, USA

Jan 2026 – Present

- Orchestrated enterprise cloud data platform integrating trading, portfolio, and market feeds using Kafka, Kinesis, Glue, and Snowflake, improving analytics efficiency by 35%.
- Re-architected Medallion data architecture across AWS S3, Snowflake, and Apache Spark, leveraging AI-driven optimization techniques, enhancing portfolio and ETF query performance by 45%.
- Streamlined real-time market tick and trade pipelines using Kafka and AWS Kinesis with AI-based anomaly detection, reducing data latency from 6 hours to under 10 minutes.
- Engineered regulatory-grade lakehouse platform using AWS S3, Apache Iceberg, and Great Expectations, achieving 100% audit traceability and improving compliance reporting turnaround by 40%.
- Spearheaded ELT orchestration workflows using Apache Airflow and dbt with ML-based failure prediction models, achieving 99.95% pipeline reliability and improving operational efficiency.
- Optimized data quality and governance framework using Python, AI-based anomaly detection, and reconciliation models, improving trading and investment reporting accuracy to 98%.

Jr. Data Engineer | HCL Technologies, India

Apr 2020 – Nov 2023

- Consolidated enterprise lending and payment data pipelines using SQL, Python, Azure Data Factory, Azure Data Lake, and Spark, increasing data availability by 35% across 20M+ records.
- Elevated reporting accuracy by 32% through integrating customer, loan, repayment, and transaction datasets, strengthening reconciliation processes, master data management, and financial reporting consistency.
- Advanced risk analytics capabilities by 30% through implementing ETL pipelines, dimensional data modeling, data warehousing, and real-time processing frameworks supporting payment monitoring operations.
- Optimized transaction ingestion and transformation workflows using Kafka, Spark, Azure Data Services, and workflow orchestration tools, reducing processing delays by 28% across 500+ datasets.
- Strengthened reporting efficiency by 40% through developing automated data validation, data cleansing, source-to-target mapping, and Power BI reporting solutions for 150+ stakeholders.
- Expanded transaction visibility by 38% through deploying real-time data integration, batch processing, and performance tuning solutions, supporting settlement tracking across 8M+ payment transactions.
- Reinforced production support reliability and governance compliance by 34% through analyzing large-scale financial datasets using SQL, Python, Azure DevOps, Git, and data quality frameworks.

EDUCATION

Masters in Data Modeling/Warehousing and Database Administration

Dec 2025

George Mason University, USA