

SUMMARY

Data Engineer with 4+ years of experience building scalable data pipelines and cloud-based architectures to support enterprise analytics. Proficient in Python, SQL, Spark, and AWS (Redshift, Glue, S3), with a strong focus on data modeling, ETL design, and performance optimization. Adept at transforming complex datasets into actionable insights and enabling data-driven decision-making across business functions.

SKILLS

Programming & Scripting: Python, Scala, SQL, JavaScript, Bash/Shell Scripting, SAS

Big Data Technologies: Hadoop, Apache Spark, PySpark, Hive, Kafka

Cloud Platforms & Services: AWS (S3, Redshift, Lambda, EC2, EMR), Azure (Data Lake Storage, Synapse)

Databases & Storage: MySQL, PostgreSQL, MongoDB, Cassandra

Data Warehousing & Modeling: Star Schema, Snowflake Schema, Dimensional Modeling

Data Visualization & Reporting: Tableau, Power BI, Tabular Models, Excel

ETL & Data Integration Tools: Apache Airflow, Informatica, Talend, dbt, Apache NiFi

Data Quality & Governance: Data Validation, Data Cleansing, Data Lineage, Databricks

Version Control & Collaboration: Git, GitHub, Bitbucket, Jira, Confluence

Containerization & DevOps: Docker, Kubernetes (EKS/GKE), Terraform, Jenkins

Methodologies: Agile, Scrum, CI/CD, Continuous Delivery, Stakeholder Collaboration, Process Improvement

PROFESSIONAL EXPERIENCE

Data Engineer, Morgan Stanley

Jul 2024 – Present

- Developed 15+ scalable ETL pipelines using PySpark and Airflow to process 5+ TB of trading and compliance data daily, ensuring consistent data availability for analytics and reporting workflows.
- Built validation routines in Python to scan trading data for integrity issues, identifying over 500 anomalies monthly and significantly improving the quality of regulatory reporting data.
- Led migration of 30+ legacy batch pipelines to AWS Glue and S3-based architecture, resulting in 25% reduction in infrastructure costs and faster recovery time in case of failures or outages.
- Optimized Redshift schemas through table partitioning, distribution keys, and compression, reducing query latency by 40% and enhancing performance for over 200 concurrent users.
- Integrated Kafka and Spark Streaming to enable near real-time ingestion of market and order data, decreasing reporting latency by 60% and supporting time-sensitive portfolio risk analytics.
- Published 20+ curated datasets to Redshift, enabling Tableau-based self-service reporting for compliance and risk teams and reducing data dependency on engineering.

Data Engineer, Hexaware Technologies

Jan 2021 – Dec 2022

- Delivered large-scale data integration pipelines using Talend and Informatica for 3 global clients, with daily transformations supporting critical KPIs for marketing, operations, and finance departments.
- Automated over 50 ETL workflows using Apache Airflow, reducing manual interventions by 90% and significantly improving pipeline reliability across production and staging environments for client projects.
- Implemented 10+ dimensional data models using Snowflake and PostgreSQL, enabling dashboard development and improved downstream analytics adoption by business users in non-technical teams.
- Deployed containerized batch data processing workloads using Docker and Kubernetes on AWS EMR, reducing deployment cycle time and increasing testing reliability during environment rollouts and upgrades.
- Wrote over 200 SQL queries per month to support ad-hoc data requests, operational reporting, and debugging needs across multiple client projects with varying data storage platforms and schema standards.
- Resolved recurring data mismatches by implementing schema validation and monitoring scripts, which led to a 70% drop in incident tickets raised over a six-month period by client business teams.

Data analyst, Cybage Software

Jul 2019 – Dec 2020

- Created more than 50 interactive dashboards in Tableau and Power BI, delivering KPIs and performance insights to client stakeholders across marketing, customer service, and operations departments on a weekly and monthly cadence.
- Extracted and transformed data from MySQL and MongoDB to support internal and external reports, handling datasets for 5 business units with requirements including time series trends, data joins, and advanced filters.
- Implemented standardized validation and data cleansing steps in Excel and SQL, improving reporting accuracy by 20% and significantly reducing rework from stakeholders during quarterly business reviews and audit checks.
- Automated Excel-based workflows using Python scripts for client-facing reporting packages, saving approximately 10 hours per week of manual spreadsheet work and reducing human error in financial data aggregation.
- Conducted exploratory and trend analysis on customer churn and subscription data, directly contributing to a 12% improvement in retention campaign efficiency through more accurate segmentation and cohort insights.
- Coordinated with product and analytics teams across 10+ reporting projects to ensure accuracy of business logic, filter conditions, and alignment with data governance practices during client-side UAT phases.

CERTIFICATES

Azure Data Engineer Associate

IBM Data Science Professional Certificate

Microsoft Certified: Power BI Data Analyst Associate

EDUCATION

Master of Computer Science

Jan 2023 – Dec 2024

University of North Texas | Denton, Texas

Bachelor of Computer Science

Jul 2016 – Jun 2020

Jawaharlal Nehru Technological University Kakinada | Kakinada, India