

SUMMARY

Results-driven Data Analyst with 6+ years of experience designing, developing, and optimizing data pipelines, dashboards, and predictive analytics across financial and healthcare domains. Skilled in Python, SQL, and cloud platforms (Azure & AWS) to build scalable and automated data ecosystems. Adept at orchestrating ETL workflows using Airflow, dbt, and Azure Data Factory, integrating modern data warehousing technologies such as Snowflake, Databricks, and Delta Lake. Proven ability to collaborate with cross-functional stakeholders, ensuring high data quality, business alignment, and regulatory compliance for Fortune 500 clients like JP Morgan Chase and NextGen Healthcare.

SKILLS

- **Programming & Frameworks:** Python (Pandas, NumPy, PySpark, Scikit-learn, Statsmodels, XGBoost, Plotly, Streamlit, FastAPI) | SQL (T-SQL, PL/SQL) | Bash | YAML
- **Data Engineering & ETL:** Apache Airflow | dbt | Azure Data Factory | AWS Glue | Snowflake | Databricks | Delta Lake | Lakehouse Architecture | Data Modeling (Star/Snowflake Schemas) | Data Quality Validation (Great Expectations)
- **Visualization & BI:** Power BI (DAX, Power Query, RLS) | Tableau | Qlik | QuickSight | Excel (PivotTables, Macros, Power Query)
- **Cloud & Big Data Platforms:** Azure (Synapse, ADF, Fabric, Data Lake) | AWS (S3, Redshift, Glue, Athena, Kinesis, Lake Formation) | Google BigQuery
- **Machine Learning & Analytics:** Regression | Classification | Time Series Forecasting | Feature Engineering | Model Interpretability (SHAP, LIME) | A/B Testing
- **Collaboration & DevOps:** Git | GitHub | Jenkins | Docker | Kubernetes | JIRA | Confluence | Agile/Scrum | CI/CD Automation

EDUCATION

- Master's in Data Science, University of Delaware, Newark, DE

CERTIFICATION

- Microsoft Certified: Power BI Data Analyst Associate
- SQL (Advanced) Certificate- hackerrank
- Dataiku Core Designer- Dataiku
- Python Data Structures- Coursera

EXPERIENCE

JP Morgan Chase

Apr 2024 – Present

Data Analytics Engineer

- Designed and deployed end-to-end ETL pipelines using Azure Data Factory, Airflow, and dbt, automating financial reporting processes and reducing manual effort by 45%.
- Developed data transformations and data models in Snowflake and Databricks, enabling unified views for regulatory and risk analytics.
- Integrated Great Expectations into ETL workflows to validate data lineage and ensure compliance with SOX and Basel III standards.
- Built interactive Power BI dashboards for CFO and Treasury teams, providing near real-time insights into liquidity and funding trends.
- Collaborated with data governance teams to implement role-based access controls (RBAC) and data masking for sensitive datasets.
- Containerized Python data validation scripts using Docker, streamlining integration within ADF pipelines and improving reusability.
- Partnered with business analysts to define key KPIs and automate data reconciliation through SQL and Airflow automation.

NextGen Healthcare

Aug 2017 – Jul 2022

Data Analyst

- Built and maintained data pipelines integrating patient, billing, and clinical data using AWS Glue, Airflow, and Snowflake, improving data availability by 40%.
- Developed statistical models (ARIMA, Prophet) for healthcare demand forecasting and patient risk stratification using Python and Scikit-learn.
- Implemented SHAP-based model interpretability, enhancing clinician trust and regulatory transparency for predictive analytics.
- Created Tableau and Power BI dashboards to monitor clinical performance and financial metrics across regional facilities.
- Migrated legacy SAS reporting workflows to Python and SQL, reducing data processing time by 30% and licensing costs by 25%.
- Automated data quality checks with Great Expectations and integrated alerts via Slack API for real-time issue notifications.
- Collaborated with cross-functional teams in Agile sprints, using JIRA and Confluence to track deliverables and UAT sign-offs.