

KRENJKUMAR KEYURBHAI PATEL

krenj.p10@gmail.com | (832) 795-1923
Snowflake Developer | Azure Data Engineer

PROFESSIONAL SUMMARY

- Cloud Data Engineer & Snowflake Developer with 4 years of experience in designing and implementing cloud-based data warehousing solutions using Snowflake, Azure Data Factory (ADF), ADLS Gen 2, and Databricks.
- Extensive expertise in building scalable data pipelines for structured and unstructured data ingestion, optimizing data transformations, and ensuring efficient analytical processing.
- Proficient in migrating on-premise legacy databases (SQL Server, Teradata, Netezza) to modern cloud-based Snowflake environments, ensuring cost efficiency and improved performance.
- Hands-on experience with Snowflake utilities such as SnowSQL, Snowpipe, Streams, Time Travel, and Zero-Copy Cloning, enhancing data governance and management.
- Strong skills in big data processing with Spark and PySpark, enabling real-time and batch data workflows for high-volume enterprise applications.
- Expertise in database modeling and schema design (Star and Snowflake Schema), optimizing query performance and indexing strategies.
- Adept in SQL performance tuning, reducing query execution times by up to 40% through clustering, caching, and efficient warehouse scaling.
- Experience integrating cloud services (AWS S3, Azure, Databricks, ADF) to develop automated data pipelines for real-time analytics.
- Strong knowledge of data governance, compliance, and security (HIPAA, GDPR) for secure handling of sensitive data in cloud environments.
- Well-versed in CI/CD automation, using Git, Azure DevOps, Jenkins, and Python scripting to ensure continuous deployment of data pipelines.
- Ability to collaborate with cross-functional teams, including data scientists, business analysts, and engineers, to translate business requirements into technical solutions.
- Highly adaptable in Agile environments, delivering data solutions on tight deadlines with an emphasis on scalability, reliability, and security.

TECHNICAL SKILLS

Cloud Platforms	Snowflake, AWS, Azure, ADLS Gen 2, AWS S3
ETL Tools	Azure Data Factory (ADF), Informatica, SnowSQL, Databricks
Big Data Technologies	Spark, PySpark, Databricks
Programming Languages	Python (NumPy, Pandas, PySpark), SQL, Java
Databases	Snowflake, SQL Server, Teradata, MySQL
Version Control	Git, Jenkins, Azure DevOps
CI/CD Tools	Jenkins, GitHub Actions
Operating Systems	Windows, Unix

PROFESSIONAL EXPERIENCE

Premier Health – Remote TX

Role: Snowflake Developer

March 2024 – Present

Developing and optimizing a healthcare analytics data warehouse using Snowflake and Azure services to enable real-time patient insights, claims processing, and clinical decision-making.

Key Responsibilities:

- Architected and implemented ETL pipelines for real-time and batch data ingestion from on-prem and cloud sources into Snowflake.
- Developed automated workflows using ADF and ADLS Gen 2, improving data integration efficiency by 40%.
- Implemented Snowpipe for continuous data ingestion, ensuring real-time availability of healthcare records for analytical teams.
- Optimized query performance and warehouse cost management, reducing query response times by 35%.
- Built custom PySpark-based data transformation jobs in Databricks, enhancing data cleansing and validation.
- Led data model optimization efforts, designing star and snowflake schema models for efficient reporting and analytics.
- Integrated Snowflake with AWS S3 for secure long-term data archival and disaster recovery.
- Automated CI/CD deployment pipelines using Azure DevOps and Jenkins, ensuring faster and error-free releases.
- Developed Python-based monitoring scripts to track pipeline health, identify data quality issues, and generate alerts.
- Ensured HIPAA and GDPR compliance, implementing role-based security policies in Snowflake.

Environment: Snowflake, ADLS Gen 2, ADF, Databricks, PySpark, AWS S3, Azure, SQL Server, Teradata.

Aurora Health – Chicago, IL

Role: Snowflake & Azure Data Engineer

Oct 2023 – Feb 2024

Developed a centralized data warehouse using Snowflake to integrate electronic medical records (EMR), patient insurance claims, and financial data for enhanced clinical and business analytics.

Key Responsibilities:

- Led the migration of hospital data from SQL Server and Teradata to Snowflake, ensuring data accuracy and security.
- Built end-to-end data pipelines in Azure Data Factory (ADF) and Databricks for automated data ingestion and transformation.
- Developed PySpark scripts for processing large-scale healthcare data, improving reporting efficiency.
- Optimized warehouse performance by tuning Snowflake clustering, partitioning, and materialized views.
- Designed high-performance data models for BI reporting and predictive analytics.
- Integrated Snowflake with AWS S3 for secure storage and disaster recovery.
- Implemented role-based access control (RBAC) policies to comply with HIPAA regulations.
- Worked closely with business teams to build Tableau dashboards for insightful healthcare analytics.
- Automated data pipeline monitoring using Python and SQL, ensuring seamless data availability.

Environment: Snowflake, ADF, ADLS Gen 2, Databricks, PySpark, SQL Server, AWS S3, Teradata.

Medtronic – Minneapolis, MN

Role: ETL Informatica Developer

Feb 2022 – Sep 2023

Designed ETL workflows for migrating legacy data warehouses into modern cloud-based Snowflake environments to support enterprise-level healthcare data analytics.

Key Responsibilities:

- Developed data pipelines using Informatica and Python for automated data extraction and transformation.
- Worked with SnowSQL and Python to build custom data processing scripts.
- Designed real-time and batch processing workflows using Spark and PySpark.
- Optimized data models to improve query response times and reporting performance.
- Assisted in migrating on-prem SQL Server databases to Azure Data Lake (ADLS Gen 2) and Snowflake.
- Implemented error-handling mechanisms and performance monitoring tools for data pipeline reliability.
- Worked with data science teams to build predictive analytics models using cleaned and transformed datasets.

Environment: Snowflake, Informatica, ADF, Databricks, Spark, PySpark, SQL Server, Teradata.

EDUCATION

Bachelor's of Computer Science & Engineering - Gujarat Technological University