

Kushal Shah | Data Engineer

Brampton, Canada | +1 (289)-541-3283 | kushalshah382@myjobsca.com | [LinkedIn](#)

SUMMARY

Data Engineer with 4+ years of experience designing, developing, and optimizing cloud-scale data platforms, Lakehouse architectures, and analytics solutions using Python, SQL, PySpark, Azure Databricks, Microsoft Fabric, Azure Data Factory, and Apache Airflow. Experienced in building Medallion Lakehouse frameworks, scalable ELT pipelines, and governed data products supporting enterprise reporting and analytics. Strong expertise in Delta Lake, Unity Catalog, Delta Live Tables (DLT), Databricks Workflows, Databricks Asset Bundles (DAB), data modeling, performance optimization, and data governance. Proven success delivering reliable, high-volume data solutions that improve data quality, accelerate processing, and support business-critical decision-making across Azure environments.

TECHNICAL SKILLS

Programming & Query Languages: Python, SQL, PySpark,

Data Engineering & Lakehouse Technologies: Azure Databricks, Apache Spark, **Apache Airflow**, Delta Lake, Delta Live Tables (DLT), Unity Catalog, Databricks SQL, Databricks Workflows, Databricks Repos, Auto Loader, Structured Streaming, Medallion Architecture, ETL Development, ELT Development, Data Modeling, Data Warehousing, Data Lakes, Batch Processing, Incremental Data Loading, Data Quality Management, Workflow Orchestration

Cloud & Data Platforms: Microsoft Azure, Azure Data Factory (ADF), Azure Databricks, Azure Synapse Analytics, Azure Data Lake Storage Gen2 (ADLS Gen2), Microsoft Fabric, Fabric Data Factory, Fabric Lakehouse, Fabric Warehouse, OneLake

Databases: PostgreSQL, SQL Server, MySQL MongoDB

Big Data & Streaming: Apache Kafka, Azure Event Hubs, Hadoop, Hive

DevOps & CI/CD: Git, GitHub, Azure DevOps, Databricks Asset Bundles (DAB), CI/CD Pipelines, Infrastructure Automation, Release Management

Analytics, Governance & Monitoring: Power BI, Tableau, Unity Catalog, Data Governance, Data Lineage, Metadata Management, Data Security, Azure Monitor, Log Analytics, Data Integration, Performance Optimization, MLflow

PROFESSIONAL EXPERIENCE

Data Engineer

Sep 2024 - Present | ON, Canada

LTIMindtree

- Developed scalable ELT pipelines using Python, PySpark, Azure Databricks, Delta Lake, and Apache Airflow, processing over 1.8TB of enterprise data daily while reducing batch execution times by 42%.
- Built and maintained dimensional data models and business-focused data marts supporting enterprise reporting across more than 120 million records, improving analytics accessibility and reporting performance.
- Automated ingestion pipelines for structured and semi-structured datasets from APIs, cloud storage platforms, and relational databases, increasing pipeline reliability to 99.8% while reducing manual processing effort.
- Optimized Azure Databricks workloads through Spark tuning, Delta optimization, partitioning strategies, workload management, and performance engineering techniques, reducing compute costs by 31%.
- Implemented automated data quality checks, validation frameworks, anomaly detection processes, and operational monitoring workflows, reducing production data issues by 48%.
- Supported Medallion Lakehouse implementations using Delta Lake, Unity Catalog, and PySpark, enabling incremental processing, historical tracking, governed data access, and trusted analytics consumption.
- Designed and maintained Unity Catalog governance frameworks, implementing role-based access controls, data lineage tracking, and centralized metadata management across enterprise Lakehouse environments.
- Implemented CI/CD deployment frameworks using Azure DevOps and Databricks Asset Bundles (DAB), improving release consistency, environment promotion, and deployment reliability across development lifecycles.

Data Engineer

May 2021 - Aug 2023 | India

Mphasis

- Developed ETL pipelines using Python, SQL, Azure Data Factory, and Azure Databricks to consolidate data from multiple enterprise systems into centralized analytical platforms.
- Modernized legacy data processing workflows through orchestration automation, Databricks migration initiatives, and transformation optimization, reducing end-to-end processing times by 45%.
- Built Delta Lake-based ingestion and transformation frameworks using Azure Databricks and PySpark, improving data reliability, scalability, and processing efficiency across enterprise reporting workloads.

- Developed scalable data warehouse and Lakehouse structures supporting reporting and analytical workloads across finance, operations, sales, and customer management functions.
- Optimized SQL workloads through indexing, query tuning, partitioning strategies, and schema enhancements, reducing report generation times by 52% across large analytical datasets.
- Developed monitoring dashboards, reconciliation processes, and validation frameworks that proactively identified data quality issues and reduced incident resolution times by 40%.
- Assisted with cloud migration initiatives, metadata management processes, and reusable framework development, reducing dataset onboarding effort by 38%.
- Partnered with cross-functional teams to gather business requirements, define transformation logic, and deliver trusted datasets for reporting and analytical use cases.

EDUCATION

Computer Engineering Technician (Diploma) Sheridan College	Sep 2025 ON, Canada
--	------------------------------

PROJECTS

City Micro-Climate Anomaly & Alert Lakehouse

- Crafted a Medallion Lakehouse on Azure Databricks, Delta Lake, Unity Catalog, and ADLS Gen2, enabling governed data processing, schema evolution, historical tracking, and scalable analytics across six cities.
- Developed real-time PySpark pipelines processing 17,000+ daily records using Auto Loader, Structured Streaming, Delta Lake, and Databricks Workflows while implementing SCD Type 2 modeling and performance optimization.
- Implemented Unity Catalog governance, centralized access controls, data lineage tracking, and automated CI/CD deployments using Databricks Asset Bundles (DAB) and Azure DevOps.
- Integrated Microsoft Fabric Lakehouse, OneLake, and Power BI semantic models to deliver near real-time environmental monitoring dashboards and automated anomaly alerting capabilities.

CERTIFICATES

Databricks Certified Data Engineer - Databricks