

Mir Muhammad Abdullah

Woodbridge, VA 22191 • (571) 250-9100 • US Citizen
mirabdullah213@gmail.com



TECHNICAL SKILLS

Programming: Python, SQL (MySQL, NoSQL), Java, R, SAS, MATLAB, C

Data Visualization: Power BI, QuickSight, Tableau, Matplotlib, Streamlit, Seaborn, Plotly, ggplot2

Machine Learning: Scikit-learn, PyTorch, TensorFlow

Big Data: Ambari, Hadoop, HBase, Spark, Pig, Hive, Cassandra, MongoDB, EMR

Cloud Platforms: AWS (SageMaker, Athena, Glue, etc.), Snowflake (Snowpark API, SQLAlchemy)

Tools: Informatica PowerCenter, Jupyter Notebook, Git, Linux, Streamlit, Ambari, Excel, ServiceNow, Cornerstone

Modeling Techniques: Regression, Classification, Clustering, NLP, Computer Vision, Neural Networks

RELATED EXPERIENCE

Tyto Athene | Herndon, VA (remote)

Data Warehouse Engineer | Aug 2024 – Present | 40 hours per week

- Led data integration initiatives for the Department of Defense's Defense Acquisition University (DAU), leveraging Informatica to create and maintain 50+ data mappings using Informatica PowerCenter, reducing data processing time by 30%
- Designed and implemented ETL processes to consolidate data from multiple sources, improving data quality and accessibility for 100,000+ DAU users
- Managed 200+ ServiceNow tickets related to data integration and warehouse operations, achieving a 100% resolution rate and improving system performance by 20%
- Collaborated on backend management of Cornerstone LMS integration, ensuring seamless data flow and reducing data inconsistencies by 40%
- Utilized DevOps practices to enhance collaboration in data warehouse development, resulting in a 35% increase in project delivery speed and a 50% reduction in deployment errors

TAA Solutions | Alexandria, VA (remote)

Data Scientist | Oct 2023 – Aug 2024 | 40 hours per week

- Developed and implemented advanced machine learning models using Python, scikit-learn, and PyTorch, resulting in a 30% improvement in predictive accuracy for client projects
- Conducted comprehensive exploratory data analysis and feature engineering on large datasets, utilizing Pandas and NumPy to uncover key insights and improve model performance
- Collaborated with cross-functional teams to translate business requirements into data science solutions, presenting findings to stakeholders using Streamlit
- Utilized AWS SageMaker to build, train, and deploy machine learning models at scale, optimizing resource utilization and reducing costs by 15%

Reliance Steel & Aluminum Co. | Cypress, CA

Data Scientist Intern | Jun 2023 – Sep 2023 | 40 hours per week

- Spearheaded a transformative project within a Fortune 300 company, leveraging Python, AWS SageMaker, SQL, Snowflake, and its Snowpark API to efficiently extract and analyze vital datasets from the company's expansive data lake, resulting in over \$10 million in annual savings through streamlined predictive analytics and customer churn modeling initiatives
- Applied advanced data engineering techniques in Python to refine and engineer data features, optimizing model performance and interpretability; conducted comprehensive RFM analysis on customer transaction data using SQL for effective customer segmentation, guiding targeted marketing campaigns and retention strategies
- Collaborated closely with cross-functional teams to define project scopes and objectives, communicated complex technical concepts effectively to stakeholders, and deployed model predictions seamlessly into a Streamlit dashboard integrated with AWS S3 for data storage and retrieval, empowering management with actionable insights to retain high-value customers and drive business growth

EDUCATION

B.S. Computational Modeling and Data Analytics

Minors: Computer Science, Statistics, and Mathematics

Virginia Tech, Blacksburg, VA

Relevant Coursework: Applied Multivariate Analysis, Intermediate Data Analytics and Machine Learning, Methods of Regression Analysis, Deep Learning

CERTIFICATIONS

AWS Machine Learning – Specialty

Dec 2023

AWS Solutions Architect – Associate

Nov 2023