

Naman Mehta

813-573-3040 | namanmehta550@gmail.com | [LinkedIn](#)

SUMMARY

Data Engineer with 7+ years of experience designing and owning production-grade data platforms on Azure and AWS. Built, governed, and optimized distributed Spark-based ELT systems supporting global logistics operations across 60+ countries, delivering measurable cost savings through scalable Lakehouse architecture, Unity Catalog governance, and CI/CD-enabled deployments.

WORK EXPERIENCE

Data Engineer, Maersk, United States

Oct 2022 – Sep 2025

- Architected and owned a production Azure Databricks Lakehouse (PySpark, SQL, Delta Lake) implementing Medallion architecture (Bronze–Silver–Gold) to integrate multi-source logistics data across **60+ countries** from 20+ production pipelines, driving **~\$100K+** in annual cost savings.
- Enforced data governance and access control using Unity Catalog, defining catalog/schema structures, managing grants across teams, and establishing data lineage visibility to support compliance and self-service analytics at scale.
- Built and maintained CI/CD pipelines using Azure DevOps and Databricks Repos with Git-based branching (feature → dev → prod), automating notebook and job deployments across environments and reducing release friction.
- Designed and owned **20+** daily batch ELT pipelines using Delta MERGE, incremental processing, and deduplication logic, orchestrated via Databricks Workflows and Azure Data Factory, meeting business-hour SLA requirements across global time zones and reducing runtime by **~15%**.
- Defined partitioning strategy, data modeling standards, and transformation frameworks to enable scalable self-service analytics via Azure Analysis Services semantic models and Power BI consumption layers with **200+ daily users**.
- Led migration of Salesforce Einstein workflows into the Databricks Lakehouse architecture, rebuilding transformation logic in PySpark/SQL and eliminating **~\$25K** in annual licensing costs.
- Owned production pipeline reliability for recurring global workflows, implementing schema evolution controls, alerting, and monitoring, alongside Spark performance tuning (query optimization, file compaction), improving job stability by **~18%**.

Analytics Engineer, SDLC Partners, United States

Nov 2020 – Oct 2022

- Built and optimized SQL-based data preparation pipelines supporting financial reporting workflows using Oracle and Alteryx.
- Refactored preprocessing logic to improve data processing reliability and reduce refresh times by **~80%** for daily datasets.
- Collaborated with reporting teams to deliver well-structured datasets supporting downstream BI consumption.

Data Engineer, Infosys Ltd, India, Pune

Sep 2015 – Jul 2018

- Engineered and optimized SQL-based data pipelines and schemas to deliver high-performance, scalable datasets for cross-functional analytics teams.
- Managed end-to-end release cycles, version control, and production troubleshooting to ensure data reliability and seamless system deployments.

EDUCATION

MS in Business Analytics and Information Systems | University of South Florida | Tampa May 2020

BS in Information Technology | Indraprastha University | Delhi May 2015

CERTIFICATIONS

Databricks Data Engineer Associate | AWS Certified Developer – Associate

TECHNICAL SKILLS

Programming Languages: Python, SQL, PySpark

Data & Lakehouse: Apache Spark, Azure Databricks, Delta Lake (MERGE, ACID, Incremental Loads), Unity Catalog, dbt, Medallion Architecture (Bronze–Silver–Gold), Data Modeling (Star Schema, OBT), CDC

Workflow & Pipelines: Databricks Workflows, Azure Data Factory, CI/CD (Azure DevOps, Databricks Repos, Git-based deployments), Airflow (working knowledge), Production Monitoring, Kafka (working knowledge)

Databases & Storage: Snowflake, Redshift, Oracle, SQL Server, MySQL, MongoDB

Cloud & Data Platforms: Azure (Data Factory, Analysis Services, DevOps), AWS (S3, EC2)

Business Intelligence & Semantic Layer: Power BI (Semantic Models, DAX), Tableau, Qlik

Leadership: Technical Leadership & Cross-Functional Collaboration, Performance Reviews, Stakeholder Alignment

Generative AI & LLM Applications: OpenAI API, LangGraph, Prompt Engineering, Structured Output Parsing