

PREETHAM REDDY GOLLAPALLI preethamreddy022@gmail.com (857) 246-8433 LinkedIn GitHub	
SUMMARY	
Highly skilled data professional with 4 years of experience in designing and implementing scalable ETL/ELT pipelines, robust data integration workflows, dashboards, visualizations, and ML models . Proficient in scripting languages like Python, SQL , cloud computing tools like AWS, Azure, Snowflake , and big data processing tools like Spark . Demonstrated success in optimizing data quality, enhancing model accuracy , and increasing data processing throughput and latency . Proven ability to perform advanced data analysis, develop insightful visualizations, and support data-driven decision-making for improved business outcomes.	
EDUCATION	
Master of Science in Engineering Management – Data Science , Northeastern University Boston, USA Relevant Coursework: Statistical Learning for Engineers, Data Mining in Engineering, Data Management in Analytics, Deep Learning and Time Series, Natural Language Processing, Artificial Intelligence, Operations Research, Engineering Probability and Statistics, Supply chain management.	
Bachelor of Engineering in Computer Science , Chaitanya Bharathi Institute of Technology Hyderabad, India Relevant Coursework: Data Structures and Algorithms, Object Oriented Programming (OOP), Software Development Life Cycle, Relational Databases (DBMS).	
SKILLS	
Scripting Languages: Python, R, SQL, NoSQL, Scala, Java, C++, SAS, Linux, shell scripting, PL/SQL, T-SQL Libraries & Frameworks: NumPy, Pandas, TensorFlow, scikit-learn, Matplotlib, Seaborn, PyTorch, PySpark Databases: MySQL, Oracle, PostgreSQL, Microsoft SQL Server, MongoDB, Cassandra, Dynamo DB, Aurora DB Distributed Systems: Apache Spark, Hadoop, HDFS, MapReduce, YARN, Flink, Presto, Trino ETL Tools: Matillion, Fivetran, Talend, Informatica, Apache NiFi, Dagster, Airbyte, SSIS, Kafka, Airflow Cloud Computing: AWS, Microsoft Azure, Google Cloud Platform Data Warehousing: Amazon Redshift, Snowflake, Big Query, Azure Synapse Analytics, Delta Lake Data Visualization Tools: Tableau, Power BI, MS Excel, Looker, Domo, MicroStrategy, Alteryx, Salesforce CRM Analytics Project Management: Scrum, Kanban, problem solving skills, written and verbal communication skills, root cause analysis, SDLC Other Tools: Ansible, Datadog, Grafana, ELK Stack, Jenkins, GitHub	
PROFESSIONAL EXPERIENCE	
Data Engineer Intern Wayfair USA - Enabled fast data availability for demand forecasting models by data ingestion of 1TB of sales, logistics, and supply chain data from Rest APIs, relational databases, CRM systems, SaaS applications using AWS Lambda, Glue ETL into Redshift, and handled automation, orchestration of data integration with Airflow. - Streamed processed data for content recommendations and customer churn, by building an ETL pipeline using AWS Kinesis, Apache Kafka for real-time data streaming, staging of 20TB of schema, data in S3 data lake, enabling 10+ data scientists and product teams to access real-time insights. - Optimized snowflake schema design, fact, and dimension tables in RDS using Kimball methodologies providing data access to analytical models in under 5 minutes, and improved data quality, the personalized recommendation, and notifications, contributing to a 15% increase in user engagement and retention. - Reduced data inconsistencies for data analysts by 50% by ensuring documentation, lineage, metadata management, and cataloging with SQL and Python based data validation, transformation, unit testing, with continuous troubleshooting and debugging, resulted in 99.9% data consistency for downstream analytics. - Resolved storage and query bottlenecks by converting JSON, CSV, and XML files into open-source formats like Parquet and Apache Iceberg and leveraging PySpark for big data distributed systems processing to optimize performance, achieving a 50% boost in scalability and data integrity. - Automated scaling for high-volume data sources using infrastructure with Terraform and managed workflows via Kubernetes, EC2 and implemented CI/CD with GitHub Actions, Docker, version control using Git, GitLab, achieving a 90% decrease in deployment time across test and production environments. - Leveraged big data tools like Hadoop, HDFS, Presto, Trino, and Hive to manage and query large-scale datasets over AWS EMR, enabling efficient data storage, processing, and complex analysis across distributed environments, which improved query performance by 40% and reduced data retrieval times for analytics. - Improved data analysis for 20,000+ records by streamlining ELT data processing pipelines by extracting data from external sources into Snowflake for data warehousing, DBT for data transformation and modeling, and FiveTran for data orchestration reducing big data processing time by 25%.	
Data Engineer KPMG India - Transformed over 100,000 market signals daily by developing ETL pipelines using tools like Informatica, Azure Data Factory, and Azure SQL Database in the KPMG Signals platform to integrate structured and unstructured data from 250+ data sources into Azure Data Lake storage for centralized data access. - Engineered financial data pipelines for the Fraud Detection System using Azure Machine Learning, and Python, Spark over Azure Databricks to process 5TB+ of transactional data, enabling real-time anomaly detection and reducing fraud-related losses by 20% through enhanced data quality. - Ensured data privacy, data governance, security, and compliance for by implementing role-based access controls and IAM policies, achieving 100% compliance with industry regulations, mitigating regulatory risks, for projects across fraud detection, AML detection, financial services, and auditing. - Implemented Financial Audit data warehouse and reporting layers in Azure Synapse Analytics using data modeling techniques such as SCD, CDC, data marts, dimensional modeling, and fact modeling that reduced report generation time by 30% and enabled quick product related decisions. - Optimized SQL Server data storage and retrieval by using T-SQL, and SQL query optimization techniques like indexing, partitioning, and caching, stored procedures, and normalization, and performed performance tuning and monitoring using Prometheus achieving a 40% reduction in query response time. - Performed data migration of legacy NoSQL databases like MongoDB, and ETL processes to Azure as part of the Tax Reporting data pipeline leveraging Cosmos DB and Azure Data Factory, building a scalable, cloud-based architecture and robust data infrastructure that reduced operational costs by 40%. - Enhanced automation of error handling and quality assurance using Python and Apache Airflow using features like logging, retries, alerting, implementing CI/CD practices using Jenkins to streamline deployment processes which reduced recovery time by 50% and improved data reliability by 30%.	
Data Analyst Connect Info solutions PVT LTD India - Integrated Excel VBA macros with Tableau and SQL to analyze logistics and transportation data, enabling progress assessment towards sustainable shipping initiatives and eco-friendly delivery methods, resulting in a 25% increase in the adoption of green delivery options within the company's supply chain. - Engineered data pipelines to build customer segmentation models using predictive modeling for personalized marketing, leveraging AWS SageMaker for machine learning, AWS Athena and QuickSight for real-time dashboards to process 5TB+ of transactional data which improved user experiences. - Developed a Python and TensorFlow based predictive inventory management model, leveraging AWS tools such as AWS SageMaker for model training and AWS Lambda for automated deployment, which optimized stock demand forecasts, reduced stock shortages by 25%, minimized overstock by 40%. - Implemented a time series forecasting model using Python libraries like ARIMA, Pandas, Numpy which enhanced demand prediction accuracy by 20% and minimized stockouts and excess inventory, resulting in a 15% reduction in operational costs and improved customer fulfillment rates. - Conducted testing and validation of data models and ETL pipelines using Python and AWS DataBrew, ensuring data accuracy and consistency across workflows. This testing process reduced data errors by 25%, increased data reliability, and provided stakeholders with confidence in analytics outputs. - Automated SQL and Python-based data cleaning routines, integrating shell scripting with AWS Step Functions and DynamoDB reducing 10 hours of weekly manual labor for the analytics team which accelerated report turnaround times by 30% enabling more timely insights for stakeholders. - Enhanced production and supply chain strategy by implementing data visualization solutions using Tableau and Power BI, creating interactive dashboards and visual reports, and conducted exploratory data analysis (EDA) to uncover critical insights, leading to a 25% improvement in inventory management. - Conducted ad-hoc reporting on customer review and feedback data using Jira and Confluence to facilitate agile collaboration across teams which directly informed the launch of 5 top-rated products in the following quarter, strengthening product-market fit and enhancing customer satisfaction.	
PROJECTS	
RAG Chatbot with Integrated ML Ops Framework , (<i>Machine Learning, AI and MLOps</i>) - Efficiently ingested up to 10K documents and optimized storage and retrieval through Delta Lake Tables and frameworks like LangChain, also scaled vector databases, graph databases and vector embedding generation and stored, prompt engineering and fine-tuning large language models (LLMs), Gen AI and GPT models such as Azure OpenAI, Llama-2 7b and Llama-2 70b, significantly improving similarity scores. - Led the creation of a CI/CD pipeline using Azure DevOps, Ansible and GitHub Actions with extensive unit and integration tests and deployed the RAG model as an API endpoint with a user-friendly interface built using Streamlit for real-time user interactions. Used Docker containers to perform containerization on the entire project and I also implemented a comprehensive monitoring framework for in-depth analysis using Elasticsearch and used MLflow to evaluate model performance.	
Predicting Electrical Grid Stability , (<i>Data Analytics and Machine Learning</i>) Developed a binary classification model using linear algebra, statistics, probability, mathematics calculus to build machine learning algorithms, achieving an accuracy of 92% in predicting the stability of electrical grids by leveraging hypothesis testing, regularization techniques, principal component analysis (PCA), hyper-parameter tuning and supervised and unsupervised learning algorithms including linear regression, classification (logistic regression, decision trees, random forests), clustering (k-means clustering), SVM and neural network and used metrics such as accuracy, precision, recall, and F1 score.	