Limeng Jiang

Jersey City, NJ | (520) 448-8168 | ljiang21@stevens.edu | LinkedIn | GitHub

EDUCATION

Stevens Institute of Technology, Hoboken, NJ — Expected: May 2026
M.S. in Data Science, GPA: 3.858
Award: Master's Scholarship Award
Coursework: Probability & Statistics, Applied Machine Learning, Natural Language Processing, Time Series Analysis

University of Arizona, Tucson, AZ — Aug 2019 – May 2023
B.S. in Information Science, dual major in Astronomy; Minors: Mathematics and Physics, GPA: 3.61
Coursework: Data Structures & Algorithms, Informatics Applications, Machine Learning, Neural Networks

SKILLS

Languages: Python (Pandas, NumPy, Matplotlib, Scikit-learn, TensorFlow, PyTorch), R (ggplot2, caret), SQL
ML/DL: Regression, Classification (SVM, k-NN), Ensemble Models (Random Forest, XGBoost), Neural Networks (CNN, RNN, LSTM, ResNet), Large Language Models (LLaMA, parameter-efficient fine-tuning)
Tools: Jupyter, VS Code, RStudio, Tableau, Git, ETL(Spark), GCP, GPU utilization, PowerPoint, Excel; basic LaTeX
Certificates: IBM Data Science, Google Data Analytics, AWS Academy Cloud Foundations + Data Engineering

PROFESSIONAL EXPERIENCE

DriverAI LLC — Peoria, AZ
Data Science Intern — Jan 2025 – May 2025

- Processed accelerometer time-series data to develop a key functionality of a a smart shopping app prototype for accurate in-store location tracking.
- Applied denoising, clustering, and piecewise modeling techniques to reduce signal noise by 15%.
- Designed and evaluated predictive models (XGBoost, LightGBM, etc.) with sensor fusion, achieving an MAE of around 0.68 standard deviation and supporting prototype deployment.

Steward Observatory — Tucson, AZ
Research Assistant Data Specialist — Oct 2022 – May 2023; Dec 2023 – May 2024

- Analyzed over 600 GB of spatiotemporal data on UA's HPC systems to uncover time-based variability in a complex physical process.
- Produced 12+ publication-quality visualizations highlighting the effects of variables such as pressure and density on system behavior.
- Built anomaly-detection models based on descriptive metrics across 1,000 time windows to replace resource-intensive multidimensional visualizations and enable faster contextual analysis.

ThirdEye Gen Inc. — Princeton, NJ
Software Engineer Intern — Jul 2023 – Oct 2023

- Benchmarked 12 open-source text-to-speech (TTS) and voice conversion models (e.g., Tortoise, RVC) across Linux and Windows.
- Evaluated GPU usage, inference speed, and fidelity to identify optimal deployment options.
- Produced technical reports and tutorials to guide the integration of models into the company's smart glass software.

RELEVANT PROJECTS

Stevens Institute of Technology — Hoboken, NJ
Scalable Airbnb Pricing Pipeline on GCP — Nov 2025 – Dec 2025

- Built an end-to-end distributed data pipeline on Google Cloud Platform (GCP) using Apache Spark to analyze Airbnb pricing factors.
- Trained and tuned Gradient Boosted Tree (GBT) regression models at scale with Spark ML, evaluating performance using RMSE and R^2.
- Performed scalability (scale-up) experiments by incrementally increasing training data size from 10% to 100% and analyzing accuracy-runtime trade-offs.

Stevens Institute of Technology — Hoboken, NJ
Drawing With LLMs — Jun 2025 – Aug 2025

- Generated 10,000 structured, semantically rich text prompts using Gen AI models (LLaMA3 and Mistral) to produce image-generating code from text descriptions.
- Fine-tuned the Code LLaMA model using the QLoRA method to generate scalable vector graphics (SVG) code from natural language prompts, reducing validation loss by 30% (from 1.15 to 0.8) despite hardware constraints.
- Deployed the fine-tuned model as a scalable web API using FastAPI to enable real-time generation of scalable vector graphics from natural language descriptions.