

Sudhanshu Gorla

Data Engineer

(945) 335-7771 ◇ sudhanshug859@gmail.com ◇ Arlington, Texas, United States ◇ [LinkedIn](#)

SUMMARY

Data Engineer with **2+ years** of hands-on experience architecting, implementing and optimizing cloud-native data platforms on AWS & Azure. Proven track record of designing automated ETL pipelines, real-time streaming and data lake/warehouse systems handling **10 TB+ daily** volumes. Expert in Python, PySpark, SQL and modern data tooling (AWS Glue, Lambda, Redshift; Azure Data Factory, Databricks, Synapse; Snowflake). Skilled in integrating IoT telemetry, predictive modelling and implementing CI/CD and governance frameworks (HIPAA-compliant). Collaborative across engineering, data science and business teams to translate requirements into actionable, measurable insights.

SKILLS

Cloud & Data Platforms AWS (Glue, Lambda, Redshift, S3, IoT Core, CodePipeline, Athena), Azure (Data Factory, Databricks, Synapse Analytics, Functions, Purview), Snowflake
Programming & Scripting: Python (Pandas, PySpark, MLflow), SQL, Shell Scripting
Data Engineering & ETL: Apache Spark, Apache Kafka, Airflow, dbt, Azure Data Factory, ETL Pipeline Design, Incremental Data Ingestion, Data Lake & Warehouse Design
Machine Learning & Analytics: AWS SageMaker, Predictive Modeling, ML Pipelines, Data Forecasting, Scikit-learn, Power BI, Tableau
DevOps & CI/CD: Jenkins, Git, Docker, Terraform, Azure DevOps, AWS CodePipeline
Data Governance & Quality: Great Expectations, Data Validation Frameworks, Metadata Management, Data Lineage, HIPAA Compliance
Monitoring & Optimization: Datadog, Query Optimization (Redshift, Snowflake), Cost Optimization (AWS & Azure), DAG Orchestration
Collaboration & Workflow: Cross-functional Team Collaboration, Requirements Gathering, Documentation, Mentoring, Agile/Scrum

EXPERIENCE

Data Engineer

Nov '24 — Present

CVS Health

Bloomfield, Connecticut, United States

- Designed & deployed end-to-end ETL pipelines on AWS (Glue, Lambda, Redshift, S3) automating ingestion/transformation of **10 TB+** patient data/day, enabling faster operational reporting.
- Integrated IoT-enabled medical devices with AWS IoT Core, enabling real-time patient monitoring and anomaly detection for proactive care.
- Developed predictive models using Python, AWS SageMaker and MLflow to forecast patient admissions with **~92% accuracy**, improving resource planning and reducing overtime.
- Streamlined data lake architecture using Amazon S3, Glue Catalog & Athena, reducing query latency by **40%** and boosting scalability.
- Built an automated Airflow ecosystem orchestrating 100+ DAGs for ingestion, transform & validation workflows — reducing manual scheduling overhead by **~60%**.
- Optimized Amazon Redshift: implemented partitioning, proper distribution keys and materialized views, improving analytics throughput by **35%**.
- Implemented cost-optimization across AWS services (S3, Redshift, Lambda) achieving an **18% reduction** in monthly cloud spend.
- Established CI/CD pipeline using Jenkins, Git & AWS CodePipeline for automated deployment of data-pipeline changes, accelerating release cycles and reducing errors.
- Ensured HIPAA-compliance through data governance frameworks and cross-functional collaboration with clinicians and data scientists.
- Authored internal documentation on workflows, design patterns and best practices — mentoring new hires in data-engineering standards.

Data Engineer

Dec '21 — Jul '23

Algocode

Pune, India

- Engineered a real-time data warehouse using Azure Data Factory, Azure Databricks, Kafka and Snowflake to support high-throughput software telemetry and event streams (100 M+ events/day).
- Migrated legacy on-premise ETL workflows (Oracle) to cloud pipelines on Azure, cutting maintenance effort by **50%** and improving deployment velocity.
- Designed schema-evolution strategies to support continuous software release updates with zero downtime / duplication.
- Built incremental ingestion frameworks processing 100 M+ application events daily with near-zero duplication.

- Implemented data lineage and metadata management via Azure Purview and dbt, enhancing governance, traceability and debugging speed.
- Created monitoring dashboards in Power BI for application performance, error analytics and KPI tracking, driving actionable business insights.
- Optimized ETL/transformation jobs using PySpark, SQL and Databricks — reduced runtime from hours to minutes, enabling faster iteration.
- Automated data-quality validation with Python, SQL test frameworks and Great Expectations, improving integrity of telemetry/logs.
- Collaborated with DevOps to containerize & deploy data services (Docker, Terraform, Azure DevOps) enabling CI/CD for pipelines.
- Conducted root-cause analysis on failed pipelines, built retry/alert mechanisms to minimise downtime and data loss.

EDUCATION

Masters of Science, Data Science, The University of Texas at Arlington	Aug '23 — May '25 Arlington, TX, United States
Bachelors of Science, Electronics and Communication Engineering, Jawaharlal Nehru Technological University Hyderabad	Aug '19 — Jun '23 Hyderabad, India

PROJECTS

- Real-Time Data Streaming Pipeline [Link](#)
- Designed and implemented an **end-to-end real-time data streaming architecture** using **Python, Apache Kafka, Spark Structured Streaming, and Cassandra**, simulating continuous data ingestion from APIs.
 - Built a **Kafka producer** to generate JSON-based user data and stream it to Kafka topics, enabling high-throughput message queuing.
 - Developed a **Spark Structured Streaming** job to consume Kafka messages, perform on-the-fly transformations, and load processed data into **Cassandra** for analytical queries.
 - Containerized the full stack using **Docker Compose**, orchestrating **Zookeeper, Kafka Broker, Spark, and Cassandra** for easy local deployment and testing.
 - Implemented **Cassandra schema design** (keyspaces, tables, partitioning) for optimized write performance and real-time read queries.
 - Ensured **fault tolerance and scalability**, allowing horizontal scaling across streaming and storage layers.
 - Added clear documentation and reproducible setup steps for **data engineers to replicate and extend** the pipeline.
 - Proposed future enhancements, including **Apache Airflow orchestration, ML-based real-time predictions, and live dashboard integration** for monitoring and analytics.

CERTIFICATIONS

Microsoft Certified: Power BI Data Analyst Associate	Nov '25
Microsoft Certified: Azure Data Fundamentals (DP-900)	Aug '25
Google Data Analytics Professional Certificate	Feb '25
AWS Academy Graduate – Cloud Foundations	Apr '22