

Rishika Reddy Baluri
Data Engineer | Azure · Databricks · Snowflake · PySpark · SQL
rishikareddybaluri@gmail.com | 410-649-9772 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Data Engineer with 5+ years of experience building scalable batch and real-time data pipelines using Azure, Databricks, and Snowflake. Experienced in designing cloud native ETL/ELT workflows, optimizing PySpark processing, and building reliable data platforms for analytics and reporting. Strong expertise in Medallion Architecture, Delta Lake, data quality validation, CI/CD automation, and incremental data processing. Proven track record improving data freshness, reducing cloud costs, and enabling business critical reporting across finance, energy, and insurance domains.

CERTIFICATIONS

AWS Certified Cloud Practitioner
Databricks Certified Data Engineer Associate
Databricks Certified Associate Developer for Apache Spark
Microsoft Certified Fabric Data Engineer Associate
SnowPro Core Certification
SnowPro Specialty Snowpark Certification

TECHNICAL EXPERTISE

Cloud & Platforms: Azure (ADF, ADLS Gen2, Databricks, Synapse Analytics), Snowflake, AWS
Data Engineering: Medallion Architecture, Delta Lake, CDC, SCD Type 2, Schema Evolution, Data Quality, ETL/ELT
Big Data & Streaming: Apache Spark, PySpark, Spark Structured Streaming, Auto Loader, Kafka, Windowing, Checkpointing
Orchestration & Transformation: Azure Data Factory (ADF), Apache Airflow, dbt, Delta Live Tables (DLT), Unity Catalog
Programming: Python, SQL, PySpark, Shell Scripting
DevOps & Governance: CI/CD, Azure DevOps, Docker, Terraform, Git, Unity Catalog, Data Contracts
Databases: Snowflake, Azure SQL, SQL Server, Oracle, Cosmos DB, Delta Lake
Visualization: Power BI, Databricks Dashboards

PROFESSIONAL EXPERIENCE

Data Engineer | Austin Energy | Texas

April 2025 – Present

- Built and maintained Azure-based data pipelines using ADF, Databricks, and Snowflake processing 1.5TB+ of data daily across 10+ source systems
- Engineered batch and API-based ingestion pipelines using ADF, Databricks, and REST APIs, improving data freshness from 4 hours to under 45 minutes
- Optimized PySpark pipelines using partitioning, caching, and incremental processing, reducing compute costs by 25% and improving runtime by 33%
- Built data validation and reconciliation checks using SQL and PySpark, reducing data quality issues by 60% and improving reporting accuracy for finance systems
- Set up pipeline monitoring and alerting using ADF alerts and logging dashboards, reducing incident detection time from 2 hours to under 15 minutes
- Worked across 10+ upstream systems handling schema changes and dependency management, reducing pipeline failures during releases by 50%
- Supported datasets powering 200+ dashboards used by 150+ business users
- Automated deployment workflows using Azure DevOps and Git, improving release reliability and reducing manual deployment effort

Data Engineer | Arrow Electronics | Colorado

April 2024 – March 2025

- Built scalable data pipelines using Databricks and ADF supporting finance and sales reporting workloads for a Fortune 500 electronics company
- Designed fact tables and SCD Type 2 dimensional models in Snowflake, reducing reporting inconsistencies by 40%
- Built API and batch ingestion pipelines using Python and ADF, eliminating 50% of manual reconciliation work
- Standardized transformation workflows using dbt across 30+ datasets, improving onboarding speed and model consistency
- Defined data contracts with upstream teams, reducing schema-related pipeline failures by 45%
- Optimized Snowflake queries and warehouse configurations to improve reporting performance and reduce compute usage

Data Engineer | 3i Infotech | India

Jan 2022 – Dec 2023

- Migrated on-premises ETL pipelines to ADF and Databricks, reducing processing time by 40%
- Built real-time streaming pipelines using Kafka and Spark Structured Streaming, reducing data latency from 45 minutes to under 5 minutes
- Developed reusable PySpark ingestion framework with schema evolution support, reducing dataset onboarding time by 50%
- Optimized Delta Lake storage using partitioning strategies, improving downstream query performance by 30%
- Improved pipeline reliability through retry logic and checkpointing, reducing failures by 35%
- Collaborated with analytics and reporting teams to deliver reliable cloud-based data solutions

Data Engineer | SunLife Financial | India

Feb 2021 – Dec 2021

- Built batch pipelines using ADF and SQL for automated financial reporting, reducing reporting cycle time from 2 days to 4 hours
- Improved data accuracy to 99.8% through validation and reconciliation checks
- Replaced manual Excel-based workflows with automated pipelines, eliminating 70% of manual processing effort
- Assisted with SQL reporting, data transformation, and ETL workflow enhancements

PROJECT

Policy Lifecycle Tracking Pipeline

Azure Blob Storage, ADLS Gen2, ADF, Databricks, Snowflake, Delta Lake, PySpark

- Built end-to-end policy lifecycle pipeline ingesting daily insurance policy data through Medallion Architecture into Snowflake dimensional tables
- Implemented Silver layer cleansing using PySpark including null handling, schema enforcement, status standardization, and validation checks
- Built Gold layer business transformations using Spark window functions to track policy lifecycle transitions and duration spent in each status phase
- Stored transformed lifecycle tracking datasets into Snowflake reporting tables for downstream analytics
- Automated raw file cleanup, Bronze archival, and audit logging workflows to improve traceability and storage efficiency

Real-Time Traffic & Roads Data Lakehouse

Azure Databricks, ADLS Gen2, Delta Lake, Auto Loader, Spark Structured Streaming, Unity Catalog, Power BI, Azure DevOps

- Built real-time Medallion Architecture lakehouse using Spark Structured Streaming and Auto Loader for incremental ingestion into Bronze Delta tables
- Engineered Silver layer transformations including derived metrics, traffic density calculations, and road category enrichment for downstream analytics
- Implemented Unity Catalog governance using catalogs, schemas, external locations, and fine-grained access controls
- Built Delta Live Tables (DLT) pipelines with automated data quality checks and lineage tracking
- Orchestrated workflows using Databricks Workflows with file-arrival and scheduled triggers
- Implemented CI/CD deployment workflows using Azure DevOps across Dev, UAT, and Production environments
- Connected Gold layer datasets to Power BI dashboards for traffic trends, vehicle analytics, and reporting

Adventure Works Azure Data Platform

ADF, ADLS Gen2, Databricks, Azure Synapse Analytics, Power BI

- Built parameterized ADF pipelines using LookUp, ForEach, and dynamic Copy Activities to ingest multiple datasets from GitHub HTTP source into Bronze layer
- Developed PySpark transformations in Databricks including date extraction, NULL handling, aggregation logic, and timestamp conversions
- Created Gold layer analytical views using Azure Synapse Analytics serverless SQL pools and OPENROWSET() over Parquet files stored in ADLS
- Built reporting-ready datasets supporting downstream analytics and visualization workflows

EDUCATION

University of Maryland Baltimore County

Master's in Information Systems

Maryland, USA

Gokaraju Rangaraju Institute of Engineering and Technology

Bachelor of Technology, Information Technology