

## SUMMARY

**Data Engineer with 5+ years of experience** building cloud data pipelines, analytics engineering frameworks, and AI-ready data platforms across telecom, healthcare, and finance. Skilled in PySpark, SQL, Databricks, Snowflake, dbt, Delta Lake, Airflow, AWS, Azure, and GCP. Experienced in building Medallion Architecture data layers, data quality frameworks, curated analytics datasets, semantic reporting layers, and AI-assisted analytics workflows using Snowflake Cortex and Power BI

## SKILLS

- **Programming:** Python, PySpark, SQL, Scala, Bash/Linux, REST APIs
- **Data Engineering:** Apache Spark, Airflow, Kafka, NiFi, Hadoop, ETL/ELT, Databricks, dbt, Parquet, JSON, Delta Lake
- **Databases:** Snowflake, PostgreSQL, MySQL, Microsoft SQL Server, Oracle
- **Cloud Platforms:** AWS (EC2, S3, EMR, Glue, Redshift, MSK, RDS), Azure (Databricks, Delta Lake), Google Cloud Platform (GKE, Cloud Composer, BigQuery)
- **Data Visualization & Statistics:** Power BI (DAX, Power Query), Tableau, Grafana, Splunk
- **Developer Tools & Workflow:** Git, Docker, Kubernetes, CI/CD, BitBucket, Jenkins, Postman, JIRA, VS Code
- **Analytics Engineering:** Data Modeling, Dimensional Modeling, Star Schema, Data Warehousing, Semantic Layer, Data Governance, Data Lineage, Data Quality
- **AI / ML Data Engineering:** Snowflake Cortex, AI-ready datasets, AI-assisted analytics, ML-ready datasets, feature engineering, unstructured data processing
- **Compliance:** HIPAA, PHI Data Handling, RBAC, Data Encryption, Access Controls.

## PROFESSIONAL EXPERIENCE

### T-MOBILE

#### Data Engineer

**Remote, US**  
**Jun 2024 - Present**

#### Project 1: Enterprise Data Pipeline Engineering for Risk Analytics and Asset Intelligence

- Built an enterprise **Databricks and Delta Lake pipeline framework** to centralize threat intelligence, vulnerability, CVE, and asset data from **10+ source systems**, giving analytics teams a trusted foundation for risk reporting and asset intelligence.
- Developed end-to-end **ETL/ELT workflows** using **PySpark, Databricks jobs, Python, REST APIs, and Azure services**, reducing manual dataset preparation by **50%** and improving the speed of source onboarding by **40%**.
- Designed Medallion Architecture using Azure Databricks and Delta Lake to organize raw, cleansed, enriched, and curated risk analytics datasets across bronze, silver, and gold layers, reducing downstream reporting rework by 35%.
- Implemented automated data quality checks for duplicate detection, schema validation, null checks, and source-to-target reconciliation, reducing data issues by 45% before reporting consumption.
- Orchestrated recurring pipeline schedules using Airflow DAGs with job dependencies, retries, failure monitoring, and automated refresh workflows to improve pipeline reliability.
- Built curated analytics datasets for leadership dashboards and security reporting, helping teams track vulnerability exposure, asset risk, and threat intelligence trends from trusted data sources.

#### Project 2: AI-Enabled analytics and Reporting Platform

- Developed **AI-ready datasets** by transforming risk, vulnerability, asset, and operational data into curated Snowflake tables and views, improving investigation readiness by **40%** for analytics, reporting, and AI-assisted review workflows.
- Integrated **Snowflake Cortex LLM functions** into SQL workflows to generate AI-assisted summaries from long-form risk content, reducing manual review time by **30%** and helping analysts prioritize high-risk assets faster. Developed AI-ready datasets by transforming risk, vulnerability, asset, and operational data into curated Snowflake tables and views for analytics, reporting, and AI-assisted review workflows.
- Built feature engineering pipelines using Python, PySpark, and Spark SQL to transform raw operational data into ML-ready feature sets for data science and AI/ML teams.
- Built reusable SQL transformation logic and semantic reporting layers to standardize business rules, risk indicators, and KPI definitions across Power BI dashboards.
- Automated recurring analytics refresh workflows using SQL procedures, scheduled transformations, Airflow jobs, and source data validations, saving 6+ hours weekly in recurring risk reporting.
- Delivered Power BI/Tableau dashboards covering vulnerability exposure, asset risk, threat trends, platform metrics, and operational KPIs, reducing manual reporting effort by 50%.

## CHARTER COMMUNICATIONS

#### Data Engineer

**Denver, US**  
**Jun 2023 - Aug 2023**

- Developed **Python automation workflows** to collect, parse, and process operational data from logs, APIs, and infrastructure systems, reducing manual monitoring effort and improving visibility for engineering teams.
- Optimized existing **Python and shell scripts**, reducing execution time by **45%** and improving reliability of recurring data collection, parsing, and validation workflows.
- Automated infrastructure data collection across **AWS EC2, S3, REST APIs, and internal systems**, reducing manual provisioning and tracking effort by **35%** using Python and Ansible.
- Built **Splunk and Grafana dashboards** to monitor logs, metrics, automation failures, and system health, enabling faster issue detection and root-cause analysis across distributed environments.
- Implemented alerting and self-healing automation logic to detect parsing failures and operational anomalies, helping teams identify issues earlier and reduce reactive troubleshooting.

## COGNIZANT TECHNOLOGY SOLUTIONS

Data Engineer

Hyderabad, India

Jan 2022 - Aug 2022

Client: JPMorgan Chase & Co, USA.

- Built Spark-based data migration workflows for 40+ banking application datasets, reducing manual reconciliation effort by 40% using Apache Spark, SQL, MSSQL, Oracle, and AWS S3.
- Built Spark-based data migration workflows for 40+ banking application datasets using Apache Spark, SQL, MSSQL, Oracle, GCP Cloud Storage, and BigQuery, reducing manual reconciliation effort by 40%.
- Developed SQL and Spark SQL transformation logic to clean, map, and synchronize branch location data from MSSQL to Oracle and BigQuery, improving migration accuracy by 35% across downstream banking systems.
- Implemented source-to-target validation checks including record counts, duplicate detection, null checks, and data mismatch analysis, reducing migration defects by 30% before production release.
- Used GCP Cloud Storage as a staging layer for migration extracts and validation files, improving traceability by 25% and supporting scalable data movement between legacy databases and cloud analytics platforms.
- Built Dataflow-based data movement workflows to load validated banking datasets into BigQuery, enabling scalable reporting and faster access for downstream analytics teams.
- Optimized SQL queries, indexes, and execution plans, improving API response time by 27% and enabling faster branch data access for 40+ internal banking applications.

## VIRTUSA CONSULTING SERVICES

Data Engineer

Chennai, India

Jun 2019 - Jan 2022

Client: IBM Watson Health, USA

- Built and maintained 8+ enterprise healthcare data pipelines processing claims, member eligibility, provider, and FHIR-based datasets using Apache Spark, Scala, Spark SQL, AWS S3, and IBM Cloud Object Storage.
- Developed scalable batch and real-time ingestion frameworks using Apache Kafka, Apache NiFi, and Spark Streaming, reducing data availability delays by 30%.
- Designed and optimized ETL/ELT workflows to transform high-volume healthcare datasets into Parquet and Delta Lake formats, improving query performance and reporting scalability.
- Implemented comprehensive data quality frameworks including reconciliation, schema validation, duplicate detection, and data profiling, improving reporting accuracy and data trust.
- Built and optimized Hive-based data models using partitioning and bucketing, significantly improving query performance for large-scale analytics workloads.
- Integrated Apache Airflow orchestration with Grafana monitoring and alerting, improving pipeline reliability and incident response capabilities.
- Supported HIPAA-compliant PHI data processing by implementing encryption, RBAC, and secure access controls across ingestion, transformation, and reporting layers.

## AWARDS, RECOGNITION, CERTIFICATIONS

- **AWS Certified Solutions Architect – Associate** | Amazon Web Services.
- Awarded **Team Recognition Badge at T-Mobile** for automating security data validation and reporting workflows, improving incident reporting turnaround time.
- Received a **Spot Award at Virtusa Consulting Services** for exceeding project delivery goals on the IBM Healthcare engagement, optimizing Scala Spark data pipelines, and improving processing efficiency while meeting critical client timelines.

## EDUCATION

UNIVERSITY OF COLORADO BOULDER

Colorado, US

Master of Science (Data Science)

Aug 2022 - May 2024

GUDLAVALLERU ENGINEERING COLLEGE

Gudlavalleru, India

Bachelor's in Computer Science

Aug 2015 - May 2019