

Rohith Reddy Yarramreddy

rohithyarramreddy17@gmail.com | (573) 647-0344

PROFESSIONAL SUMMARY

Data Analyst with 2+ years of experience analyzing large datasets, developing interactive dashboards, writing complex SQL queries, and automating reporting workflows. Proficient in SQL, Python (Pandas, NumPy), Power BI, Tableau, and Advanced Excel for transforming data into actionable business insights. Experienced in ETL processes, data quality validation, cloud analytics platforms (Azure, Google Cloud, AWS), and collaborating with cross-functional teams to support strategic business decisions. Strong background in statistical analysis, KPI reporting, trend analysis, and business intelligence. Proven ability to collaborate within Agile/Scrum teams to deliver scalable, secure, and analytics-ready data solutions that support business intelligence, reporting, and operational decision-making.

TECHNICAL SKILLS

Programming & Query Languages: Python (Pandas, NumPy, Matplotlib, PySpark), SQL, R, Java

Data Engineering & ETL: ETL/ELT Pipelines, Apache Airflow, Apache Kafka, Change Data Capture (CDC), Workflow Orchestration, Data Transformation, Data Integration, Data Quality Validation, Data Governance, REST APIs, FastAPI, Power Automate

Databases & Data Modeling: Google BigQuery, PostgreSQL, MySQL, MongoDB, Azure SQL, Oracle ERP, Data Warehousing, Data Marts, Dimensional Modeling, Star Schema, Snowflake Schema

Business Intelligence & Analytics: Power BI, Tableau, Advanced Excel (PivotTables, Power Query, VLOOKUP/XLOOKUP), KPI Reporting, Dashboard Development, Data Storytelling

Cloud: Microsoft Azure (Data Factory, Synapse Analytics, Fabric, Azure SQL), Google Cloud Platform (BigQuery, Cloud Storage, Cloud Composer), AWS (S3, Redshift)

Tools: Git, GitHub, Docker, Jenkins, REST APIs, FastAPI

Machine Learning & AI: Scikit-learn, TensorFlow, PyTorch, LangChain, RAG, AI Agents, OpenAI, Cursor

Methodologies: Agile, Scrum, SDLC, Cross-functional Collaboration

EXPERIENCE

Data Engineer | Missouri University of Science and Technology

Feb 2024 – Dec 2025

- Designed and developed scalable ETL/ELT pipelines using Python, SQL, Apache Airflow, and Apache Kafka, integrating Oracle ERP and NetDB into centralized cloud data warehouses, reducing data processing latency by 45%.
- Built and optimized SQL transformations, dimensional models, and data marts in Google BigQuery and Azure SQL, supporting enterprise analytics workloads while maintaining 99.9% system availability during 40% data growth.
- Automated data ingestion, workflow orchestration, and recurring reporting using Apache Airflow, Python, Power Automate, and Google Apps Script, improving data freshness by 60% and reducing manual effort by 50%.
- Developed interactive Power BI dashboards, automated KPI reports, and executive reporting solutions, enabling stakeholders to monitor operational performance and make data-driven business decisions.
- Performed data quality validation, business trend analysis, and root cause analysis to improve reporting accuracy and operational efficiency.
- Collaborated with cross-functional business and technical teams in an Agile environment to deliver secure, analytics-ready datasets and improve forecasting accuracy by 15%.

Data Analyst | Rane Group

Aug 2022 – Dec 2023

- Analyzed large operational datasets using SQL and Python (Pandas, NumPy) to identify business trends, performance patterns, and operational insights, improving reporting accuracy by 20%.
- Developed interactive Power BI dashboards, KPI scorecards, and analytical reports that enabled stakeholders to monitor business performance and support strategic decision-making.
- Wrote complex SQL queries for data extraction, transformation, validation, and ad hoc analysis across multiple business functions.
- Automated recurring reports, implemented data quality validation processes, and reduced manual reporting efforts through optimized analytical workflows.
- Conducted root cause analysis, collaborated with cross-functional teams, and presented analytical findings that improved operational reporting and business decision-making.
- Utilized Advanced Excel (PivotTables, Power Query, VLOOKUP) for data cleaning, reporting, reconciliation, and business analysis.

PROJECTS

Real-Time Accident Severity Prediction System | Python, Microsoft Fabric, Azure Synapse, Power BI

- Architected a scalable ETL pipeline processing 7.7M+ records using Microsoft Fabric and Azure Synapse to integrate structured and API-based datasets into a unified analytics platform.
- Developed automated Power BI dashboards for accident hotspot analysis and weather-impact reporting while implementing end-to-end data quality validation.

Customer Segmentation & Predictive Analysis | Python | SQL | Power BI | Scikit-learn

- Processed and validated 50K+ customer records using Python and SQL, improving data quality and analytical reliability for downstream reporting.
- Developed predictive customer segmentation models using clustering and statistical analysis techniques, achieving 85% prediction accuracy and delivering actionable business insights.

Multi-LLM Claim Automated Evaluation Pipeline | Python, Pandas, Gemini API, Seaborn

- Engineered scalable parallel data ingestion pipelines processing 10,000+ records, implementing automated validation and retry mechanisms.
- Developed automated claim evaluation workflows reducing manual annotation effort by 95%, while generating analytical dashboards to measure model performance and reliability.

EDUCATION

Master of Science in Computer Science | Missouri University of Science and Technology | GPA: 3.71/4.0

Jan 2024 – Dec 2025

Courses: Data Mining, Database Systems, Information Visualization, Analysis of Algorithms, Large Language Models, Distributed Systems, Parallel & Distributed Computing

CERTIFICATIONS

Databricks Certified: Generative AI Fundamentals, Prompt Engineering Fundamentals | IBM Certified: Data Fundamentals | Microsoft Certified: Programming with Python