

SAI KUMAR MORA

Data Engineer

Profile

- **AWS Data Engineer** with 4+ years of experience designing, implementing, and optimizing large-scale data processing pipelines, cloud-based infrastructures, and real-time analytics solutions.
- Leading the development and automation of robust **EL/ETL pipelines** in **BigQuery**, streamlining **data ingestion, transformation, and curation** while collaborating with **cross-functional teams** to enable **data-driven decision-making**.
- Proficient in Python for data manipulation and analysis, **integrating AWS libraries** like Boto3, PySpark, and DBT for efficient data handling and transformation.
- Extensive experience in managing and optimizing databases, including **PostgreSQL, SQL Server, AWS Redshift, and BigQuery**, ensuring seamless data storage, retrieval, and performance optimization.
- Hands-on experience in **AWS based ETL solutions, including AWS Glue, EMR, and Databricks**, as well as building modular, reusable transformation **pipelines with DBT in BigQuery**.

Employment History

AWS Data Engineer at CitiBank, Dallas, Texas

February 2023 — Present

- Interacted with GSAR and VIN 360 product lines and **business partners** to understand data engineering opportunities, tooling and needs.
- Collaborated with Data Engineering and **Data Architecture** to design and build templates, pipelines and data products including automation, transformation and curation using best practices.
- Developed and maintained robust **EL/ELT/ETL pipelines** to integrate disparate data sources into **BigQuery for Business Intelligence and Analytics teams**, ensuring timely and reliable data availability.
- Collaborated with **cross-functional teams** to design and implement ETL Pipelines to enable data-driven solutions.
- Develop custom cloud solutions and pipelines with **GCP native tools – Data Proc, Composer, Cloud Storage, DBT and Big Query**.
- Operationalize and automate data **best practices**: quality, auditable, timeliness and complete.
- Participated in **design reviews** to accelerate the business and ensure scalability.
- Worked with Data Engineering and Architecture and Data Platform Engineering to implement **strategic solutions**.

Details

1306 s finley road
LOMBARD, 60148
United States
3312436867

Saikmoraa@gmail.com

Date of birth

01/30/1997

Links

[Linkedin](#)

Skills

Ability to Work in a Team

Microsoft Office

Apache Spark

SQL

Python

Snowflake

AWS S3

AWS Redshift

DBT

Apache Kafka

- Orchestrated complex workflows and automated end-to-end data pipelines using **Apache Airflow** and **Tekton**, ensuring optimized resource utilization and minimal downtime.
- Designed and implemented advanced SQL queries for data transformation, validation, and analysis across platforms like **Teradata**, **BigQuery**, and **GCP Cloud Storage**, achieving performance optimization and reduced query latency.
- Containerized and deployed scalable **ETL/ELT** solutions using **Docker** to streamline development, testing, and deployment processes, ensuring consistency across environments.
- Leveraged **DBT** to build modular, reusable, and auditable data transformation pipelines in **BigQuery**, enabling business stakeholders to derive actionable insights efficiently.
- Implemented GCP-native services such as **DataProc** for large-scale batch processing and **Cloud Composer** for orchestrating workflows, reducing manual interventions and improving pipeline reliability.
- Optimized Python-based data workflows to process terabytes of structured and unstructured data, improving pipeline execution time and ensuring seamless integration with cloud-native services.
- Developed monitoring and alerting mechanisms for data pipelines, leveraging GCP tools and logging frameworks to ensure data accuracy, completeness, and timeliness in production environments.

Data Engineer at Amazon Development Center, Hyderabad

December 2020 — July 2022

- Developed and optimized **big data processing tasks** using AWS EMR and PySpark, implementing scalable data analytics and machine learning models to efficiently handle large datasets.
- Managed and optimized AWS EMR clusters running **Kafka** and Apache Spark for large-scale data processing, designing and executing big data strategies to ensure effective data flow and storage management.
- Led the deployment and management of **AWS Redshift** for data warehousing solutions, collaborating with BI teams to build and maintain scalable analytics platforms that integrated seamlessly with AWS QuickSight and existing AWS data systems.
- Designed and implemented real-time data ingestion pipelines using **AWS Kinesis**, capturing and processing events efficiently while ensuring high availability, fault tolerance, and optimized throughput for large-scale data ingestion.
- Developed and maintained **ETL processes** using AWS Glue, designing automated data transformation pipelines, implementing monitoring and logging mechanisms to ensure performance tracking and quick issue resolution.

- Designed and managed high-performance **AWS RDS instances**, including PostgreSQL and MySQL, focusing on optimal database tuning, query optimization, and data integrity for scalable data storage and retrieval.
- Administered **AWS RDS for MySQL**, ensuring performance optimization, security compliance, and automated backup strategies using AWS Backup to support business continuity and regulatory compliance.
- Utilized **SQLAlchemy** for database operations on AWS, enhancing data integration and interoperability, and maintaining ORM models for efficient and structured database management.
- Developed high-performance APIs with **FastAPI**, deployed on AWS Elastic Beanstalk, ensuring asynchronous processing and high concurrency, while integrating AWS Cognito for security and AWS API Gateway for documentation.
- Implemented and managed data workflows using **Apache Airflow** on AWS, automating pipeline execution across AWS S3 and AWS Glue, ensuring data consistency and process reliability.
- Developed and maintained **CI/CD pipelines** using AWS CodePipeline and AWS CodeBuild, automating testing, build, and deployment processes for seamless software releases and updates.
- Collaborated with development teams to enforce **code quality standards** and integrate security best practices from the initial stages of development using AWS CodeCommit.
- Managed **AWS SQS** for enterprise-level messaging solutions, optimizing message processing and queue management, ensuring reliable data communication across AWS services.
- Oversaw AWS EC2 and **S3 operations**, ensuring performance optimization, cost-effectiveness, and security compliance, while implementing data storage strategies focused on access control and lifecycle management.
- Used **AWS CloudFormation** to define and deploy cloud infrastructure, automating setup processes to ensure consistency, scalability, and cost efficiency.
- Deployed and managed **AWS ECS** for containerized applications, enabling fast rollouts, seamless updates, and secure deployments, while conducting regular security assessments for compliance with industry standards.
- Designed and optimized **prompt engineering** strategies for large language models on AWS SageMaker, leveraging AI-driven solutions to enhance application performance and business insights.

Environment: AWS EMR, PySpark, AWS RDS, SQLAlchemy, FastAPI, Apache Airflow, AWS Redshift, AWS Quicksight, AWS S3, AWS Glue, AWS Kinesis, AWS SQS, AWS CodePipeline, AWS CodeBuild, AWS CodeCommit, AWS ECS, AWS CloudFormation, AWS EC2, AWS S3, SageMaker.

Associate BigData Analyst at Globallogic Technologies, Hyderabad

September 2018 — December 2020

- Gather structured and unstructured data from various sources: databases, APIs, logs, flat files, or cloud storage.
- Use basic **SQL queries, Python scripts, or data ingestion tools** (e.g., Apache Kafka).
- Remove duplicates, handle missing values, standardize formats, and correct errors.
- Use **Excel, Python (Pandas), R, or SQL** for preprocessing.
- Apply basic transformation logic to make data analytics-ready.
- Learn and support the use of **Hadoop, Hive, Spark, and HDFS**.
- Execute basic **HiveQL** queries and assist with Spark data processing tasks.
- Understand the basics of distributed data storage and computation.
- Create charts, pivot tables, and dashboards to present findings to non-technical stakeholders.
- Automate recurring reports and maintain dashboards.
- Programming: C and Java
- SQL & Query Engines: MySQL, HiveQLBig
- Data Frameworks: Hadoop, Spark, HiveData,
- Data Visualisation: Power BI, Excel.

🎓 Education

Masters, Trine University, Angola, Indiana

August 2022 — December 2024

Master's in Business Analytics is a specialized graduate program that blends **data science, statistics, and business strategy** to equip students with the tools to make **data-driven business decisions**. The program is designed for individuals aiming to become **data-savvy professionals** who can translate complex data into actionable insights.

💼 Internships

Big data Certification at Microsoft, Hyderabad

July 2018 — September 2018

⚙️ Courses

AWS Data Engineering, Prepzee

January 2022 — August 2022

Microsoft Technology Associate Certified, Microsoft Development

August 2018 — August 2018