

Saiteja Siripurapu

Senior Business Data Analyst | Data Platforms | Analytics Engineering

saisiri1420@gmail.com | +1 2818167937 | [Linkedin](#)

PROFESSIONAL SUMMARY

- Accomplished Business Data Analyst with over 9+ years of progressive experience in delivering end-to-end analytics solutions across diverse enterprise environments, focused on transforming large-scale enterprise data into structured insights for operational reporting, financial analysis, and decision support. Worked extensively across cloud data ecosystems, analytical modeling, data quality engineering, and BI delivery. Operationalized cloud-based data pipelines to ingest, standardize, and transform claims, transactions, and operational datasets from multiple enterprise systems into curated analytical structures supporting reporting and downstream analytics consumption.
- Built SQL-based analytical transformations across distributed warehouse and lakehouse environments to reconcile transactions, validate claims, and align datasets with reporting logic, ensuring consistency across source systems and analytical layers used for reporting.
- Developed multi-layered BI dashboards by integrating heterogeneous data sources into unified semantic reporting structures, enabling analysis of healthcare utilization trends, investment performance metrics, and operational KPIs across business units.
- Implemented structured quality validation frameworks to detect anomalies, missing attributes, referential inconsistencies, and schema deviations across datasets, strengthening reliability for regulatory reporting, financial analysis, and operational decision support.
- Organized and standardized data storage structures across cloud-based data lakes and warehouse environments, ensuring consistent dataset formatting, efficient retrieval patterns, and stable access for analytics, reporting, and downstream data processing workflows.
- Applied data governance and security controls to manage controlled access to sensitive healthcare and financial datasets, maintaining auditability, enforcing access boundaries, and supporting regulated data handling requirements across analytical ecosystems.
- Monitored end-to-end data pipeline execution and system performance signals to identify ingestion failures, processing delays, and workflow disruptions, ensuring continuity and reliability of analytical data delivery across enterprise reporting systems.
- Analyzed cloud resource utilization patterns across data processing workloads to identify inefficiencies in compute and storage consumption, supporting optimization of execution behavior and improving cost efficiency of analytical processing environments.
- Supported healthcare analytics operations by processing claims, enrollment, and provider datasets to generate structured insights on utilization trends, care coordination outcomes, and regulatory reporting requirements across operational healthcare systems.
- Contributed to investment analytics workflows by reconciling transaction-level and portfolio-level datasets, preparing structured analytical views, and supporting performance attribution analysis used in financial reporting and portfolio evaluation processes.
- Assisted consulting analytics engagements by preparing clean, structured datasets for predictive modeling, segmentation analysis, and churn evaluation studies, enabling data-driven insights for business performance assessment and strategic analysis use cases.
- Maintained version-controlled analytical assets including transformation logic, and reporting definitions to ensure reproducibility, structured collaboration, and consistency across data engineering and analytics workflows spanning multiple enterprise environments.
- Partnered with business stakeholders through targeted workshops and interviews to uncover operational bottlenecks, design technically viable solutions, and deliver actionable insights that directly informed strategic decision-making.
- Enabled adoption of governed self-service analytics by defining certified datasets, clarifying KPI ownership, and communicating reporting scope and refresh expectations to finance and operations teams, reducing ad-hoc analyst dependency.
- Represented analytics requirements within Agile and SAFe delivery frameworks, ensuring reporting dependencies, data readiness, and validation checkpoints were identified early and aligned with product and platform release cycles.
- Designed and maintained metadata standards, data lineage mappings, and stewardship documentation linking reported metrics to source systems, supporting regulatory reviews and audit confidence without disrupting day-to-day analytical workflows.
- Applied domain knowledge across financial services compliance, insurance claims processing, retail customer analytics, and IT service reporting to design metrics and dashboards that respected regulatory, operational, and data availability constraints.

TECHNICAL SKILLS

- Reporting & Visualization:** Power BI, Looker, Tableau, QuickSight, Excel
- Databases & Querying:** SQL, Azure Synapse Lakehouse, Snowflake, Redshift, Athena, Power Query
- Warehousing & Storage:** Microsoft Fabric, Azure Data Lake Storage Gen2, Amazon S3, Ad-hoc Analysis, Trend Analysis.
- Automation & Scripting:** PowerApps, Power Automate, RPA, Azure Logic Apps, SQL Scripting, Excel VBA.
- Programming Languages & Libraries:** Python, Polars, Pandas, NumPy, scikit-learn
- ETL & Data Integration:** AWS Glue, AWS Lambda, Azure Data Factory, Azure Synapse Analytics
- Governance & Security:** Azure Purview, Microsoft Purview, Azure Key Vault
- Cloud Platforms & ML:** AWS SageMaker, Microsoft Azure, AWS
- Business Analysis & Strategy:** Requirements Elicitation, Design Thinking, BPMN, UML, Gap Analysis.
- Monitoring & Operations:** SAFe, Agile Scrum, Jira Portfolio, Confluence, SDLC, KPI & OKR Modeling. Azure Monitor, CloudWatch, ACM, Git, GitHub, Jupyter Notebooks, Root Cause Analysis.

WORK EXPERIENCE

Senior Business Data Analyst | Centene Corporation, St. Louis, MO

Jul 2025 – Present

Project Title: Enterprise Healthcare Analytics & Claims Performance Reporting Platform

Project Overview: Worked on a healthcare analytics modernization initiative to centralize claims, member, provider, and enrollment data from multiple healthcare systems into Microsoft Fabric and Azure Synapse. The solution enabled business users, operations teams, and leadership to monitor claims performance, healthcare utilization, member trends, provider effectiveness, and regulatory reporting through self-service Power BI dashboards.

- Collaborated with business stakeholders, claims teams, and healthcare operations leaders to gather reporting requirements and define KPIs for claims processing, member utilization, and provider performance.
- Analyzed healthcare claims, member enrollment, and provider datasets using SQL to identify utilization trends, operational bottlenecks, and opportunities for process improvement.
- Developed Power BI dashboards and executive scorecards to track claims turnaround time, denial rates, provider performance, member growth, and healthcare utilization metrics.
- Created SQL-based analytical datasets and business rules to standardize reporting across multiple healthcare systems.
- Performed data profiling, reconciliation, and validation activities to improve data quality and reporting accuracy.
- Partnered with data engineering teams to support ETL development using Microsoft Fabric, Azure Data Factory, and Azure Synapse.
- Conducted root cause analysis on reporting discrepancies and collaborated with source system owners to resolve data issues.
- Defined business requirements, documented functional specifications, and translated stakeholder needs into actionable analytics solutions.
- Implemented data governance practices including data lineage, KPI definitions, and reporting standards to improve trust in business reporting.
- Automated regulatory and operational reporting workflows using Python and SQL to generate standardized utilization, claims, compliance reports, ensuring consistent delivery of validated outputs for internal monitoring and external regulatory submission requirements.
- Maintained version-controlled repositories containing SQL scripts, Python notebooks, data transformation logic, and reporting artifacts, supporting code consistency, analytical reproducibility, and collaborative development practices across analytics initiatives.
- Automated recurring operational and regulatory reports using SQL, Power BI, and Python, reducing manual effort and improving reporting efficiency.
- Collaborated with investment operations, portfolio management, reporting teams to translate analytical requirements into structured data models, supporting the delivery of investment analytics, operational reporting, and portfolio performance measurement solutions.
- Investigated and resolved ETL performance bottlenecks by analyzing Synapse SQL execution plans and optimizing Python-based transformation logic, improving processing efficiency for large-scale data workflows and reducing latency in reporting pipelines.
- Documented data models, pipeline workflows, and analytical specifications to support maintainability of Microsoft Fabric-based DataOps processes and facilitate onboarding, knowledge transfer, and consistent usage of healthcare data environments.
- Facilitated Agile ceremonies including backlog refinement, sprint planning, user story reviews, and stakeholder demonstrations.

Environment: Python (Polars, Pandas), SQL, Microsoft Fabric, Azure Synapse Lakehouse, Power BI, Looker, Azure Monitor, Azure Purview, Jupyter Notebooks, Excel, Azure Cost Management

Senior Business Data Analyst | SEI Investments Company, Colorado, US

Sep 2023 - Jun 2025

Project Title: Investment Performance & Portfolio Analytics Platform

Project Overview: Supported an enterprise investment analytics platform that consolidated portfolio, transaction, benchmark, and market data into a centralized reporting environment. The platform enabled portfolio managers, investment operations teams, and executives to analyze portfolio performance, asset allocation, investment risk, benchmark comparisons, and client reporting metrics.

- Gathered and analyzed business requirements from portfolio managers, investment operations teams, and reporting stakeholders.
- Developed SQL queries and analytical datasets to support portfolio performance reporting, transaction reconciliation, and investment analytics.
- Built Power BI dashboards to monitor portfolio returns, benchmark performance, asset allocation, risk indicators, and operational KPIs.
- Performed data analysis on portfolio holdings, transactions, and market data to identify trends and support investment decision-making.
- Collaborated with business users to define reporting metrics, KPI calculations, and dashboard requirements.
- Conducted data validation and reconciliation activities to ensure consistency between portfolio accounting systems and reporting platforms.
- Partnered with ETL and data engineering teams to support ingestion of investment data into Microsoft Fabric and Azure Data Lake.
- Documented business requirements, process flows, and reporting logic to support regulatory and operational reporting needs.
- Performed root cause analysis for data discrepancies and reporting issues, driving timely resolution.
- Supported migration of legacy reporting solutions to Microsoft Fabric while ensuring reporting continuity and user adoption.
- Worked within Agile Scrum teams to deliver analytics enhancements, dashboard improvements, and reporting solutions.
- Configured Azure Key Vault integrations to securely manage secrets, credentials, and connection strings used by analytical workloads, strengthening access controls and supporting secure interaction between cloud-based data platforms and reporting environments.

- Monitored data pipeline execution, dataset refresh activities, and platform health metrics using Azure Monitor, establishing automated alerting mechanisms to identify processing failures, data anomalies, and workflow interruptions affecting critical investment data feeds.
- Applied Microsoft Purview data classification and lineage capabilities to document sensitive financial data assets, improve traceability of data transformations, and support compliance with internal data governance and privacy management requirements.
- Validated data quality across source platforms, Lakehouse datasets, and reporting layers by implementing automated validation checks, reconciliation procedures, and exception reporting processes designed to identify discrepancies within financial and market data feeds.
- Supported migration activities associated with transitioning legacy reporting solutions to Microsoft Fabric, validating data integrity, testing reporting outputs, and ensuring continuity of analytical and operational reporting processes throughout cloud adoption efforts.
- Utilized Python to process and clean datasets, preparing them for performance attribution analysis & risk modeling in Jupyter Notebooks.
- Conducted RCAs for processing failures identified in logs, implementing fixes to prevent recurrence & improve the reliability of delivery.

Environment: Python (Pandas, Polars), SQL, Microsoft Fabric, Azure Data Lake Storage (ADLS Gen2), Synapse Lakehouse, Power BI, Looker, Jupyter Notebooks, Azure Monitor, Microsoft Purview, Git, Azure Key Vault

Business Data Analyst | Boston Consulting Group, Austin, US

Mar 2021 - Aug 2023

Project Overview: Worked with multiple client engagements to deliver customer analytics, marketing insights, operational reporting, and business intelligence solutions. Leveraged SQL, Python, AWS data platforms, and Power BI to transform large datasets into actionable business insights supporting strategic decision-making.

- Partnered with business stakeholders to understand analytical requirements and translate business questions into data-driven solutions.
- Performed exploratory data analysis using SQL and Python to identify customer behavior patterns, churn drivers, and business opportunities.
- Developed interactive Power BI and Tableau dashboards to track customer engagement, campaign performance, and operational KPIs.
- Analyzed large datasets from multiple business systems and transformed them into actionable insights for executive leadership.
- Designed and executed ad-hoc analyses to support business planning, customer segmentation, and performance optimization initiatives.
- Collaborated with data engineering teams to prepare analytical datasets from AWS Glue, Snowflake, and S3 environments.
- Automated recurring reporting processes using Python and AWS services, reducing manual reporting effort.
- Conducted data quality assessments and validation checks to improve confidence in business reporting.
- Presented analytical findings and recommendations to business stakeholders and client leadership teams.
- Participated in Agile project delivery, requirement workshops, sprint planning, and stakeholder review sessions.
- Developed predictive ML models within SageMaker to forecast customer churn & optimize marketing strategies for diverse engagements.
- Processed large-scale raw datasets using Glue and Python, transforming structured and unstructured data for analysis in S3 data lakes.
- Designed and executed SQL queries via Athena to perform ad-hoc analysis on data stored in data lakes, supporting rapid decision-making.
- Built visualizations in Tableau and Power BI to present analytical findings to stakeholders, translating data into actionable business insights.
- Managed extraction & loading processes from Snowflake warehouses into analytical environments, ensuring availability for ongoing projects.
- Automated recurring data processing tasks using Lambda functions, reducing manual intervention and ensuring timely updates of KPIs.
- Collaborated with case teams to define requirements and design approaches, ensuring alignment with objectives and client expectations.
- Monitored AWS resource usage and costs using CloudWatch, optimizing resource allocation for data processing jobs to maintain efficiency.
- Utilized GitHub for version control of analysis scripts and model code, maintaining a history of changes and facilitating team collaboration.
- Deployed ML models to production endpoints, enabling clients to integrate predictive capabilities into their operational workflows.

Environment: SQL, Python (Pandas, NumPy, scikit-learn, XGBoost, SHAP), Jupyter Notebooks, AWS (Athena, Glue, S3, Lambda, CloudWatch, SageMaker), Power BI, Tableau, Snowflake, Advanced Excel, Git, GitHub

Business Data Analyst – Operational Analytics | NIIT Technologies, Pune, India

Sep 2018 - Dec 2020

- Extracted operational data from Redshift using SQL to generate recurring performance reports for stakeholders and operational monitoring.
- Developed Tableau dashboards to visualize KPIs & service-level metrics, supporting tracking of performance trends & reporting requirements.
- Processed and cleaned raw datasets using Python to structure data for analysis and reporting workflows in Jupyter Notebook environments.
- Performed ad-hoc analysis using AWS Athena on data stored in S3, supporting business queries and reporting requirements across projects.
- Validated consistency across multiple systems, identifying discrepancies and supporting resolution before loading into reporting warehouses.
- Automated Excel reporting workflows using Python scripts, reducing manual effort in recurring reporting and improving report efficiency.
- Supported configuration and execution of workflows using Glue for movement of data between on-premise systems and cloud environments.
- Analyzed operational data using SQL and Excel to identify trends, support performance review cycles, assist process improvement activities.
- Assisted in migration of legacy reporting dashboards to QuickSight, ensuring continuity of reporting and validation of output accuracy.
- Collaborated with business users to gather reporting requirements and translate them into structured specifications for reporting solutions.

- Maintained documentation for data models, reporting logic, transformation rules to support knowledge sharing and operational continuity.
- Monitored data pipeline executions and resolved data quality and load-related issues to ensure stability of reporting workflows.
- Pre-processed raw data from multiple client sources, addressing missing values and outliers to ensure high-quality inputs for analysis.

Environment: SQL, Python (Pandas, NumPy), Jupyter Notebooks, Redshift, Athena, S3, QuickSight, Tableau, Excel, Git, GitHub

Associate Business Data Analyst | Hexaware Technologies, Bengaluru, India

Aug 2016 - Aug 2018

- Extracted, aggregated structured client, datasets using SQL, delivering accurate insights for reporting, operational decision-making.
- Built, maintained interactive Power BI, Tableau dashboards for client onboarding, SLA monitoring, project performance metrics.
- Performed data cleansing, normalization, validation using Excel functions to ensure consistency and accuracy across reports.
- Defined and standardized KPIs and reporting metrics for IT project tracking, SLA adherence, client performance evaluation.
- Applied descriptive statistical methods in Excel to identify outliers, detect process bottlenecks, improve service delivery efficiency.
- Documented workflows, system dependencies using BPMN, enabling audit-ready process visualizations, operational transparency.
- Developed SQL scripts for recurring operational, financial data extracts, reducing reporting effort, supporting monthly audits.
- Consolidated report definitions, transformation logic, calculation methodologies in Confluence for cross-team traceability, KT.
- Managed historical report data and version-controlled outputs in AWS S3, ensuring secure storage and retrieval for audits.
- Tracked report requests, enhancements, and defect resolutions using Jira, maintaining SLA compliance for deliverables.
- Partnered with business stakeholders to translate operational, financial gaps into actionable insights via visual analytics, trend analysis.

Environment: SQL, Excel, Power BI, Tableau, Data Cleansing & Validation, KPI Development, Data Visualization, Descriptive Statistics, BPMN, Jira, Confluence, AWS S3

EDUCATION DETAILS:

Bachelor of Technology

Aug 2012 - May 2016

Electronics and Communication Engineering | Anurag Group of Institutions, Hyderabad, India