

Santhosh G

Gsanthosh8737@gmail.com

984-262-8737

Data Engineer

PROFESSIONAL SUMMARY:

- Over 5 years of experience designing and developing database solutions and Business Intelligence applications using Microsoft technologies such as Power BI, Excel, SSIS, and SQL Server.
- Proficient in automating data workflows with Python and Apache Airflow and creating ETL processes using SSIS, Azure Data Factory, and AWS services like Glue, EMR, and Kinesis.
- Extensive experience designing and optimizing Snowflake data warehouses, including handling semi-structured data, continuous ingestion with Snowpipe, and leveraging advanced analytics capabilities.
- Hands-on expertise with cloud ecosystems, integrating diverse data sources into Azure and AWS platforms, and utilizing services like Redshift, S3, and Athena for real-time analytics.
- Skilled in creating efficient data models and schemas, managing Snowflake databases, and optimizing query performance using clustering keys and materialized views.
- Experienced in implementing Snowflake security features such as role-based access control, data encryption, and secure data sharing across accounts.
- Proficient in integrating Snowflake with BI tools like Power BI to deliver optimized reporting and actionable data insights.

TECHNICAL SKILLS:

BI Tools	Power BI, Power View, Power BI Desktop, Power Pivot, SQL, Tableau, SSRS
Databases	SQL Server, Azure SQL Database, TSQL, MySQL, Oracle, Redshift, Snowflake.
Programming Languages	SQL, T-SQL, DAX, Python, Java, PL/SQL, Spark
Integration Tools	SSIS
AZURE and AWS	Azure Data Factory (ADF) V2, Azure SQL, Azure Data factory, Azure Data Lake, Lambda, DynamoDB, S3, EC2, Redshift, VPC, Step functions, Glue, SNS, CloudWatch
Operating Systems	Windows, LINUX, Unix
Messaging tools	Kafka, EMQ, RabbitMQ
Bigdata Technologies	Hadoop, HDFS, Hive, Pig, Spark, Machine Learning, Pandas, NumPy, Informatica, Snowflake, Snow SQL, Data bricks, Kafka, Cloudera

EDUCATION:

❖ Master's in Computer Science, Wright State University.

CERTIFICATIONS:

✓ AWS Cloud Practitioner

✓ AWS Solutions Architect - Associate

PROFESSIONAL EXPERIENCE:

1. Client : Omega HealthCare Role : Data Engineer

Jun 2024 – Present

- Build and maintain scalable pipelines for processing large healthcare datasets, ensuring accuracy, compliance, and integration from diverse sources.
- Manage and optimize databases, data warehouses, and data lakes using big data tools like Hadoop and Spark.
- Deploy and optimize cloud-based data infrastructure on AWS, Azure, or GCP with low-latency, fault-tolerant solutions.
- Collaborate with cross-functional teams to design robust data models, support analytics with tools like Power BI and Tableau, and enable machine learning pipelines.
- Leverage healthcare domain expertise, including standards like ICD, HL7, and FHIR, to ensure data quality, validation, and regulatory compliance.
- Skilled in creating and maintaining complex data models and schemas through Python scripting.
- Apply technical skills in Python, SQL, and Spark while adhering to best practices in CI/CD, testing, and emerging data engineering trends.

Environment : *Python, SQL, Scala, Java, Apache NiFi, Talend, Informatica, AWS Glue, Apache Spark, Hadoop, Hive, Pig, PostgreSQL, MySQL, MongoDB, Cassandra, AWS (Redshift, S3, Lambda), Azure (Data Factory, Synapse), GCP (BigQuery, Dataflow), Star Schema, Snowflake Schema, Dimensional Modeling, Tableau, Power BI, QlikView, Jenkins, GitLab, Docker, Kubernetes, Grafana, Datadog, AWS CloudWatch, ICD, CPT, HL7, FHIR, HIPAA, GDPR, EMRs, EHRs, Claims Data, cross-functional teamwork, and agile methodologies.*

2. Client : Revature (Remote) Role : Data Engineer

April 2023 – Jun 2024

- Implemented real-time data ingestion pipelines with AWS Lambda and Amazon Kinesis, processing over 500GB of data weekly, and automated predictive analytics workflows by integrating AWS SageMaker with data pipelines.
- Designed scalable data pipelines with AWS Glue, GCP Dataflow, and Azure Data Factory, reducing processing times by 25% and enabling efficient batch processing through serverless ETL workflows.
- Leveraged Delta Lake's ACID transactions and Apache Spark's in-memory computing for consistent and high-performance data processing and analytics.
- Deployed cost-effective data lake solutions on Amazon S3 with serverless querying using Amazon Athena, reducing query times and enabling analysis of petabyte-scale datasets.
- Built interactive dashboards with Amazon QuickSight to deliver actionable insights from datasets exceeding 5TB, driving business decisions.
- Enhanced cloud resource security with AWS IAM best practices, implemented centralized logging with CloudWatch and AWS Lambda, and managed container orchestration using Kubernetes.
- Utilized Hadoop, Spark (Python and Scala), and multi-cloud platforms (AWS, GCP, Azure) to manage big data workflows and develop scalable, high-performance data engineering solutions.

Environment : *AWS Lambda, Amazon Kinesis, AWS Glue, GCP Dataflow, Azure Data Factory, AWS SageMaker, Delta Lake, Apache Spark (Python, Scala), Amazon S3, Amazon Athena, Amazon QuickSight, Kubernetes, IAM roles and policies, Amazon CloudWatch, Hadoop ecosystem, serverless ETL workflows, machine learning workflows, big data analysis.*

3. Client : Seneca Global (Remote)
Role : Data Engineer

Jun 2019 – July 2021

- Architected and implemented end-to-end data warehousing solutions on Azure, integrating diverse data sources and enabling seamless analysis and decision-making.
- Designed and optimized ETL pipelines with Azure Data Factory, improving data processing efficiency by 30%, and utilized Azure Data Lake Storage Gen2 for secure, scalable, and cost-effective data storage.
- Leveraged Apache Flink and Kafka for real-time stream processing, enabling low-latency analytics, stateful computations, and timely insights from high-velocity data streams.
- Established robust data governance frameworks and implemented data security measures, including Azure Key Vault and compliance with regulatory standards, to maintain data quality, integrity, and security.
- Optimized data ingestion and processing pipelines with technologies like Hadoop MapReduce and Snowflake, improving query performance and reducing execution times.
- Collaborated with stakeholders to translate business requirements into scalable data models and advanced analytics solutions, achieving a 25% improvement in query performance and actionable insights.
- Developed interactive dashboards and reports using Power BI and enabled real-time data processing with Azure Stream Analytics, supporting data-driven decision-making across the organization.

Environment : *Azure, Snowflake, Apache Kafka, Azure Data Lake Storage Gen2, Hadoop MapReduce, Apache Flink, Azure Stream Analytics, Azure Data Factory, Power BI, Data Governance, Data Security, ETL Pipeline Optimization, Real-Time Data Processing, Data Modeling, Query Performance Optimization.*