Subhash Chandra Gannamraju

+14699243965 | subhashgannamraju99@gmail.com | Github | LinkedIn

PROFESSIONAL EXPERIENCE

Data Engineer- Eight Eleven Group LLC (Brooksource Inc) | Dallas | TX | USA **November 2025 - Present**
Client: Cigna/(Allegiance)

- Designed and implemented end-to-end ETL tasks for the conversion of data from SQL Server to Snowflake using PySpark, Python, and SQL with the intention of aiding large-scale analysis in the health domain.
- Utilized the Qlik Replicate feature of data replication and change data capture (CDC) capabilities to enable the consumption of data from the source in real time in the Snowflake environment.
- Developed and automated the workflows on the raw, intermediate, and reporting levels of Snowflake to facilitate data preparation for analytics and reporting.
- UNIX shell scripts developed for work automation and file management using Linux servers through the use of VS Code and PuTTY.
- Highly efficient pipelines with Apache Airflow, integrated with the Azure DevOps CI/CD pipeline environment for managing version control, code reviews, and deployment processes.
- Implemented and maintained business intelligence dashboards and Enterprise reports on Power BI & Qlik Sense, with a design based on optimized data models on Snowflake to meet business requirements.

Data Analyst- Talent Value Inc. (CompEdge) | Boulder | CO | USA **May 2025 - October 2025**

- Engineered scalable ETL pipelines combining web scraping (Python, Selenium, BeautifulSoup) with PySpark and SQL, automating ingestion of compensation datasets and expanding coverage by 40%.
- Designed data cleaning, curation, and anomaly detection workflows to transform raw scraped data into structured, ML-ready inputs, reducing prep time by 30%.
- Automated ingestion into cloud platforms (GCP BigQuery, Dataflow, Cloud Storage; Azure Data Factory for workflow integration), ensuring scalable and reliable data access.
- Developed dimensional/star schema models and curated SQL data marts, improving analytic precision by 20% and reducing reporting latency by 25%.
- Implemented data quality validation, RBAC-based security, and encryption protocols, achieving 98% accuracy and ensuring compliance across enterprise datasets.

Materials Data Engineering Coop- Entegris | Hillsboro | OR | USA **July 2024 - December 2024**

- Designed and optimized ETL pipelines using SQL Server, Snowflake, SAP, and PySpark, improving processing efficiency across financial and manufacturing datasets.
- Enhanced data modeling and schema design for enterprise reporting, enabling faster KPI tracking and improving manufacturing analytics.
- Built end-to-end data pipelines from Snowflake and SQL Server into Power BI dashboards, fully automating reporting and improving procurement accuracy by 30%.
- Applied data integration best practices to unify LIMS, ETQ Reliance, SAP, and other enterprise data sources into a single cloud-based reporting environment.
- Conducted predictive analytics and statistical testing (ANOVA, T-tests, control charts) on ETL-driven datasets to improve product quality and reduce out-of-control counts by 15%.
- Automated workflows with Power Automate to optimize supplier data handling, cutting manual workload by 18 hours/week.

Data Analyst- The University of Texas at Dallas | Dallas | TX | USA **January 2024 - July 2024**

- Engineered end-to-end ETL pipelines to extract demographic, healthcare, crime, education, housing, and traffic datasets for 4 Texas counties (Dallas, Denton, Collin, Anderson).
- Automated data collection using Python (Selenium, BeautifulSoup) and R web scraping frameworks, achieving 95% data completeness and 98% accuracy in community indicators.
- Designed geospatial ETL workflows with R and GeoAnalytics, integrating census and traffic data with location-based mapping for enhanced regional insights.
- Structured raw data in PostgreSQL with optimized schema design, reducing query latency by 25% for downstream reporting.
- Developed data cleaning and transformation pipelines (Python + SQL) to standardize heterogeneous public datasets, cutting preprocessing time by 30%.
- Built interactive dashboards in Power BI, R Shiny, and ArcGIS, increasing community engagement by 5% within 3 months and providing actionable insights on lifestyle trends.

Business Intelligence Developer- Diagonal Consulting LLP | India **December 2019 - December 2022**

- Designed and implemented full end-to-end ETL solutions to integrate data from multiple sources, such as SAP, SQL Server, Excel, and streaming data, to load data warehouses and business intelligence solutions for near-real-time analytics.
- Developed and optimized dimensional models (star schema, snowflake schema) for the health, banking, energy, and public sectors to measure different KPIs, thus speeding up reporting by 30-35%.
- Implemented expert-level SQL logic with a focus on stored procedures, window functions, indexes, and partitioning to meet business and regulation-related analytics requirements.
- Built incremental loading mechanisms, metadata driven ETL solutions, and automation engines by leveraging Python, PySpark, Azure Data Factory, AWS Glue, and Spark Streaming that translated to a 40% reduction in reporting cycles.
- Migrated large-scale reporting jobs to Snowflake, Azure Synapse, and Amazon Redshift, resulting in improvements in scalability, performance, and reducing operational costs up to 20%.
- Provided executive and operational dashboards for execution in tools such as Power BI, QlikSense, and Tableau, which triggered a 20-30% increase in adoption based on drill-down analytics and self-service reports.
- Established data validation, reconciliation, and governance processes with an accuracy rate of 95% to 98% in healthcare outcomes, financial audits, and regulatory reporting.
- Designed and implemented security and monitoring solutions using Azure Monitor, Log Analytics, and Key Vault to protect sensitive data related to healthcare, banking, and the government while maintaining a 99.9% system availability rate.
- Worked along with other cross-functional teams on Agile Delivery Models, implementing business requirements into scalable business intelligence solutions, and successfully delivered over 10 enterprise reports.

TECHNICAL SKILLS & CERTIFICATIONS

- **Programming Languages:** Python (Advanced), R, SQL, PySpark, Unix/Linux Shell Scripting, Scala and Java.
- **Cloud Data Platforms:** Snowflake (Expert), AWS (S3, Redshift, Glue, Lambda, EMR), Azure (Data Factory, Synapse).
- **Databases:** MySQL, PostgreSQL, AWS, Amazon Redshift, SAP, Microsoft SQL server, Mongo DB, Snowflake, Oracle.
- **Data Visualization:** Qlik Sense, Tableau, Power BI, Looker, Python (Matplotlib, Seaborn), Google Analytics, Alteryx.
- **Data Modeling:** Dimensional Modeling, Star/Snowflake Schemas, Data Vault, ER Diagrams.
- **Certification:** Received Applied Machine Learning Certification from University of Texas at Dallas in 2024.

PROJECT EXPERIENCE

Loan Approval and Rejection Prediction Machine Learning Model **April 2024**

- Built a K-Nearest Neighbor model achieving 81.4% accuracy, cutting processing time by 20% and improving decision efficiency by 25%.
- Integrated model insights into Power BI dashboards, improving loan approval accuracy by 30% through data driven visualization.

Hotel Booking Cancellation Prediction using R **December 2023**

- Designed a logistic regression model with 85.3% accuracy and 85.2% balanced accuracy, reducing cancellation risk by 14.7%.
- Identified key predictors such as special requests, lead time, and market segment to optimize booking strategies.

Online Shoppers Purchase Intentions Prediction using Python **November 2023**

- Successfully developed SVM model with Kernel functions and Decision Tree to predict online shoppers' intentions with an impressive accuracy of 89.7% by featuring data cleaning, feature engineering and model training.

Risk Analytics on Truck Fleet Data Python **August 2023**

- Performed linear regression analysis to quantify accident risks, finding a +2.5 accident increase per 100 trucks.
- Identified Ford trucks and lane departure as highest-risk contributors, guiding preventive safety measures.

Predicting Bank Failures **May 2023**

- Built a Random Forest model achieving 96.3% accuracy and 96.1% balanced accuracy for bank failure prediction.
- Highlighted capital adequacy, asset quality, and earnings as key predictors, enabling proactive risk management.

EDUCATION

Master of Science, Business Analytics & Artificial Intelligence **January 2023 - December 2024**
The University of Texas at Dallas | USA **GPA: 3.6/4.0**
Relevant Coursework: Database foundations for Business Analytics, Big Data Analytics, Data Visualization, AWS Solution Architecture, Cloud Computing & Applied Machine Learning.

Bachelor of Technology, Computer Science & Engineering **July 2016 - September 2020**
Jawaharlal Nehru Technological University Kakinada | India **GPA: 3.3/4.0**
Relevant Coursework: Data Structures, Database Management System, Software Engineering, Data Mining, Big Data Analytics, Fundamentals of Machine Learning.