

Supriya Tumuganti

United States | ✉ tsupriya0108@gmail.com | ☎ 945-713-4594

www.linkedin.com/in/supriya-tumuganti/

PROFESSIONAL SUMMARY

Data Engineer with 4 years of experience building and optimizing scalable batch data pipelines using Databricks, PySpark, and SQL. Strong experience in cloud data pipelines, data modeling, and data warehousing with Snowflake. Proven ability to design ETL workflows, optimize Spark jobs, and support analytics and BI use cases in enterprise environments.

TECHNICAL SKILLS

- **Languages:** SQL, Python
 - **Big Data Processing:** Apache Spark (PySpark), Databricks
 - **Data Warehousing:** Snowflake
 - **Cloud Platforms:** AWS (S3, EC2), Azure (Azure SQL, Azure Data Factory)
 - **ETL & Orchestration:** Databricks Jobs, Talend, SSIS
 - **Databases:** SQL server, DB2, Oracle
 - **BI & Analytics:** Power BI
 - **Tools:** VS Code, Git, JIRA, CI/CD
-

EDUCATION

Master of Science in Computer and Information Science

Southern Arkansas University – Magnolia, AR

August 2024 - December 2025

Bachelor of Technology in Computer Science and Engineering

Balaji Institute of Technology and Science - India.

August 2015-May 2019

PROFESSIONAL EXPERIENCE

Cognizant Technology Solutions – Hyderabad, India
Software Engineer (September 2019 – November 2023)

Databricks-Based Spark ETL Platform

- Designed and built Spark-based batch ETL pipelines in Databricks using PySpark, processing large number of records per day from structured and semi-structured data sources.
- Enhanced Databricks notebooks and Spark transformations, increasing data processing capacity to 1M+ records per hour and enabling faster business insights.
- Leveraged Databricks Lakehouse architecture, integrating cloud object storage with Spark processing to support unified analytics and downstream consumption.
- Implemented data quality checks and validation logic within ETL pipelines using Python (pandas, NumPy), improving data accuracy by ~15% and enhancing trust in analytical outputs.
- Optimized Spark jobs using partitioning strategies, caching, and SQL tuning, reducing pipeline execution time by approximately 35% and improving cluster resource utilization.
- Co-ordinated end-to-end batch data pipelines with monitoring and failure handling, improving pipeline reliability for downstream analytics workloads.
- Collaborated with 2+ cross-functional teams to analyze performance bottlenecks and apply algorithmic optimizations for large-scale data processing.

Cloud Data Warehouse & Analytics Integration

- Developed and maintained Snowflake data warehouse models (fact and dimension tables) to enable scalable analytics and Power BI reporting across business teams.
- Applied data governance principles by implementing column-level data masking policies in Snowflake to enforce GDPR compliance and secure sensitive data.
- Built efficient data transformations and tuned SQL queries in Snowflake to enable fast analytical queries on large datasets.
- Spearheaded data quality initiative employing Python (pandas, NumPy), identifying and fixing 12 data integrity issues, and improving overall accuracy in data pipelines by 15%.
- Developed and facilitated Power BI dashboards using refined data models to provide actionable insights for executive stakeholders and business decision-making.
- Minimized data errors in reporting by orchestrating batch workflows using Databricks Jobs, which ensured 100% data accuracy on key reports used by executive stakeholders.