Tejani Shree

New Jersey | +1 (857)-230-5776 | tejanishree6@gmail.com | LinkedIn

SUMMARY

AI/ML Engineer with 3+ years of experience across data analytics, machine learning, and generative AI projects in India and the United States. Hands-on experience with Python, SQL, PyTorch, TensorFlow, scikit-learn, XGBoost, RAG, LangChain, LangGraph, Hugging Face, and AI agents. Experienced in developing, evaluating, deploying, and monitoring machine learning and LLM applications using FastAPI, MLflow, Docker, Kubernetes, AWS, Redis, Grafana, DeepEval, and Guardrails AI. Background includes financial analytics, demand forecasting, network troubleshooting, model evaluation, and production AI workflows.

TECHNICAL SKILLS

Programming & Data Science: Python, SQL, NumPy, pandas, SciPy, Jupyter, PySpark, Apache Spark, Exploratory Data Analysis, Statistical Analysis, Data Visualization

Machine Learning & Deep Learning: scikit-learn, XGBoost, LightGBM, PyTorch, TensorFlow, Supervised Learning, Unsupervised Learning, Regression, Classification, Clustering, Feature Engineering, Hyperparameter Tuning, Model Evaluation

Generative AI & LLM Engineering: Generative AI, Large Language Models (LLMs), Transformers, Hugging Face, Retrieval-Augmented Generation (RAG), Prompt Engineering, Embeddings, Semantic Search, Vector Databases, Fine-Tuning, PEFT, LoRA, NLP

Agentic AI & LLM Applications: AI Agents, Agentic AI, LangChain, LangGraph, LlamaIndex, Tool Calling, Function Calling, Multi-Agent Systems, LLM Evaluation, RAG Evaluation, DeepEval, Guardrails AI

MLOps, Deployment & Monitoring: MLOps, MLflow, Experiment Tracking, Model Registry, Model Serving, Model Monitoring, Model Drift, FastAPI, REST APIs, Docker, Kubernetes, Git, GitHub Actions, Jenkins, CI/CD, Grafana

Cloud, Databases & Data Engineering: AWS, Amazon SageMaker, Amazon Bedrock, Azure Machine Learning, Azure OpenAI, Google Vertex AI, Databricks, Kafka, ETL Pipelines, Data Pipelines, Snowflake, PostgreSQL, MySQL, Redis, AWS S3

EXPERIENCE

Cisco Systems | USA **May 2025 – Present**

AI/ML Engineer

- Designed LLM-powered RAG workflows using Python, LangChain, LLM APIs, vector embeddings, semantic search, and Redis vector storage to ground Cisco network-support responses in enterprise knowledge sources, improving answer relevance by 21% across internal evaluation datasets.
- Engineered agentic AI workflows with LangGraph, LLM tool/function calling, prompt engineering, and FastAPI to automate multi-step network diagnostics, reducing manual investigation time by 24% for operations teams.
- Integrated Hugging Face Transformer and NLP models with Cisco telemetry and unstructured knowledge sources, applying PyTorch and contextual embeddings to support semantic understanding and faster root-cause analysis across complex enterprise network incidents.
- Evaluated LLM and RAG applications using DeepEval, retrieval and generation metrics, and Guardrails AI to measure faithfulness, relevance, latency, and safety, increasing evaluation pass rates by 17%.
- Developed and evaluated deep learning and NLP models using PyTorch and TensorFlow, applying model fine-tuning, hyperparameter optimization, and validation techniques to improve task-specific model performance.
- Integrated LLM APIs into FastAPI-based AI services to support production inference, prompt orchestration, structured responses, and downstream agent workflows.
- Monitored production LLM/AI services using MLflow and Grafana, tracking inference latency, model behavior, retrieval quality, and application performance to identify issues before affecting network operations.
- Deployed and maintained production AI/ML applications using Docker, Kubernetes, GitHub Actions, and AWS, standardizing CI/CD workflows for scalable model inference and reliable cloud-native releases.

Intex Technologies | India **Aug 2022 – Feb 2023**

Machine Learning Engineer

- Built demand forecasting models using Python, pandas, scikit-learn, and XGBoost to predict Smart TV sales across regional product categories, improving forecast accuracy by 16%.
- Developed and trained supervised machine learning classification models using sales, inventory, seasonal, and product features to identify slow-moving electronics, helping planning teams reduce excess stock levels by 12%.
- Engineered predictive features and optimized model performance through feature selection, hyperparameter tuning, cross-validation, and model evaluation, improving validation performance by 10% and supporting more reliable production planning.
- Prepared and transformed historical sales and inventory data using SQL, Python, and pandas, automating recurring data pipelines and improving the consistency of machine learning training datasets.
- Compared model performance across forecasting and classification approaches using appropriate validation metrics and error analysis, selecting models that provided more reliable predictions for business planning workflows.
- Deployed ML prediction services using FastAPI and Docker, exposing model inference through APIs so analysts and planners could access forecasts consistently for inventory and replenishment decisions.

Mphasis | India **Mar 2021 - Jul 2022**

Data Analyst

- Analyzed customer and transaction data using Python, SQL, pandas, and Jupyter to support 360-degree risk profiling, improving reporting accuracy by 18% for banking teams.
- Developed data validation checks and exploratory analysis workflows across multiple source systems to identify quality issues, reducing manual reconciliation effort by 22% for analysts.
- Evaluated customer attributes through statistical analysis and feature selection to strengthen risk segmentation, helping business teams identify higher-value credit opportunities and improve decision quality.
- Implemented ETL and data preparation routines for structured financial datasets, standardizing downstream reporting inputs and cutting recurring data preparation time by 25% across monthly cycles.
- Created Tableau dashboards and SQL-based reports for customer risk, product trends, and data quality metrics, giving stakeholders clearer insights for faster portfolio decisions.

EDUCATION

Master of Science in Data Analytics **Apr 2025**
Northeastern University, MS

Bachelor of Technology in Software Development **Apr 2022**
Jai Hind College, University of Mumbai, India

PROJECTS

Credit Card Data Analysis **Sept 2023 – Dec 2023**
Tech: PySpark, Apache Spark, Pandas, AWS S3, MySQL, Tableau, SQL

- Constructed PySpark ETL pipelines to ingest and transform 6,146 credit-card records, reducing batch processing time by 30% and enabling faster downstream fraud and reporting workflows.
- Derived predictive features (credit-utilization, delinquency flags, rolling averages) using Pandas and Spark to improve model inputs, boosting fraud classifier accuracy by 10% and reducing false alerts.
- Developed interactive Tableau dashboards for three senior stakeholders to monitor behavior patterns and segment growth opportunities, cutting manual analysis time and informing targeted retention strategies.

Kickstarter Success Prediction **Apr 2023 – May 2023**
Tech: Python, scikit-learn, XGBoost, Pandas, feature engineering, Tableau

- Trained ensemble classification models using scikit-learn and XGBoost on engineered features to predict campaign outcomes, achieving 99.5% validation accuracy and informing funding decisions.
- Formulated hypothesis-driven feature engineering (backer momentum, reward tiers, launch timing) that raised model AUC and improved performance by 12%, guiding resource allocation.
- Visualized model insights with narrative dashboards and charts to present actionable recommendations, improving project selection efficiency by 15% for product and investment teams.

CERTIFICATIONS

- Data Visualization and Storytelling Basics