Name: Veera Venkata
Role: Data Engineer

Contact Info:
veeravenkata0143@gmail.com

+1(603)5122416

PROFESSIONAL SUMMARY:

- **5+ years of experience** across the Big Data lifecycle, with a focus on **analytics, design, development, implementation, testing, and deployment**. **Expert in building scalable enterprise ETL (Extract, Transform, Load) and data warehouse solutions** for robust data management.
- **AWS:** Adept at utilizing EC2, S3, EMR, Athena, Redshift, Glue, Step Functions, Lambda, Aurora, CloudWatch, SNS, ECS, Fargate, and CloudFormation to design and implement data solutions.
- **Azure:** Hands-on experience with Azure SQL Server, Data Factory, Databricks, Synapse Analytics, ML Studio, SQL Data Warehouse, Cosmos DB, and DevOps for building data analytics solutions.
- **Hadoop distributions:** Building highly reliable and scalable big data solutions on platforms like Cloudera.
- **Data warehousing:** Designing and implementing data warehouses like Azure Synapse, Redshift, and Snowflake to meet diverse reporting and analytics needs.
- **Big data databases:** Knowledge of HBase, NoSQL databases (MongoDB, DynamoDB, CosmosDB), and optimizing data storage and query performance in cloud data warehouses.
- **Apache Spark:** Expertise in writing Spark Streaming API for real-time data processing on big data clusters.
- **Statistical analysis:** Deep understanding of statistical modeling, multivariate analysis, model testing, problem analysis, comparison, and validation.

TOOLS AND TECHNOLOGIES:

Hadoop/Big Data Technologies	Hadoop, MapReduce, HDFS, YARN, Oozie, Hive, Sqoop, Spark, Nifi, Zookeeper and Cloudera Manager, Horton Works.
Azure Cloud Services	Azure Data Factory, Azure Synapse Analytics, Data Lake, Blob Storage, HDInsight, Azure Databricks, Azure Data Analytics, Azure functions
AWS Cloud Services	EC2, EMR, Redshift, S3, Databricks, Athena, Glue, AWS Kinesis, Cloud Watch, SNS, ,lSQS, SES
NO SQL Database	HBase, Dynamo DB, Mongo DB
ETL/BI	PowerBI, Tableau, Snowflake, Informatica, SSRS, SSAS, QlikView, Qlik Sense.
Hadoop Distribution	Horton Works, Cloudera.
Programming & Scripting	Python, Scala, SQL, Shell Scripting, Kafka
Operating systems	Linux (Ubuntu, Centos, RedHat), Windows (XP/7/8/10).
Databases	Oracle, MY SQL, Teradata, PostgreSQL, SQL Server.
Machine Learning Libraries	Scikit-Learn, Pandas, NumPy, Seaborn, Matplotlib, Ggplot2, Beautiful Soup, NLTK, Spacy, Text Blob
Devops	Jenkins, Git, GitLab, Cloud Formation, A/B Testing
Version Control	Bitbucket, Git Lab, Git Hub, Azure DevOps

EDUCATIONAL QUALIFICATION:

- ✓ Master's in computer information systems from New England Collage, Henniker, NH- U.S.A

PROFESSIONAL EXPERIENCE:

Client: Varo Bank, San Francisco, CA **Aug 2022 – Present**
Role: Data Engineer
Roles & Responsibilities:

- **Optimized Performance & Scalability (Hadoop & Spark): Leveraged partitioning, bucketing, and advanced optimization techniques** within Hive tables to significantly improve data processing speed and resource utilization.
- **Possess in-depth experience managing HDFS, YARN, Spark, and MapReduce components on-premises and in the cloud.** Utilized Cloudera Manager and cloud-based platforms for efficient cluster configuration, resource monitoring, and diagnostics.
- **Big Data Ingestion, Processing & Analytics: Skilled at ingesting data from diverse sources** including relational databases, NoSQL stores, and dynamic files into HDFS and cloud storage (e.g., AWS S3) using efficient file formats (ORC, Parquet, Avro).
- **Developed production-ready ETL pipelines** using Apache PySpark to automate data extraction, transformation, and loading processes, leveraging Spark SQL and Data Frames APIs for efficiency and fault tolerance.
- **Performed in-depth analysis of Big Data clusters using various tools** like Pig, Hive, HBase, Spark, and Sqoop to extract valuable insights.
- **Developed HQL queries, mappings, tables, and external tables in Hive** for efficient data analysis across multiple domains. Optimized query execution through partitioning, compilation strategies, and cost-based optimization techniques.
- **Cloud Migration & Automation Expertise: Successfully migrated data from on-premises environments to cloud platforms like AWS EMR and Azure HDInsight** using scripting languages (e.g., Shell) for automation and efficient resource utilization.
- **Invoked Python scripts to perform data transformations on massive datasets** within cloud-based streaming platforms like AWS Kinesis or Azure Stream Analytics.
- **Automated data movement and processing workflows** using Apache NiFi or cloud-native data orchestration services, streamlining data pipelines across on-premises and cloud environments.
- **Data Integration & Extraction Specialist: Adept at using Sqoop and cloud-native data integration tools** to load data from multiple relational databases (SQL, DB2, Oracle) and NoSQL stores into HDFS and cloud storage for further processing and analysis within Hive tables or cloud data warehouses.
- **Successfully migrated data from diverse data sources like Teradata into HDFS and cloud storage** using Sqoop or cloud-native connectors, bridging the gap between disparate data sources for holistic analysis.

Environment: ADLS, ADF, Azure Synapse Analytics, Azure Databricks, Spark, Hive, Azure Monitor, Azure Functions, DMS, Matplotlib, Power BI, Azure Power BI, Kafka, Collibra, Azure Event Hubs, Azure Cosmos DB, ARM templates.

Client: American Equity, West Des Moines, IA **Sep 2019 – July 2022**
Role: Data Engineer
Roles & Responsibilities:

- **Architected and Delivered Scalable Data Solutions:** Designed and implemented data lakes and data warehouses on both on-premises and cloud platforms (AWS, Azure) using a variety of technologies like HDFS, S3, Redshift, Snowflake, Glue, and Data Lake Storage.
- Developed and optimized data pipelines for efficient data ingestion, transformation, and loading (ETL) using tools like Sqoop, AWS Glue, Azure Data Factory (ADF), and PySpark.
- Possess expertise in various data storage solutions including relational databases (RDS, SQL Server), NoSQL databases (DynamoDB, Cosmos DB), and data warehouses (Redshift, Snowflake) to ensure efficient data storage and retrieval based on project requirements.
- Experience in designing and implementing streaming data pipelines using technologies like Apache Kafka, Kinesis, and Spark Streaming for real-time data processing and analytics.
- **Advanced Analytics & Real-Time Processing:** Utilized various big data processing frameworks like Apache Spark, Spark SQL, and Hive for complex data analysis and manipulation.
- Implemented machine learning algorithms and data mining techniques to extract valuable insights from large datasets.
- Developed data visualization dashboards and reports using tools like Power BI, Tableau, and Matplotlib for effective data communication.
- Experience in cloud-based analytics platforms like AWS Quick Sight and Azure Synapse Analytics for interactive data exploration and visualization.

- **Data Governance & Automation:** Implemented data governance frameworks (Collibra) to ensure consistent data definitions, access control, and data quality throughout the data lifecycle.
- Utilized automation tools like Apache NiFi, AWS Lambda, and Azure Functions to automate data processing tasks and workflows, improving efficiency and scalability.
- Experience with Infrastructure as Code (IaC) tools like ARM templates for automating infrastructure provisioning and deployment in the cloud.
- Employed CloudWatch and similar monitoring tools to ensure operational visibility, proactive management, and performance optimization of data pipelines and infrastructure.

Environment: Azure Data Factory, Databricks, Azure HDInsight, Azure Events Hub, Azure SQL Databases, Azure Cubes, Azure Analysis Services, Cosmos DB, Scala, Python, PySpark, JSON, Jenkins, Git, Hive, T-SQL, Spark SQL, U-SQL, HIPAA compliance, Azure Data Lake.

Client: M&T Bank, Buffalo, NY **Jun 2018 – Aug 2019**
Role: Data Engineer
Roles & Responsibilities:

- **Architected and Delivered Scalable Data Solutions: Designed and implemented data pipelines using Azure Data Factory (ADF) for data ingestion, transformation, and loading (ETL) into Azure Data Lake Storage (ADLS).** Optimized data flows for efficient loading and minimized costs.
- **Developed data warehousing solutions using Azure Synapse Analytics for data analysis and reporting.** Utilized PolyBase for efficient data access and Azure Data Catalog for robust metadata management.
- **Migrated tables to production using Azure Synapse Analytics and leveraged Azure Databricks clusters for advanced analytics and insights.** Implemented real-time data pipelines using Databricks and Delta Lake for Change Data Capture (CDC).
- **Advanced Analytics & Real-Time Processing: Optimized Snowflake SQL queries and table structures for cost reduction.** Designed and implemented a large-scale data warehouse on Snowflake for a major retail client, enabling efficient storage and analysis of terabytes of customer data. Utilized Snow Pipe, Snow SQL, and Tasks for streamlined data ingestion, transformation, and automation.
- **Successfully migrated data pipelines to Snowflake, implementing optimized schema data models for efficient warehousing and analysis.** Achieved significant query execution time reduction through application optimization.
- **Built real-time streaming pipelines utilizing Azure Databricks, Azure Event Hubs, Kafka, and Spark Streaming for faster data ingestion and processing.** Enabled ad-hoc analysis with Azure Synapse Analytics and Delta Lake.
- **Data Governance & Quality Champion: Implemented a data quality framework utilizing Azure Data Factory for orchestration, Azure Functions for serverless logic, and Azure SQL Database for data storage.** Reduced duplicate records by 12% and increased data accuracy by 8% for a major retail client through automated data profiling, validation checks, and anomaly detection.
- **Led the integration of CI/CD pipelines utilizing Azure DevOps and version control with Git, adhering to Scrum and Agile methodologies.**
- **Effectively utilized Unity Catalog to establish a unified governance layer across multiple data lakes and warehouses, ensuring consistent access controls, metadata management, and data lineage tracking within the project.** Implemented Unity Catalog's lineage tracking capabilities for real-time monitoring and analysis across warehouses.

Environment: ADLS, ADF, Azure Synapse Analytics, Snowflake, Azure Databricks, Spark, Azure Event Hubs, Spark Streaming, Azure Data Factory, Azure Functions, Azure SQL Database, Azure Cosmos DB, Azure DevOps, Git, Azure Purview.