

Summary

Results-driven Data Scientist with 4+ years of experience in machine learning, ETL pipelines, and data analysis. Expert in Python, R, SQL, and cloud platforms like AWS and Azure. Specializes in developing predictive models using NLP, Generative AI, and LLMs. Proficient in optimizing data warehouses, enhancing data delivery, and leveraging cloud services for scalable solutions. Adept at data visualization to drive strategic growth, transform complex datasets into actionable insights, and improve decision-making and operational efficiency.

- Technical Skills
- **Programming Languages:** Python, R, SQL, NoSQL, HTML, CSS, C++, Java
 - **Frameworks & Tools:** Flask, MS SQL, MySQL, PostgreSQL, MongoDB, Git, GitHub, Docker, Statistics, AWS, Azure, Power BI, Tableau, Kubernetes, FlaskAPI, Eclipse, Google Colab, A/B Testing, PyTorch, TensorFlow, Hugging Face, OpenAI API
 - **Generative AI & Deep Learning:** Llama AI, Mistral AI, GPT-3, GPT-4, BERT, GANs, Variational Autoencoders (VAEs), Transformer Models, Siamese Networks, LSTM, RNN, CNN, FFNN, DALL-E, Stable Diffusion
 - **Data Processing & Streaming:** PySpark, Apache Kafka, Pandas, NumPy, Apache Airflow, Snowflake, ETL, MS Excel, Matplotlib, Seaborn, Plotly, scikit-learn, statsmodels
 - **Machine Learning Algorithms:** Random Forest, Linear Regression, SVM, Decision Tree, ARIMA, XGBoost, NLTK
 - **Statistical Analysis and Modeling:** ARIMA models, PCA, ANOVA, descriptive statistics, probability distributions, Regression models, LDA, TF-IDF, Mixed-effects models, Longitudinal data analysis, hypothesis testing, correlation analysis

- Professional Experience
- Data Scientist, Palantir Technologies

08/2023 - Present | Remote, USA

 - Developed an NL-to-SQL platform using a fine-tuned version of LLaMA 3.1 to translate natural language queries into SQL for easy database management and user-friendly data access.
 - Designed and implemented ETL pipelines with Apache Airflow and AWS Glue to automate the extraction, transformation, and loading of structured and unstructured financial data for SQL querying, ensuring data integrity and consistency.
 - Utilized AWS Lambda and Apache Kafka for real-time data streaming and low-latency execution of model inference functions, facilitating smooth and efficient SQL query generation and execution.
 - Deployed the fine-tuned LLaMA 3.1 model using AWS SageMaker and Docker for scalable model training and deployment, integrated with Kubernetes clusters for real-time SQL generation with improved latency and reliability.
 - Employed advanced prompt engineering techniques to optimize the fine-tuned LLaMA 3.1 (70B) model, implementing one-shot and two-shot learning strategies, resulting in a 40% increase in the accuracy of translating natural language queries into SQL.
 - Leveraged AWS Step Functions for orchestrating multi-step workflows involving NL-to-SQL translation, query execution, and automated response generation, enhancing overall system performance and user experience.

- Data Scientist, Infosys

01/2019 - 12/2021 | Hyderabad, India

 - Crafted AI-powered financial advisory platform integrating personalized investment suggestions by analyzing user profiles and market data using Python for data manipulation and model training, achieving a 30% performance boost.
 - Designed ETL pipelines using Apache Airflow and Azure Data Factory to automate the ETL process of financial data from sources like market feeds, user profiles, and news articles, for clean, accurate, and timely data preparation for analysis.
 - Employed BERT and GPT-4 models to analyze financial news and social media sentiment, generating actionable insights.
 - Leveraged Hugging Face Transformers for fine-tuning BERT to capture context and GPT-4 for delivering personalized investment advice.
 - Developed ML models for risk assessment and portfolio optimization, improving prediction accuracy by 25%, using XGBoost for classification and regression tasks. Performed hyperparameter tuning to refine model performance and robustness.
 - Deployed models using Azure Machine Learning and Docker for scalable training and deployment, managing large-scale data inputs and real-time predictions using Azure Kubernetes Service, improving latency by 25%.
 - Integrated Azure Functions and Azure Event Hubs for real-time data processing, allowing low-latency execution of model inference functions. used Azure Logic Apps for data workflows and automated response generation, enriching user experience.

- Education
- Master of Science, Harrisburg University

Information Systems Engineering and Management

2022 – 2024 | Harrisburg, USA
- Bachelor of Technology, MVSR Engineering College

Computer Science and engineering

2015 – 2019 | Hyderabad, India