Yash Bhaskarwar

Los Angeles Metro Area, CA (Open to Relocation) | (951)548-9583 | bhaskarwaryash12@gmail.com
GitHub | LinkedIn | Portfolio

Summary

Data Engineer with 3+ years of experience building and operating ETL pipelines for analytics and reporting systems. Specialized in transforming manual data workflows into reliable, automated pipelines and designing scalable data platforms that deliver accurate, timely data for business and research use cases.

Education

University of California, Riverside — **Sep 2023 – Mar 2025**
Master of Science in Computer Science — *Riverside, CA*

Technical Skills

Programming: Python, SQL, Bash
Data Engineering: ETL/ELT, Apache Airflow, PostgreSQL, Data Modeling, Data Pipelines
Data & Databases: PostgreSQL, Feature Engineering, ML Pipelines
Cloud Platforms: AWS (S3, Glue, Lambda, DynamoDB), Data Lake Architecture
DevOps & Tools: Docker, Jenkins, CI/CD, Git, GitHub Copilot

Experience

Sustainable Mobility & Environment Laboratory (SMEL), UCR — **Nov 2025 – Present**
Data Engineer — *Riverside, CA*

- Developed Python data pipeline architecture with data quality framework to ingest and standardize structured datasets, enabling reliable, repeatable analytics workflows
- Built Python data workflows for EMFAC vehicle emissions datasets, processing 30+ years of data, automating modeling pipelines and enabling reproducible scenario analysis

TATA Consultancy Services (TCS) — **Oct 2020 – Aug 2023**
Software Engineer - Data Platforms — *Remote*

- Built Python ETL data pipeline architecture supporting 10+ analytics workflows, integrating multi source healthcare and financial data and reducing data quality failures by 30%
- Optimized SQL based data models on PostgreSQL datasets with 5M+ records, reducing query latency by 20% and improving dashboard performance for business users
- Built AWS ETL pipelines with S3, Glue, Lambda, and DynamoDB to automate batch ingestion and transformation of structured and semi-structured data, reducing infrastructure costs by $300K annually
- Automated CI/CD pipelines using Jenkins and Git for data services and batch jobs, reducing deployment time from about 1 hour to under 15 minutes
- Added Python monitoring and logging to key data pipelines, detecting failures earlier and reducing repeated pipeline reruns during production support

Projects

StreamCart Realtime ETL System | *Python, Redpanda (Kafka API compatible), PostgreSQL, Docker* **GitHub**

- Built real time streaming ETL pipeline using Redpanda and PostgreSQL to ingest, validate, and enrich order events, handling up to 600 events per second for event level analytics
- Implemented idempotent writes and health checks for fault tolerance, containerizing services with Docker to ensure reliable and repeatable pipeline execution

Automated Transit Data Pipeline (GTFS Dataset) | *Python, PostgreSQL, Airflow, XGBoost* **GitHub**

- Built end to end batch ETL data pipeline processing 2M+ records, designing PostgreSQL star schema with incremental loads to improve query performance by 8x
- Orchestrated ingestion and daily predictions using Airflow DAGs, integrating XGBoost achieving 4–5 minute MAE and 89% tolerance accuracy, leveraging AI-assisted tool to prototype and validate data transformations