

Yashraj Bhandare +1(510)-396-6915 yashrajbajobguide@gmail.com CA (Open to Relocation) LinkedIn	
SUMMARY	Senior Data Engineer with 7+ years of experience designing scalable batch and real-time data platforms supporting analytics, experimentation, and machine learning workloads. Expert in Databricks, Spark, Snowflake, AWS, Python, SQL, and cloud-native ETL architectures processing millions of events daily. Strong background partnering with Product, Engineering, and Data Science teams to build trusted datasets, KPI frameworks, and governed data platforms that drive business decisions. Experienced in AI-powered data engineering including RAG, LLM applications, vector search, and intelligent data quality automation.
SKILLS	Programming & Query Languages: Python, SQL, Scala, R, NoSQL, Shell Scripting Big Data & Streaming: Apache Spark, PySpark, Apache Kafka, Apache Flink, Hadoop, Hive, HDFS, Databricks (Delta Lake, Unity Catalog), Impala, ZooKeeper Cloud Platforms: AWS (S3, Glue, EMR, Redshift, Lambda, Athena, Kinesis, EC2, IAM, CloudWatch, Elasticsearch), Microsoft Azure (Azure Data Factory, Data Lake, Synapse Analytics, Blob Storage, Stream Analytics), Google Cloud Platform (BigQuery, Pub/Sub) ETL, Data Integration & Orchestration: Apache Airflow, dbt, AWS Glue, Azure Data Factory, Apache NiFi, Talend, SSIS, Informatica, ETL/ELT Pipeline Development Data Warehousing & Databases: Snowflake, Databricks, BigQuery, PostgreSQL, SQL Server, MySQL, MongoDB, Cassandra, DynamoDB Data Modeling & Governance: Data Modeling, Data Lineage, Data Contracts, Data Quality Monitoring, SLA Management, Collibra, RBAC, Row-Level Security (RLS), Privacy & Compliance Business Intelligence & Analytics: Tableau, Power BI, Amazon QuickSight, SSRS, Adobe Analytics, Birst Machine Learning & AI: Scikit-learn, TensorFlow, Feature Engineering, Regression, Classification, Clustering, OpenAI GPT, Hugging Face Transformers, LangChain, Retrieval-Augmented Generation (RAG), Prompt Engineering, AI Agents, Vector Search, Responsible AI DevOps & Infrastructure: Git, GitHub, GitHub Actions, Docker, Kubernetes, Terraform, Jenkins, CI/CD Pipelines
PROFESSIONAL EXPERIENCE	
Intuit	Jan 2025 – Feb 2026 CA, USA
Data Engineer – Technical Data Analyst	- Architected and owned end-to-end real-time and batch data pipelines using PySpark, Databricks, SQL, and Delta Lake, processing millions of authentication events daily across three Identity & Authentication business domains. - Led the migration of enterprise analytics workloads from Amazon Redshift to Databricks Delta Lake, reducing cloud infrastructure costs by over \$100K annually while improving query performance, scalability, and pipeline reliability. - Designed scalable data models, ETL/ELT workflows, and trusted analytics datasets that powered product analytics, executive reporting, experimentation, KPI measurement, and business intelligence initiatives. - Established enterprise KPI frameworks, data contracts, SLA monitoring, and governance standards, creating a single source of truth and improving data consistency across Product, Engineering, and Analytics teams. - Partnered closely with Product Managers, Software Engineers, Data Scientists, and Analytics teams to enhance event instrumentation, customer journey tracking, experimentation datasets, and observability for A/B testing, funnel analysis, and guardrail metrics. - Built automated data quality validation, freshness monitoring, SLA tracking, and alerting frameworks, proactively detecting pipeline failures and improving downstream data reliability. - Developed interactive Tableau dashboards providing real-time visibility into authentication health, user behavior, experimentation results, and executive KPIs, significantly reducing ad hoc reporting requests. - Designed AI-powered data quality automation using RAG, LLM applications, prompt engineering, vector search, and agent orchestration, reducing manual data investigation efforts by approximately 30%, while implementing CI/CD pipelines with GitHub Actions to standardize deployments and engineering best practices.
Wyndham Hotels & Resorts	Mar 2022 – Jan 2025 NJ, USA
Data Engineer – Business Intelligence	- Designed and owned cloud-native data architecture processing 100M+ weekly records using AWS Glue, Spark, Python, and Snowflake, powering near-real-time analytics and reducing infrastructure costs by \$150K+ annually. - Architected and optimized enterprise ETL/ELT pipelines integrating booking, loyalty, customer, pricing, and behavioral data from multiple internal and third-party sources into centralized analytics platforms. - Developed scalable data models, semantic reporting layers, and standardized KPI frameworks, enabling consistent measurement of customer acquisition, retention, conversion, engagement, and executive business performance. - Partnered with Product Managers, Marketing, Operations, and Business stakeholders to define enterprise KPIs, deliver actionable insights, and improve operational decision-making by 20%. - Built interactive Tableau dashboards and performed funnel analysis, customer segmentation, retention analysis, and conversion optimization, increasing lead conversion by 10%, improving campaign ROI by 30%, and reducing manual reporting effort by 40%. - Optimized Snowflake performance through clustering, query tuning, CTE optimization, and efficient join strategies while integrating third-party pricing and market intelligence data, reducing dashboard latency by 30% and pricing decision time by 13 hours. - Established enterprise data quality validation, governance, pipeline observability, and security frameworks using Row-Level Security (RLS), Role-Based Access Control (RBAC), and privacy-compliant policies, reducing production incidents by 15% and improving data reliability.
Web DreamWorks Pvt. Ltd.	Jan 2018 – Dec 2020 India
Jr. Data Engineer	- Developed real-time event-driven data pipelines using Apache Kafka, Apache Spark, and Python, processing high-volume e-commerce transactions and improving data processing efficiency by 35%. - Built and optimized distributed ETL pipelines using Python, Pandas, NumPy, and PySpark, processing 50M+ records weekly across multiple business systems to support analytics and reporting. - Designed scalable data preprocessing, feature engineering, and machine learning pipelines using regression, clustering, and classification models, contributing to 20% revenue growth and 10% higher upsell conversion. - Designed and implemented A/B testing frameworks and statistical analysis to evaluate product features and customer behavior, increasing user engagement by 15% and transaction completion by 10%. - Collaborated with cross-functional teams to transform business requirements into data-driven solutions, delivering actionable insights through data analysis, customer segmentation, and predictive modeling. - Implemented infrastructure monitoring and workflow orchestration using Prometheus, Kubernetes, and Apache Airflow, reducing ETL workflow downtime by 50% through proactive alerting and resource optimization.
EDUCATION	
Master of Science in Computer Science	Dec – 2022
Pace University, New York, NY – USA	
CERTIFICATIONS	
AWS Certified Data Engineer	(Link)
AWS Certified Cloud Practitioner	(Link)
Databricks Certified Generative AI Engineer	(Link)